

# A “Hitchhiker’s” Guide to Fast and Efficient Data Reconstruction in Erasure-coded Data Centers

K. V. Rashmi, Nihar Shah, D. Gu,  
H. Kuang, D. Borthakur, K. Ramchandran



# Need for Redundant Storage in Data Centers


- Frequent unavailability events in data centers
  - unreliable components
  - software glitches, maintenance shutdowns, power failures, etc.
- Redundancy necessary for reliability and availability

# Popular Approach for Redundant Storage: Replication

- Distributed file systems used in data centers store multiple copies of data on different machines
- Machines typically chosen on different racks
  - to tolerate rack failures

E.g., Hadoop Distributed File System (HDFS) stores 3 replicas by default

# HDFS



# Massive Data Sizes: Need Alternative to Replication

- Small to moderately sized data: disk storage is inexpensive
  - replication viable
- No longer true for massive scales of operation
  - e.g., Facebook data warehouse cluster stores multiple tens of Petabytes (PBs)


“Erasure codes” are an alternative

# Erasure Codes in Data Centers

- Facebook data warehouse cluster
  - uses Reed-Solomon (RS) codes instead of 3-replication on a portion of the data
  - *savings of multiple Petabytes of storage space*

# Erasure Codes


## Replication



Overhead      2x

Fault tolerance: tolerates any one failure

## Reed-Solomon (RS) code



2x

tolerates any two failures

In general, erasure codes provide orders of magnitude higher reliability at much smaller storage overheads

# Outline


- Erasure Codes in Data Centers
  - HDFS
- Impact on the data center network
  - Problem description
- Our system: “Hitchhiker”
- Implementation and evaluation
  - Facebook data warehouse cluster
- Literature




# Outline

- Erasure Codes in Data Centers
  - HDFS
- Impact on the data center network
  - Problem description
- Our solution: “Hitchhiker”
- Implementation and evaluation
  - Facebook data warehouse cluster
- Literature

# Erasure codes in Data Centers: HDFS-RAID




Overhead: 3x




Overhead: 1.4x

(10, 4) Reed-Solomon code



# Erasure codes in Data Centers: HDFS-RAID




Overhead: 3x

Cannot tolerate  
many 3-failures



Overhead: 1.4x



(10, 4) Reed-Solomon code

- Any 10 blocks sufficient
- Can tolerate any 4-failures


# Outline

- Erasure Codes in Data Centers
  - HDFS
- Impact on the data center network
  - Problem description
- Our system: “Hitchhiker”
- Implementation and evaluation
  - Facebook data warehouse cluster
- Literature

# Impact on Data Center Network

Network Layer

Reconstruction Operations



Storage Layer

- **Degraded Reads**

- requesting currently unavailable data
- on-the-fly reconstruction

- **Recovery**


- periodically replace unavailable blocks
- to ensure desired level of reliability

# Impact on Data Center Network


RS codes significantly increase network usage during reconstruction

# Impact on Data Center Network

## Replication



## Reed-Solomon code




Network transfer & disk IO

= (#data-blocks) x (size of data to be reconstructed)

In (10, 4) RS, it is 10x


# Impact on Data Center Network



Burdens the already oversubscribed  
Top-of-Rack and higher level switches



# Impact on Data Center Network: Facebook Data Warehouse Cluster



- Multiple PB of Reed-Solomon encoded data
- Median of **180 TB transferred across racks** per day for RS reconstruction  $\approx$  **5 times** that under 3-replication

# RS codes: The Good and The Bad

- Maximum possible fault-tolerance for given storage overhead
  - storage-capacity optimal
  - (“*maximum-distance-separable*” in coding theory parlance)
- Flexibility in choice of parameters
  - Supports any number of data and parity blocks
- Not designed to handle reconstruction operations efficiently
  - negative impact on the network

# RS codes: The **Goal** and The Bad

- Maximum possible fault-tolerance for given storage overhead
  - storage-capacity efficiency
  - (“*maximum-distance*” theory parlance)
- Flexibility in choice of parameters
  - Supports any number of data and parity blocks
- Not designed to handle reconstruction operations efficiently
  - negative impact on the network

**Maintain**

**Improve**

# Goal

To build a system with:

Maintain

Same (optimal) storage requirement and fault tolerance

Same (complete) flexibility in choice of design parameters

Improve

Reduced data transfer across network and reduced IO from disk during reconstruction

# Hitchhiker

Is a system with:

Maintain

Same (optimal) storage requirement and fault tolerance ☒

Same (complete) flexibility in choice of design parameters ☒

Improve


25 to 45% less network transfers and disk IO during reconstruction ☒

# Outline

- Erasure Codes in Data Centers
  - HDFS
- Impact on the data center network
  - Problem description
- **Our system: “Hitchhiker”**
- Implementation and evaluation
  - Facebook data warehouse cluster
- Literature


# At an Abstract Level

## HITCHHIKER



# Hitchhiker's Erasure Code: Toy Example


Start with the RS code





# Intermediate Code


Add information from first group on to  
parities of the second group



No extra storage


# Storage-optimality of Intermediate Code

Retains failure tolerance of RS codes:  
can tolerate failure of any 2 nodes




# Final Code

Invertible operation *within* a block




# Final Code

Invertible operations *within* blocks do not change storage or fault tolerance



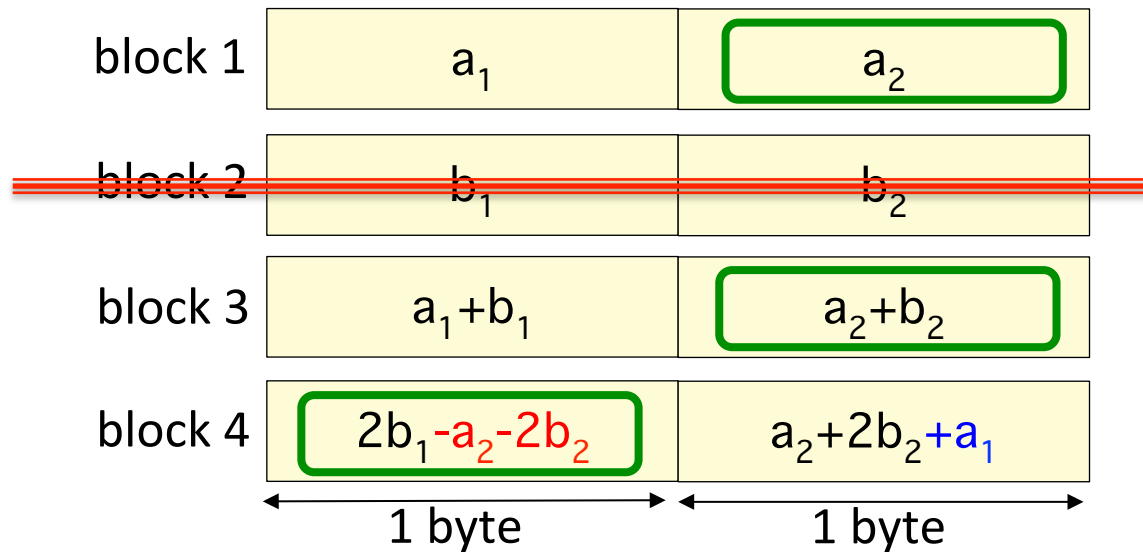
# Efficient Reconstruction

Data transferred: only 3 bytes  
(instead of 4 bytes as in RS)



# Efficient Reconstruction

Data transferred: only 3 bytes  
(instead of 4 bytes as in RS)



# Hitchhiker's Erasure Code

- Builds on top of RS codes
- Uses our theoretical framework of “Piggybacking”\*
- Three versions
  - XOR
  - XOR+
  - non-XOR

\* K.V. Rashmi, Nihar Shah, K. Ramchandran, “A Piggybacking Design Framework for Read-and Download-efficient Distributed Storage Codes”, in IEEE International Symposium on Information Theory, 2013.


## Hop-and-couple (disk layout)

- Way of choosing which bytes to mix
  - couples bytes farther apart in block
  - to minimize fragmentation of reads during reconstruction
- Translate savings in network-transfer to savings in disk-IO as well
  - By making reads contiguous



RS vs Hitchhiker from the Network's Perspective...


# Data Transfer during Reconstruction in RS-based System



**Transfer: 10 full blocks**  
Connect to 10 machines

# Data Transfer during Reconstruction in Hitchhiker


Reconstruction of data blocks 1-9:



Transfer: 2 full blocks + 9 half blocks (= 6.5 blocks total)  
Connect to 11 machines

# Data Transfer during Reconstruction in Hitchhiker

Reconstruction of block 10:



Transfer: 13 half blocks (= 6.5 blocks total)

Connect to 13 machines

# Outline

- Erasure Codes in Data Centers
  - HDFS
- Impact on the data center network
  - Problem description
- Our system: “Hitchhiker”
- Implementation and evaluation
  - Facebook data warehouse cluster
- Literature


# Implementation & Evaluation Setup (1)

- Implemented on top of HDFS-RAID
  - erasure coding module in HDFS based on RS
  - used in the Facebook data warehouse cluster
- Deployed and tested on a 60 machine test cluster at Facebook
  - verified 35% reduction in the network transfers during reconstruction

# Implementation & Evaluation Setup (2)

- Evaluation of timing metrics on the Facebook data warehouse cluster in production
  - under real-time production traffic and workloads
  - using Map-Reduce to run encoding and reconstruction jobs, just as HDFS-RAID

# Decoding Time



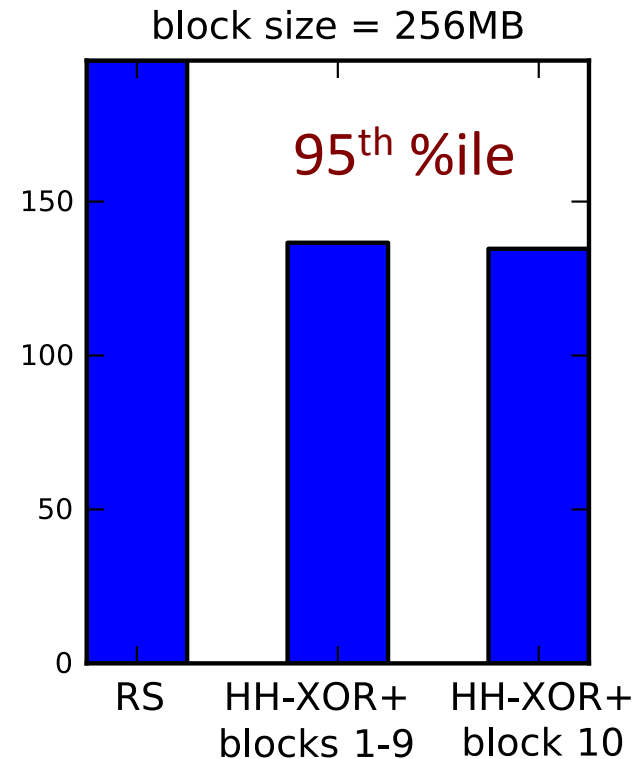

36% reduction

- RS decoding on only half portion of the blocks
- Faster computation for degraded reads and recovery
- XOR versions: 25% lesser than non-XOR




# Read & Transfer Time

System	Data transfer	Connectivity (#machines)
RS	2.56 GB	10
HH blocks 1-9	1.67 GB	11
HH block 10	1.67 GB	13



- Read & transfer time **30% lower** in Hitchhiker (HH)
- Similar reduction for other block sizes as well

# Encoding Time



72% higher

Benefits outweigh higher encoding cost in many systems (e.g., HDFS):

- encoding is one time operation
- often run as a background job
- does not fall along any critical path

# Outline

- Impact on the data center network
  - Problem description
- Our system: “Hitchhiker”
- Implementation and evaluation
  - Facebook data warehouse cluster
- Literature

# Existing Systems

- Need additional storage
  - Huang et al. (Windows Azure) 2012, Sathiamoorthy et al. (Xorbas) 2013, Esmaili et al. (CORE) 2013
    - Add additional parities to reduce download
  - Hu et al. (NCFS 2011)
- Highly restricted parameters
  - Khan et al. (Rotated-RS) 2012:  $\#parity \leq 3$
  - Xiang et al., Wang et al. 2010, Hu (NCCloud) et al. 2012:  $\#parity \leq 2$
  - Hitchhiker performs as good or better for these restricted settings as well

# Hitchhiker: Summary

## Code metrics:

Storage requirement	Same (optimal)
Supported parameters	All
Fault tolerance	Same (optimal)

## Reconstruction:

Network transfers	35% less
Disk IO	35% less
Data read and transfer time (median)	31.8% less
Data read and transfer time (95th %ile)	30.2% less
Computation time (median)	36.1% less

## Encoding:

Encoding time (median)	72.1% more
------------------------	------------

Thanks!

# Backup Slides

# Hop-and-Couple

- Technique to pair bytes under Hitchhiker's erasure code
- Makes disk reads during reconstruction contiguous

