

# Traffic Engineering with Forward Fault Correction (FFC)

Hongqiang “Harry” Liu,

Srikanth Kandula, Ratul Mahajan, Ming Zhang,  
David Gelernter (Yale University)

Microsoft®  
**Research**

# Cloud services require large network capacity




Cloud Services

Growing traffic



Cloud Networks



Expensive  
(e.g. cost of WAN: \$100M/year)


# TE is critical to effectively utilizing networks

## Traffic Engineering




- Microsoft SWAN
- Google B4
- .....

- Devoflow
- MicroTE
- .....

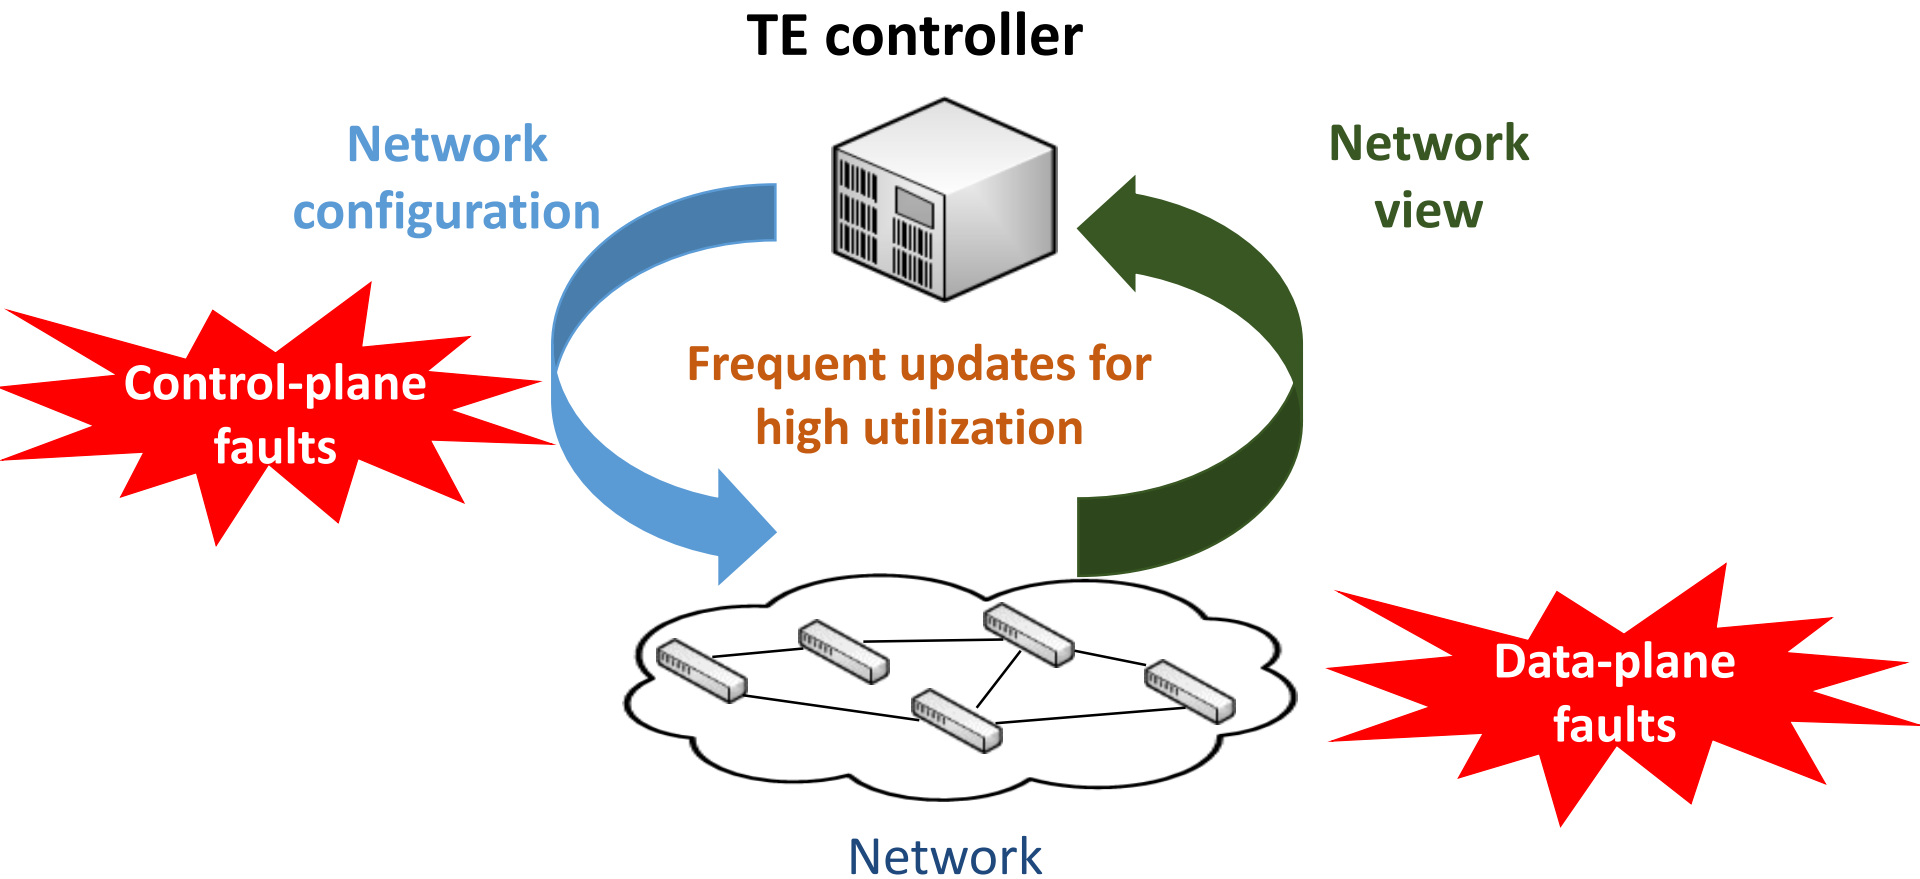


WAN Network



Datacenter Network

But, TE is also vulnerable to faults



# Control plan faults

## Failures or long delays to configure a network device

TE Controller




TE configurations



RPC failure

Switch



Firmware bugs

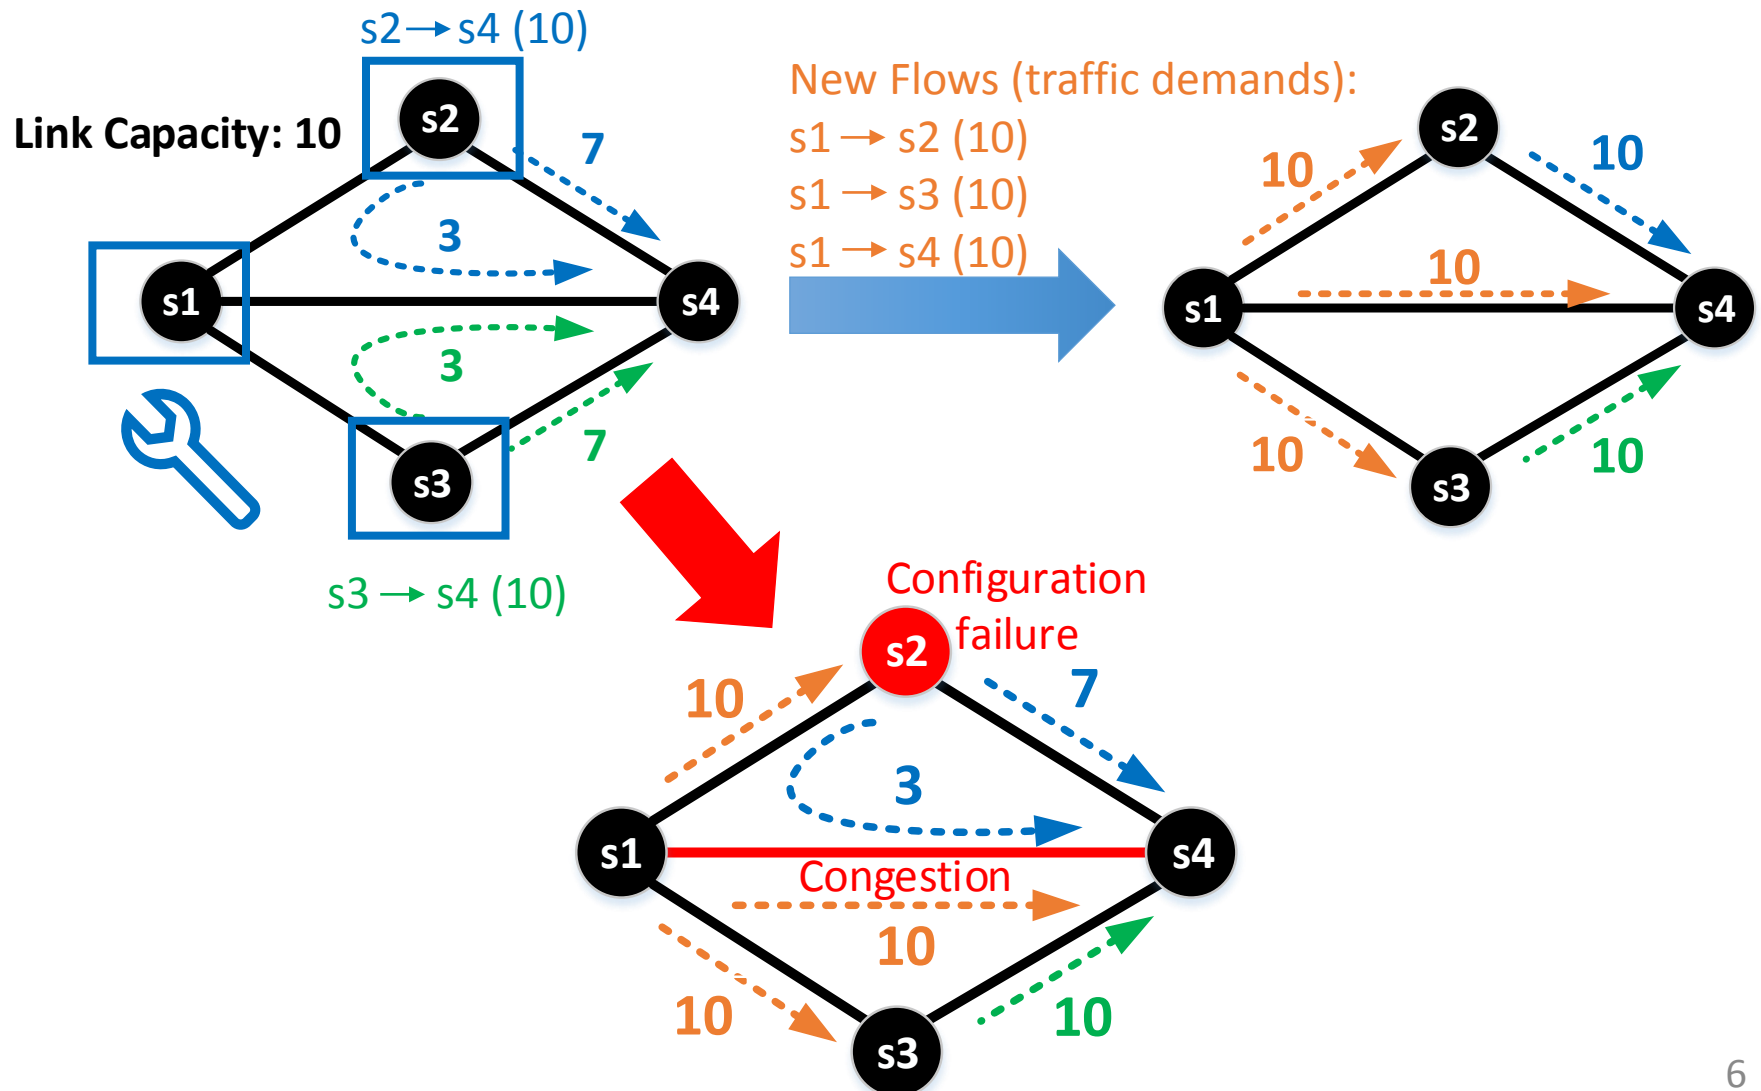


Overloaded CPU




Memory shortage

# Congestion due to control plane faults

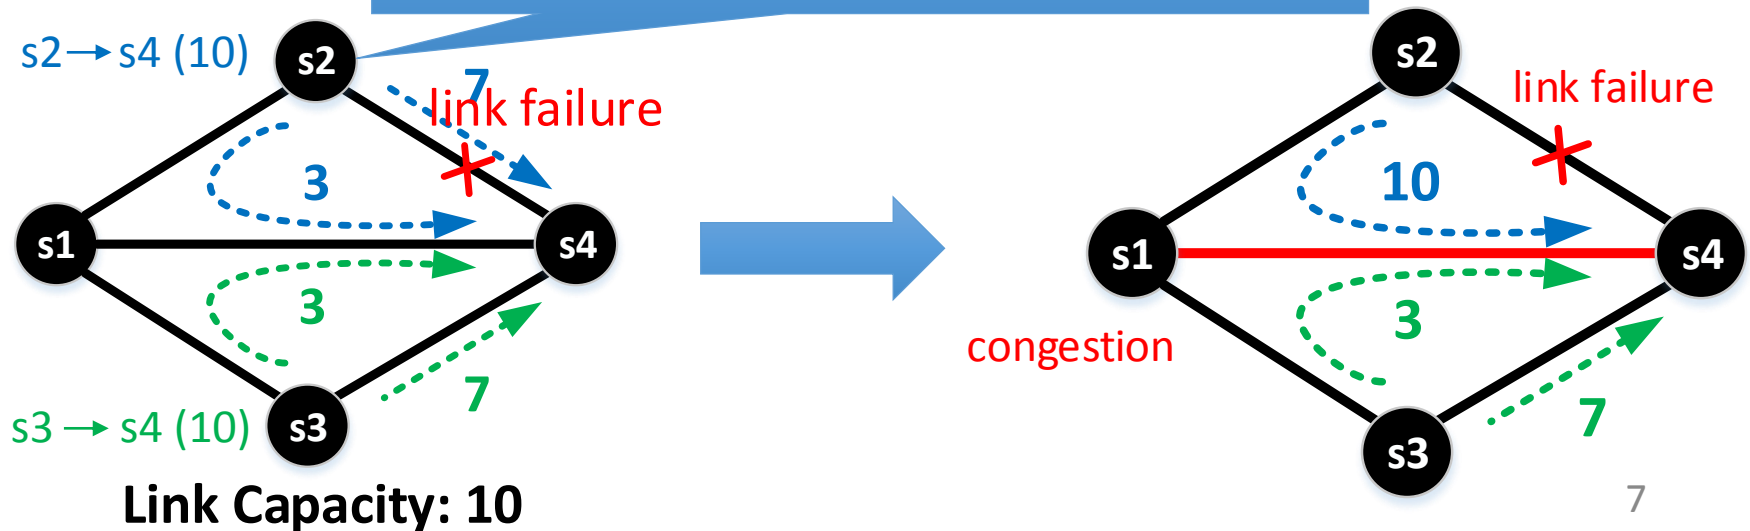


# Data plane faults

## Link and switch failures



**Rescaling:** Sending traffic proportionally to residual paths




# Control and data plane faults in practice

## In production networks:

- Faults are common.
- Faults cause severe congestion.

Control plane:  
fault rate = **0.1%** -- **1%** per TE update.

Data plane:  
fault rate = **25%** per 5 minutes.





# State of the art for handling faults

- Heavy over-provisioning

Big loss in throughput

- Reactive handling of faults

- Control plane faults: retry
- Data plane faults: re-compute TE and update networks

Cannot prevent  
congestion

Slow  
(seconds -- minutes)


Blocked by control  
plane faults



How about handling congestion  
**proactively?**


# Forward fault correction (FFC) in TE

- [Bad News] Individual faults are **unpredictable**.
- [Good News] Simultaneous #faults is **small**.



FEC guarantees no information loss  
under up to  $k$  arbitrary packet drops.

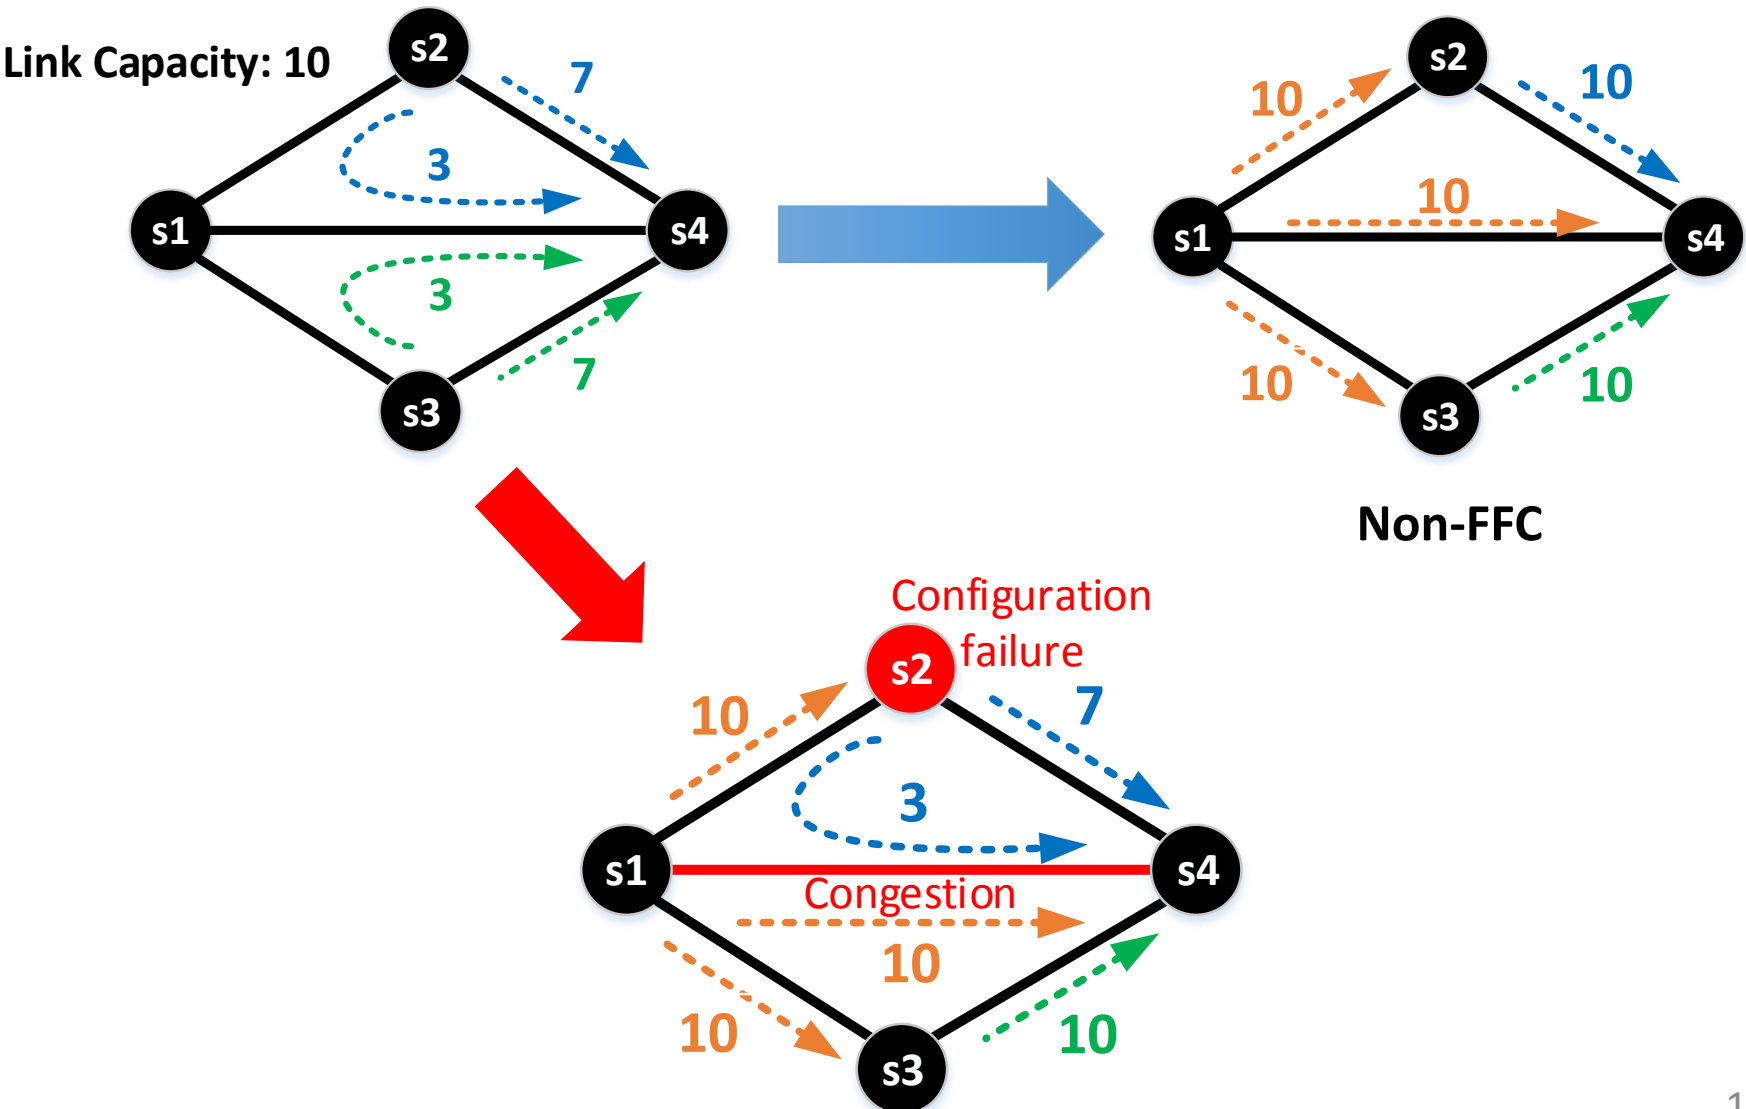
with careful data encoding



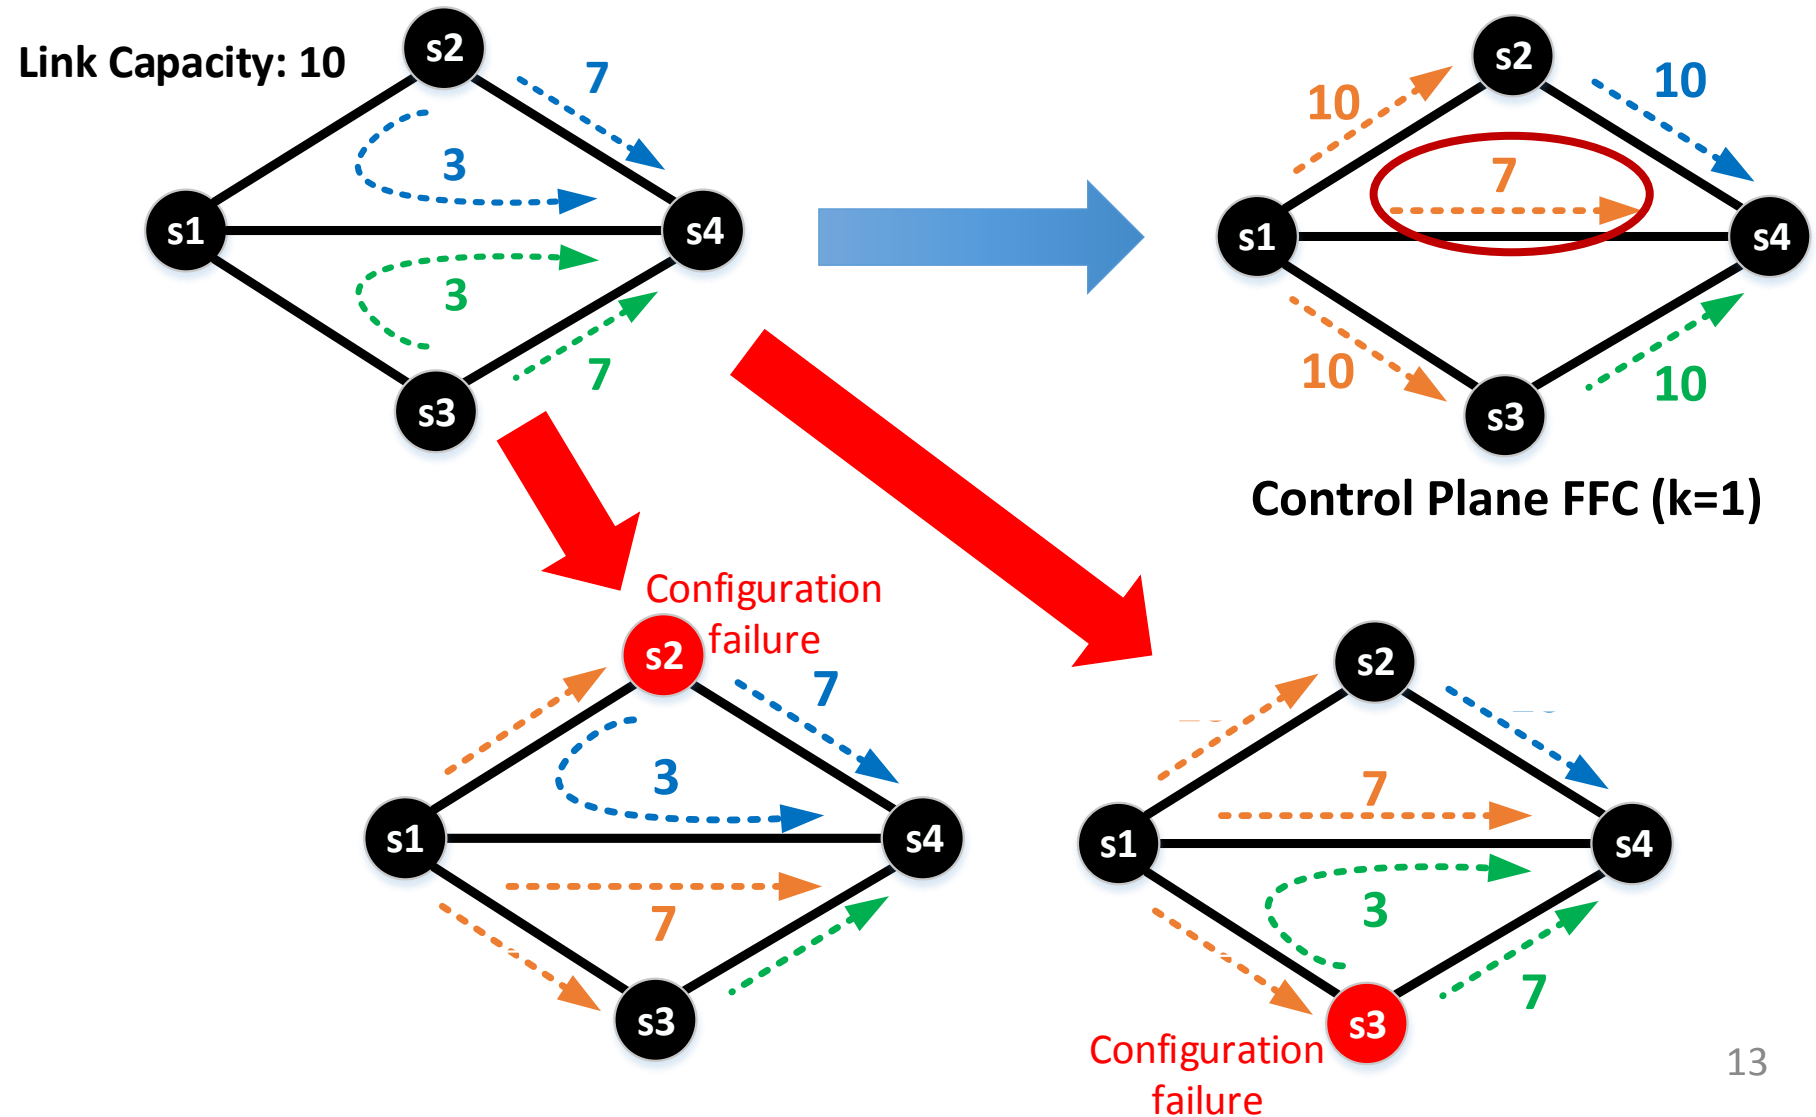
FFC guarantees no congestion  
under up to  $k$  arbitrary faults.

with careful traffic distribution


# Example: FFC for control plane faults



# Example: FFC for control plane faults




# Trade-off: network utilization vs. robustness




Non-FFC

*Throughput: 50*



K=1  
(Control Plane FFC)

*Throughput: 47*



K=2  
(Control Plane FFC)

*Throughput: 44*

# Systematically realizing FFC in TE

## **Formulation:**

*How to merge FFC into existing TE framework?*

## **Computation:**

*How to find FFC-TE efficiently?*

# Basic TE linear programming formulations

**TE decisions:**  $\left\{ \begin{array}{l} \text{Sizes of flows} \\ \text{Traffic on paths} \end{array} \right.$

**TE objective:** Maximizing throughput

**Basic TE constraints:**  $\left\{ \begin{array}{l} \text{Deliver all granted flows} \\ \text{No overloaded link} \\ \dots \end{array} \right.$



**FFC constraints:**

No overloaded link up to  $\left\{ \begin{array}{l} k_c \text{ control plane faults} \\ k_e \text{ link failures} \\ k_v \text{ switch failures} \end{array} \right.$

**LP formulations**

$b_f$

$l_{f,t}$


max.  $\sum_{\forall f} b_f$

s.t.  $\forall f: \sum_{\forall t} l_{f,t} \geq b_f$   
 $\forall e: \sum_{\forall f} \sum_{\forall t \ni e} l_{f,t} \leq c_e$   
...





# Formulating control plane FFC



$f_1$ 's load in old TE

$f_2$ 's load in new TE

Fault on  $f_1$ :  $l_1^{old} + l_2^{new} + l_3^{new} \leq \text{link cap}$

Fault on  $f_2$ :  $l_1^{new} + l_2^{old} + l_3^{new} \leq \text{link cap}$

Fault on  $f_3$ :  $l_1^{new} + l_2^{new} + l_3^{old} \leq \text{link cap}$

$\begin{pmatrix} 3 \\ 1 \end{pmatrix}$

Challenge: too many constraints

With  $n$  flows and FFC protection  $k$ :


**#constraints** =  $\binom{n}{1} + \dots + \binom{n}{k}$  for each link.

# An efficient and precise solution to FFC

## Our approach:

A **lossless** compression from  $O(\binom{n}{k})$  constraints to  $O(kn)$  constraints.

Total load under faults  
 $\leq$  link capacity



Total additional load due to faults  
 $\leq$  link spare capacity

$x_i$ : additional load due to **fault- $i$**

Given  $X = \{x_1, x_2, \dots, x_n\}$ , FFC requires that  
the sum of **arbitrary  $k$**  elements in  $X$  is  $\leq$  link spare capacity

$O(\binom{n}{k})$




Define  $y_m$  as the  **$m$ th largest** element in  $X$ :

$$\sum_{m=1}^k y_m \leq \text{link spare capacity}$$

$O(1)$

Expressing  $y_m$  with  $X$ ?

# Sorting network



$$y_1 + y_2 \leq \text{link spare capacity}$$


- **Complexity:**  $O(kn)$  additional variables and constraints.
- **Throughput:** optimal in control-plane and data plane if paths are disjoint.

# FFC extensions

- Differential protection for different traffic priorities
- Minimizing congestion risks without rate limiters
- Control plane faults on rate limiters
- Uncertainty in current TE
- Different TE objectives (e.g. max-min fairness)
- ...

# Evaluation overview

- Testbed experiment
  - FFC can be implemented in commodity switches
  - FFC has no data loss due to congestion under faults
- Large-scale simulation



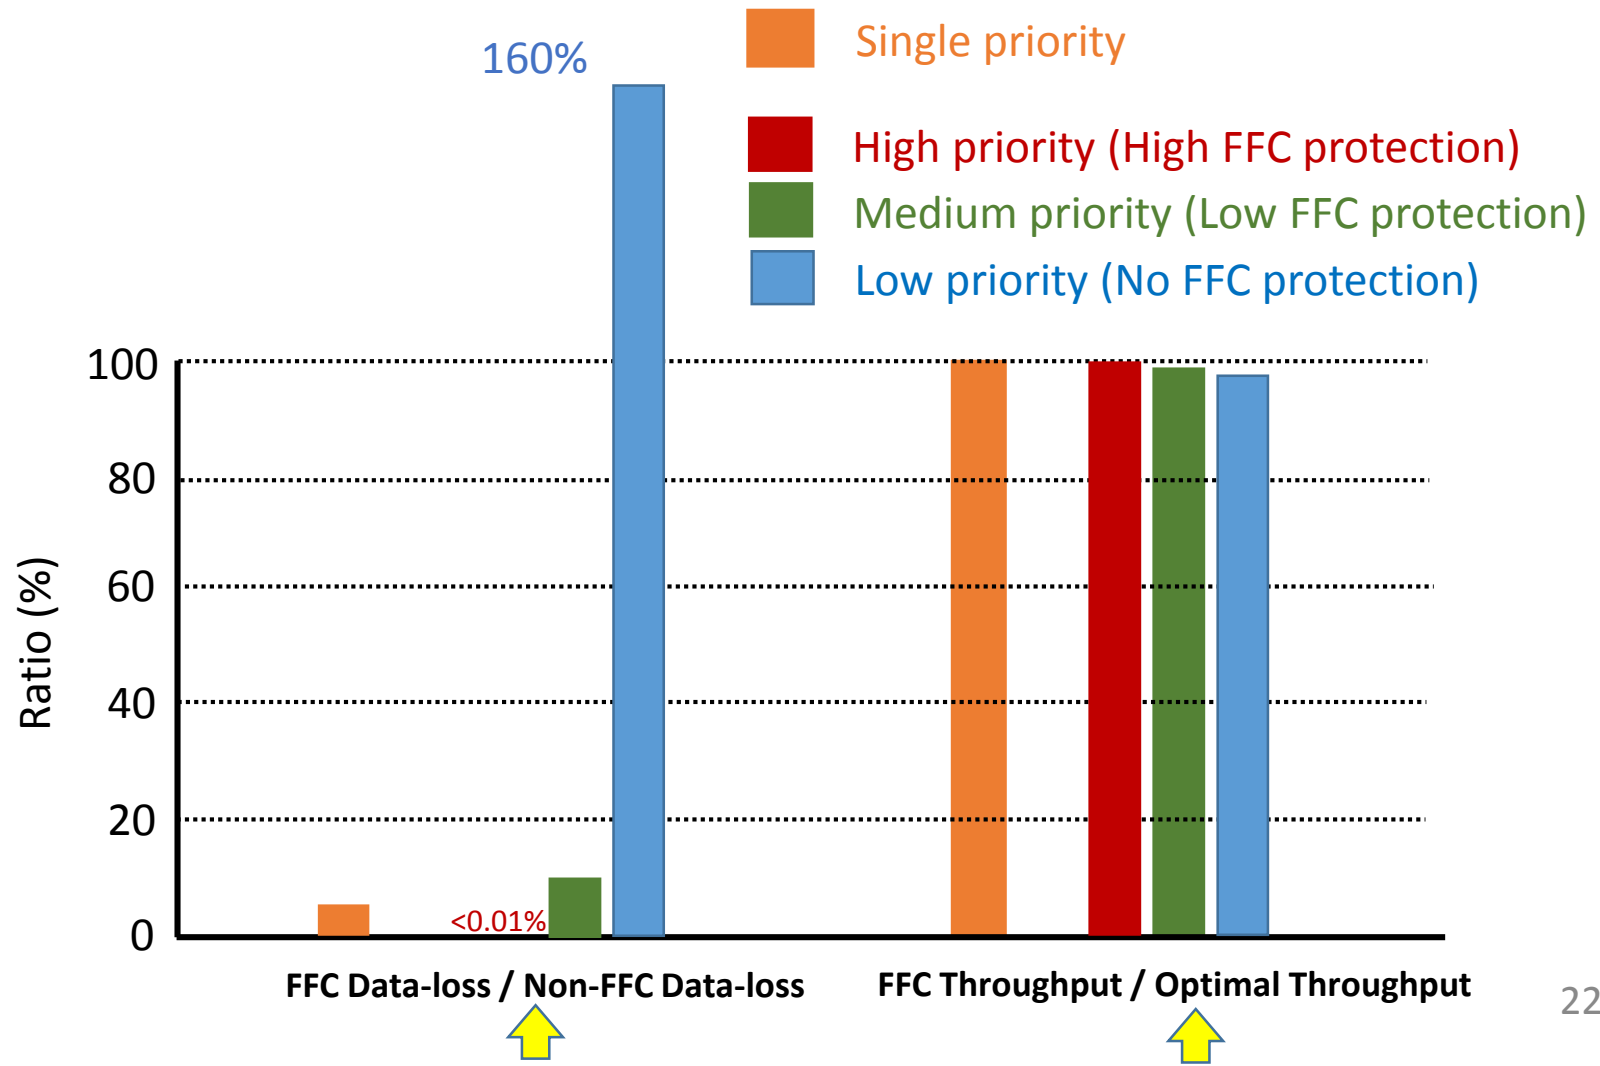
Injecting faults based  
on real failure reports

A WAN network with  $O(100)$   
switches and  $O(1000)$  links

Single priority traffic in a  
well-provisioned network

Multiple priority traffic in a  
well-utilized network

# FFC prevents congestion with negligible throughput loss



# Conclusions

- Centralized TE is critical to high network utilization but is vulnerable to control and data plane faults.
- FFC proactively handle these faults.
  - Guarantee: no congestion when **#faults  $\leq k$** .
  - Efficiently computable with low throughput overhead in practice.

***FFC***

High risk of  
congestion



Heavy network  
over-provisioning