SDX: A Software-Defined Internet Exchange

Arpit Gupta

Laurent Vanbever, Muhammad Shahbaz, Sean Donovan,
Brandon Schlinker, Nick Feamster, Jennifer Rexford,
Scott Shenker, Russ Clark, Ethan Katz-Bassett

Georgia Tech, Princeton University, UC Berkeley, USC

The Interdomain Ecosystem is Evolving ...

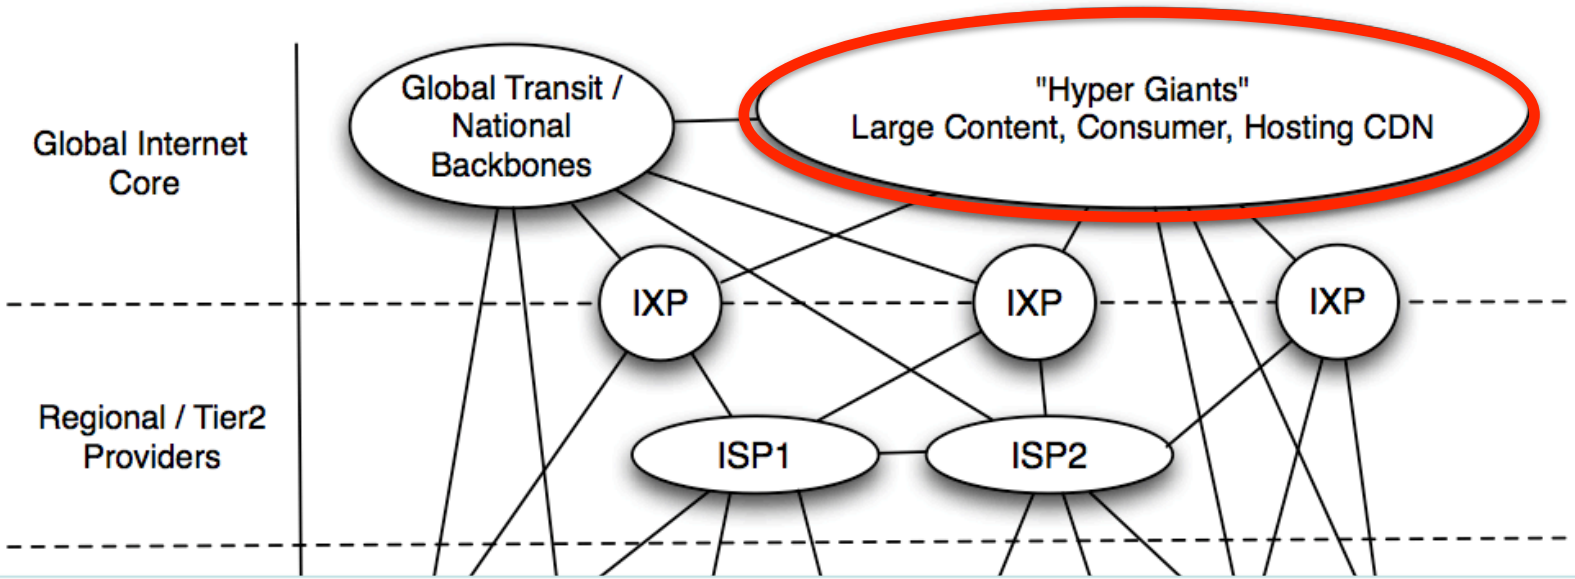

Flatter and densely interconnected Internet*

*Labovitz et al., *Internet Inter-Domain Traffic*, SIGCOMM 2010

…But BGP is Not

- Routing **only on destination IP prefixes**
  (No customization of routes by application, sender)
- Can only influence **immediate neighbors**
  (No ability to affect path selection remotely)
- **Indirect** control over data-plane forwarding
  (Indirect mechanisms to influence path selection)

How to overcome BGP's limitations?

SDN for Interdomain Routing

- Forwarding on **multiple header fields** (not just destination IP prefixes)
- Ability to **control entire networks** with a single software program (not just immediate neighbors)
- **Direct control** over data-plane forwarding (not indirect control via control-plane arcana)

How to incrementally deploy SDN for Interdomain Routing?

Deploy SDN at Internet Exchanges

- **Leverage:** SDN deployment even at single IXP can yield benefits for tens to hundreds of ISPs
- **Innovation hotbed:** Incentives to innovate as IXPs on front line of peering disputes
- **Growing in numbers:** ~100 new IXPs established in past three years*

*https://prefix.pch.net/applications/ixpdir/summary/growth/

Background: Conventional IXPs

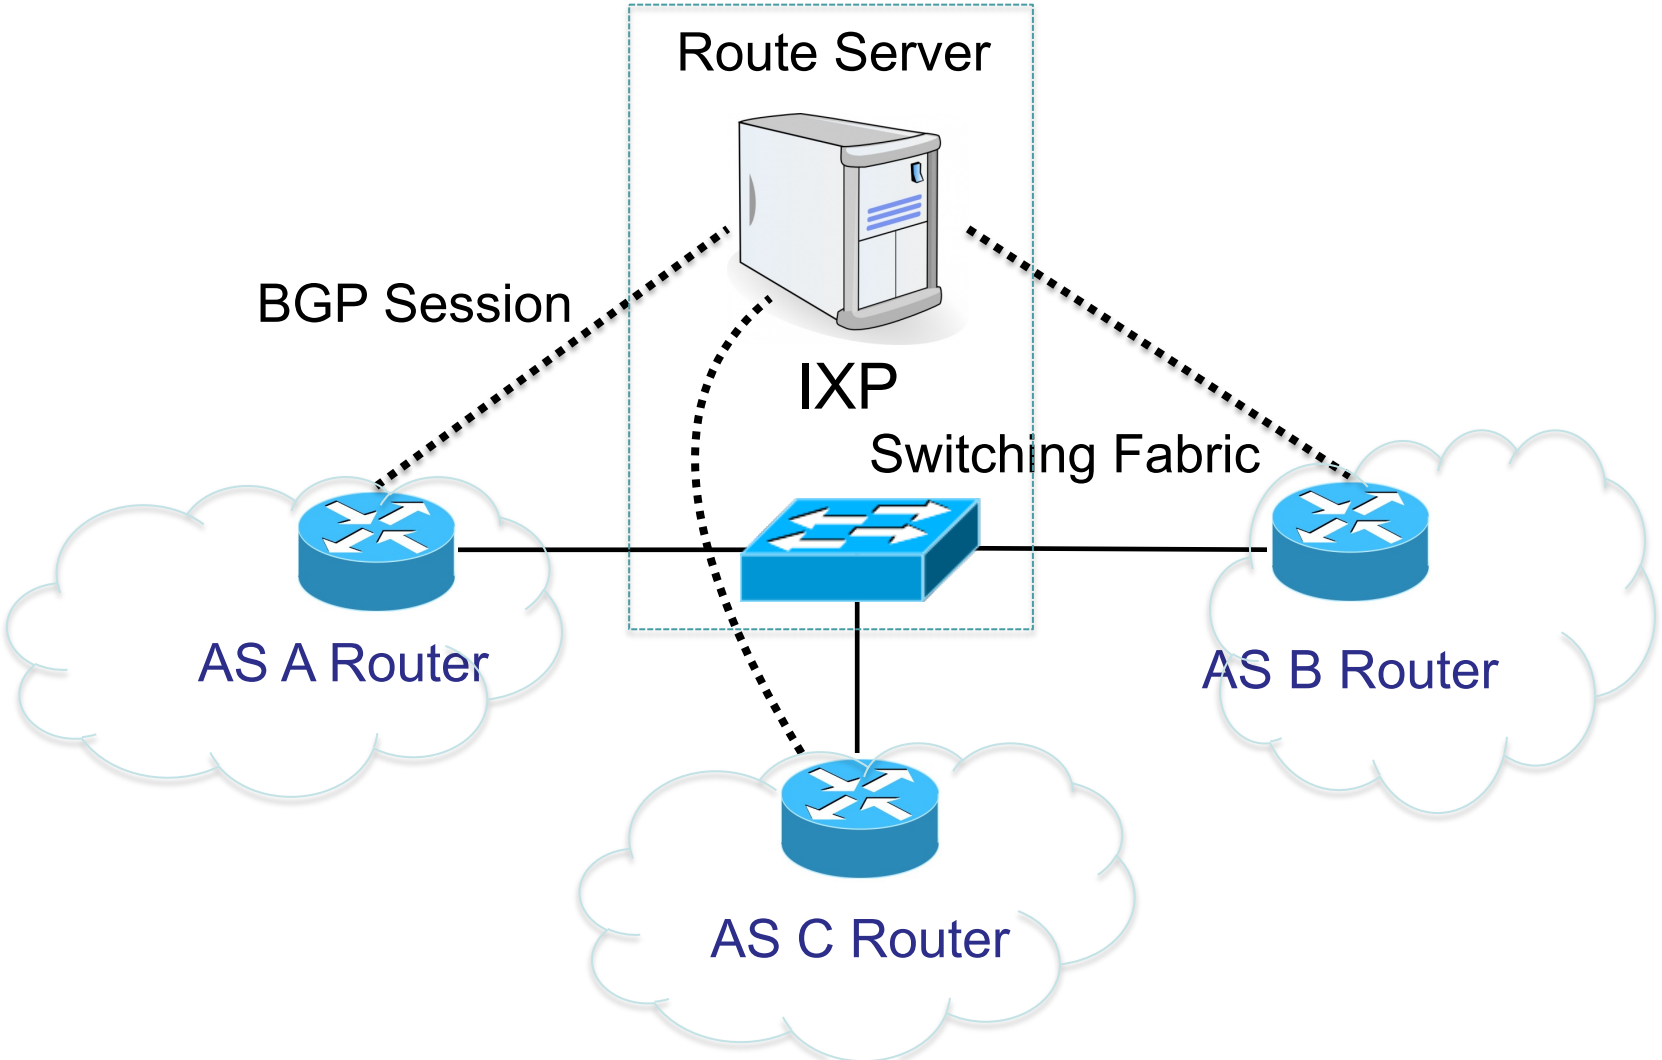

SDX = SDN + IXP

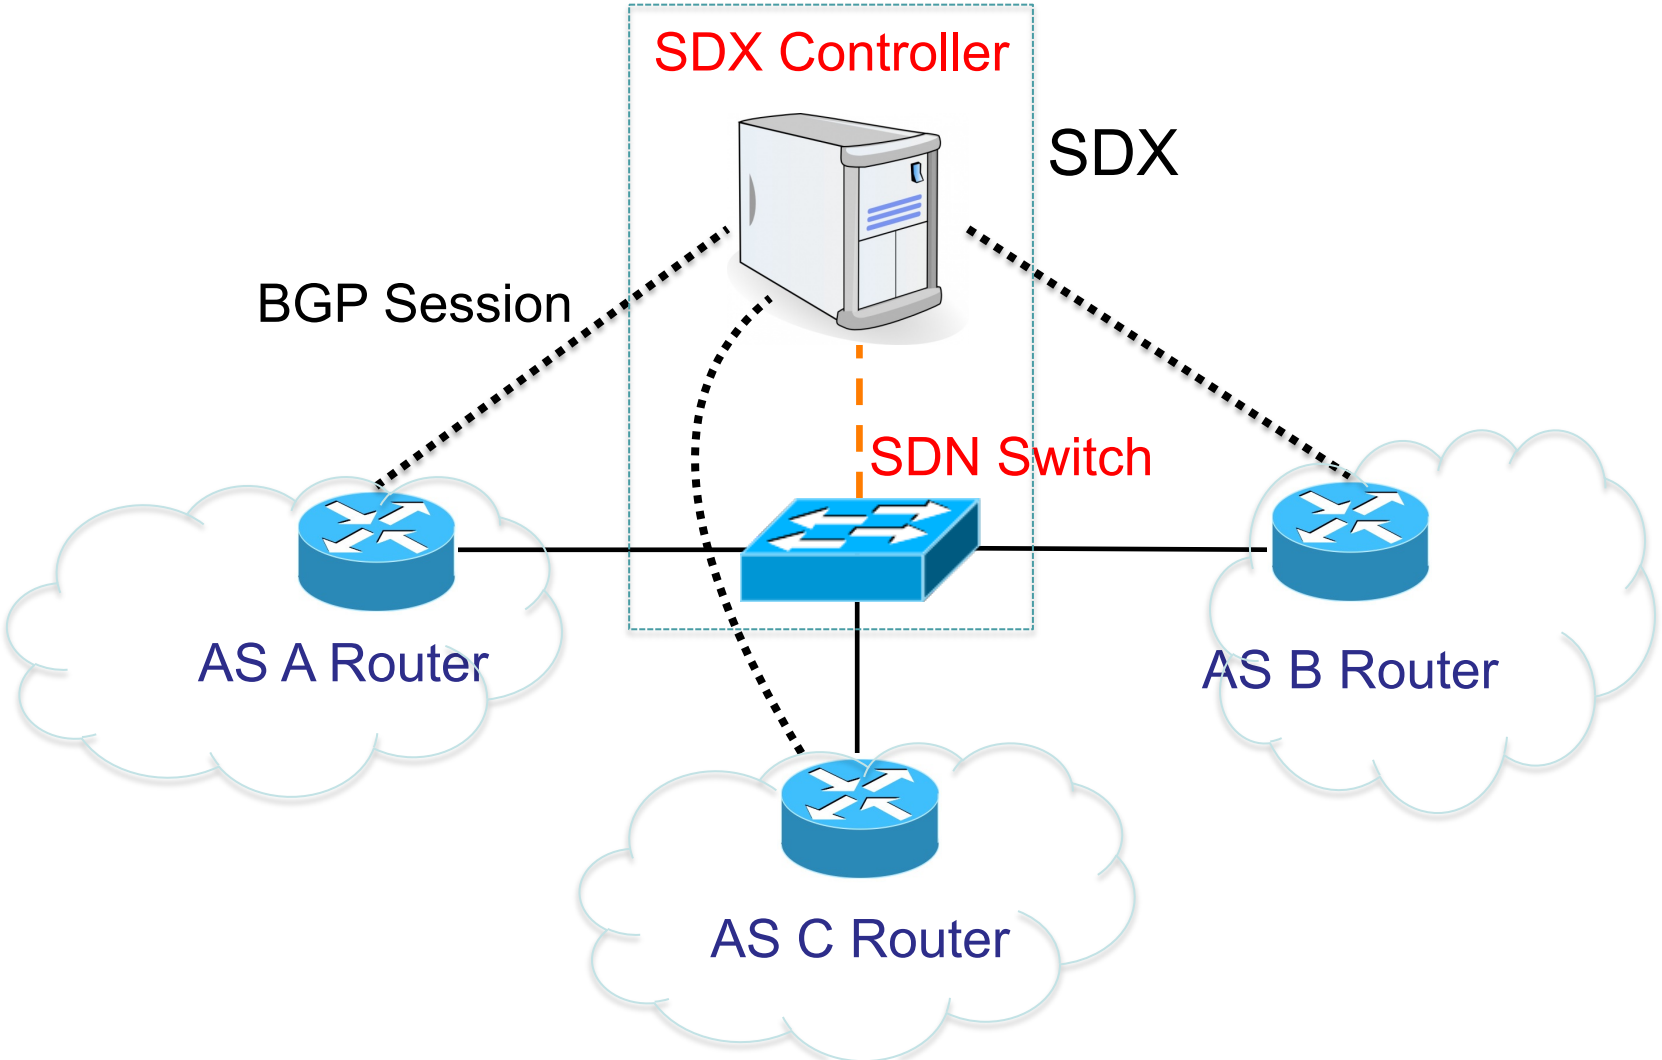
SDX Controller
SDX
BGP Session
SDN Switch
AS A Router
AS B Router
AS C Router

SDX Opens Up New Possibilities

- More flexible **business relationships**
  - Make peering decisions based on time of day, volume of traffic & nature of application
- More direct & flexible **traffic control**
  - Define fine-grained traffic engineering policies
- Better **security**
  - Prefer “more secure” routes
  - Automatically blackhole attack traffic

Use Case: Inbound Traffic Engineering

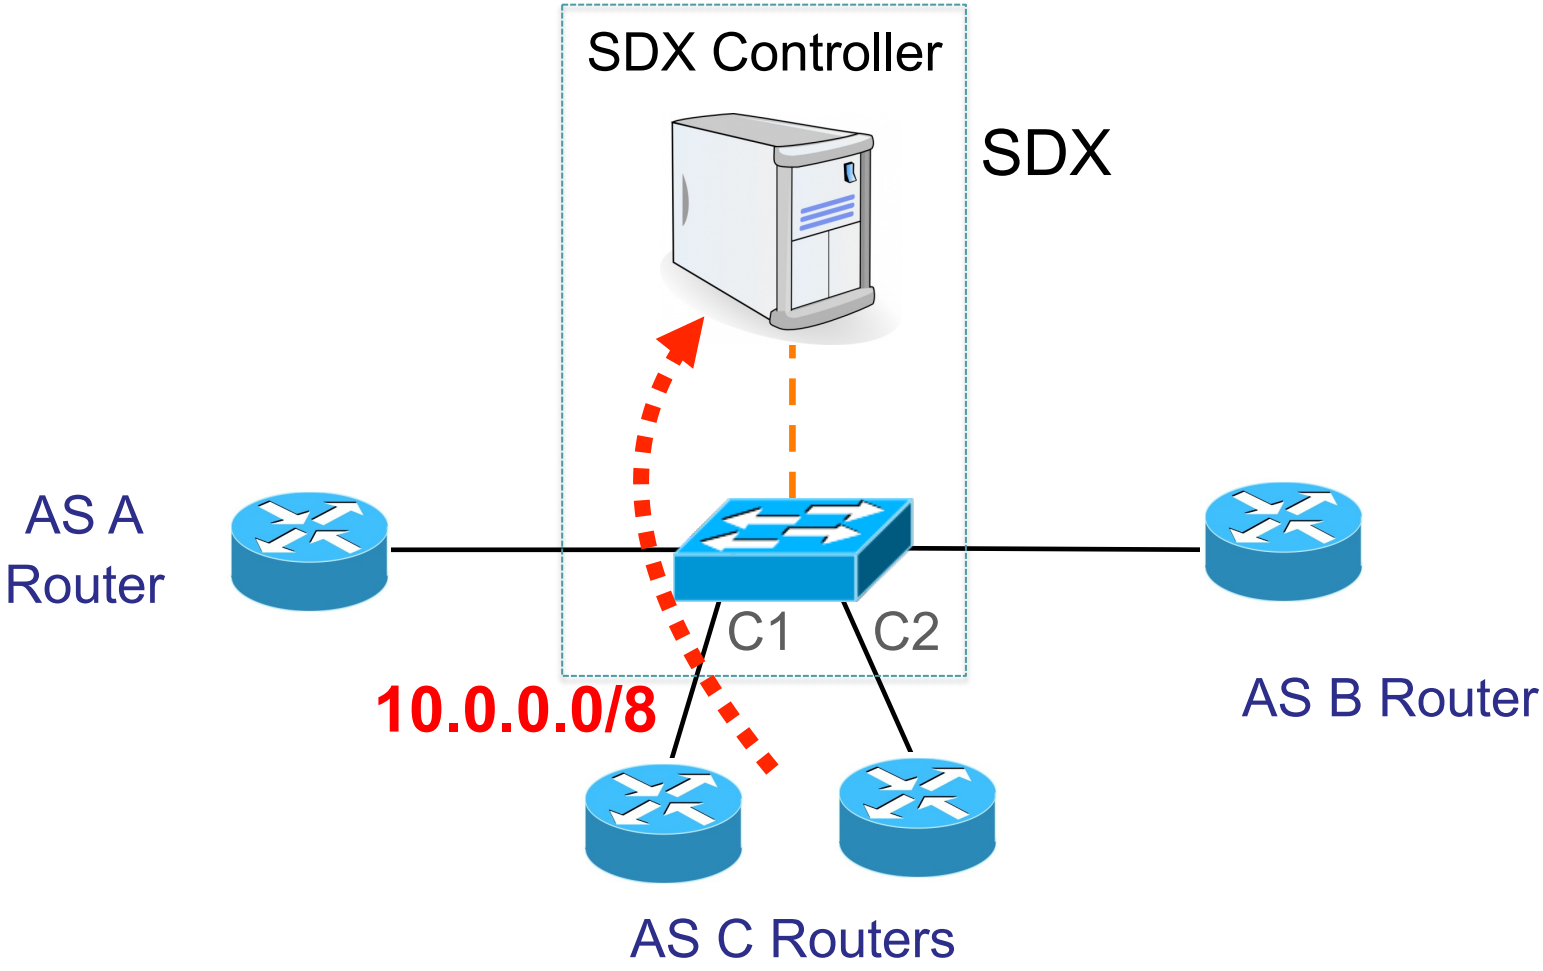

Use Case: Inbound Traffic Engineering

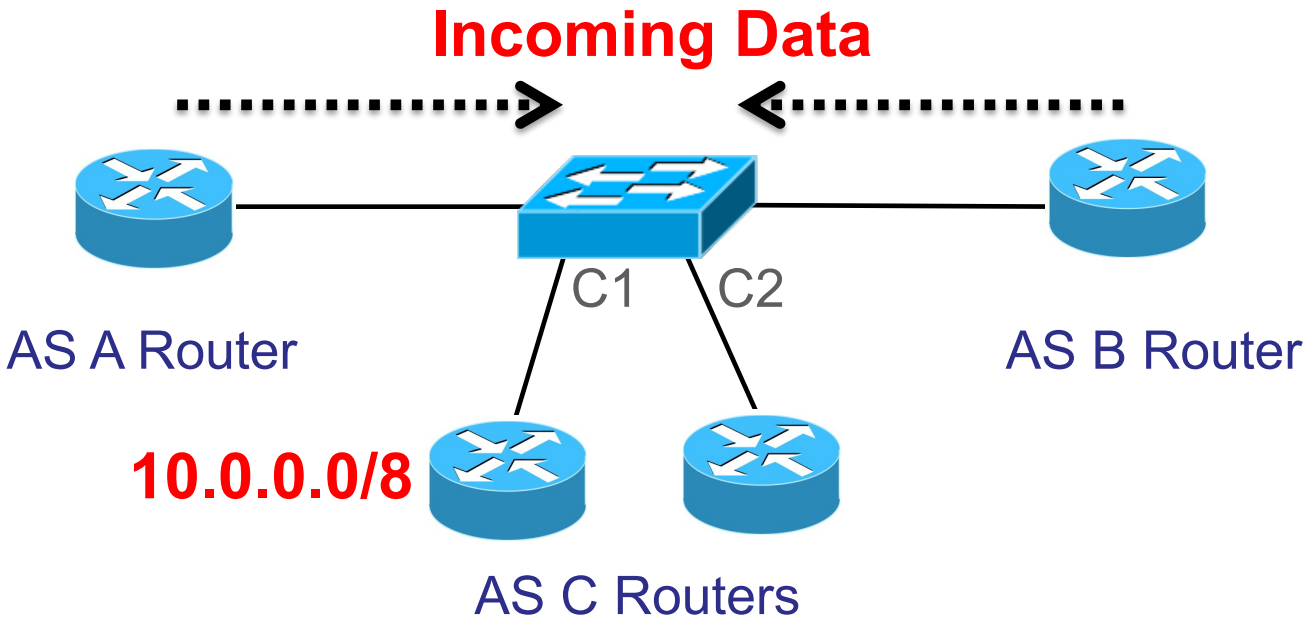

Incoming Traffic	Out Port	Using BGP	Using SDX
dstport = 80	C1		

Use Case: Inbound Traffic Engineering

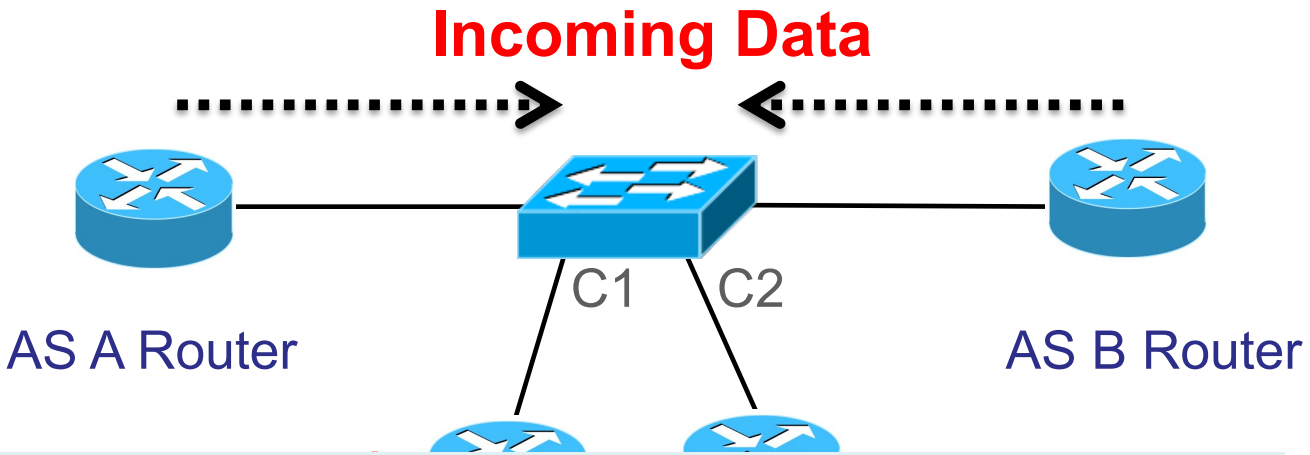

Fine grained policies not possible with BGP

Incoming Traffic	Out Port	Using BGP	Using SDX
dstport = 80	C1	?	

Use Case: Inbound Traffic Engineering

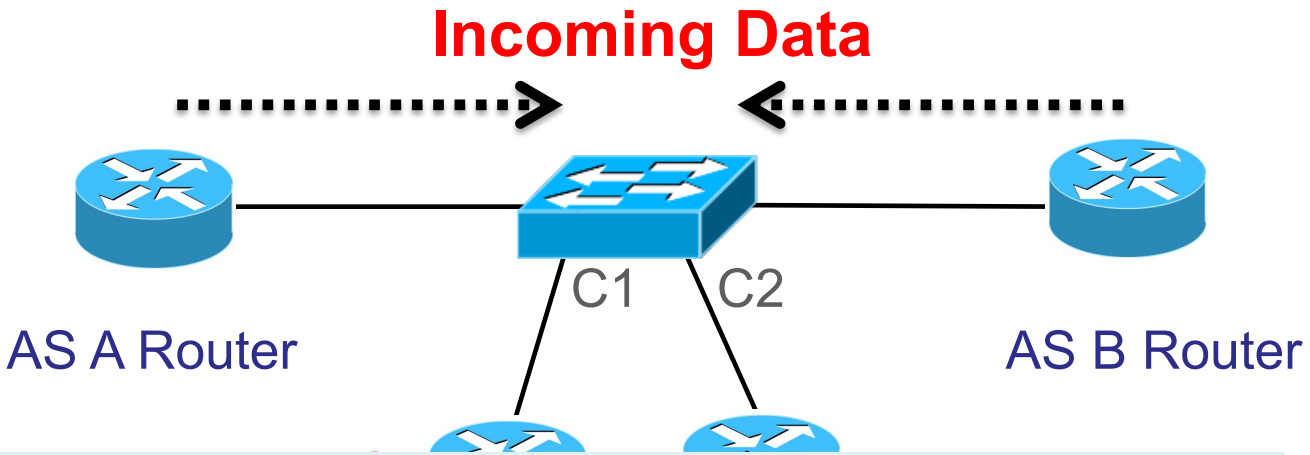

Enables fine-grained traffic engineering policies

Incoming Traffic	Out Port	Using BGP	Using SDX
dstport = 80	C1	?	match(dstport =80)→ fwd(C1)

Building SDX is Challenging

- Programming **abstractions**
  - How networks define SDX policies and how are they combined together?
- **Interoperation** with BGP
  - How to provide flexibility w/o breaking global routing?
- **Scalability**
  - How to handle policies for hundreds of peers, half million prefixes and matches on multiple header fields?

Building SDX is Challenging

- Programming **abstractions**
  - How networks define SDX policies and how are they combined together?
- **Interoperation** with BGP
  - How to provide flexibility w/o breaking global routing?
- **Scalability**
  - How to handle policies for hundreds of peers, half million prefixes and matches on multiple header fields?

Directly Program the SDX Switch

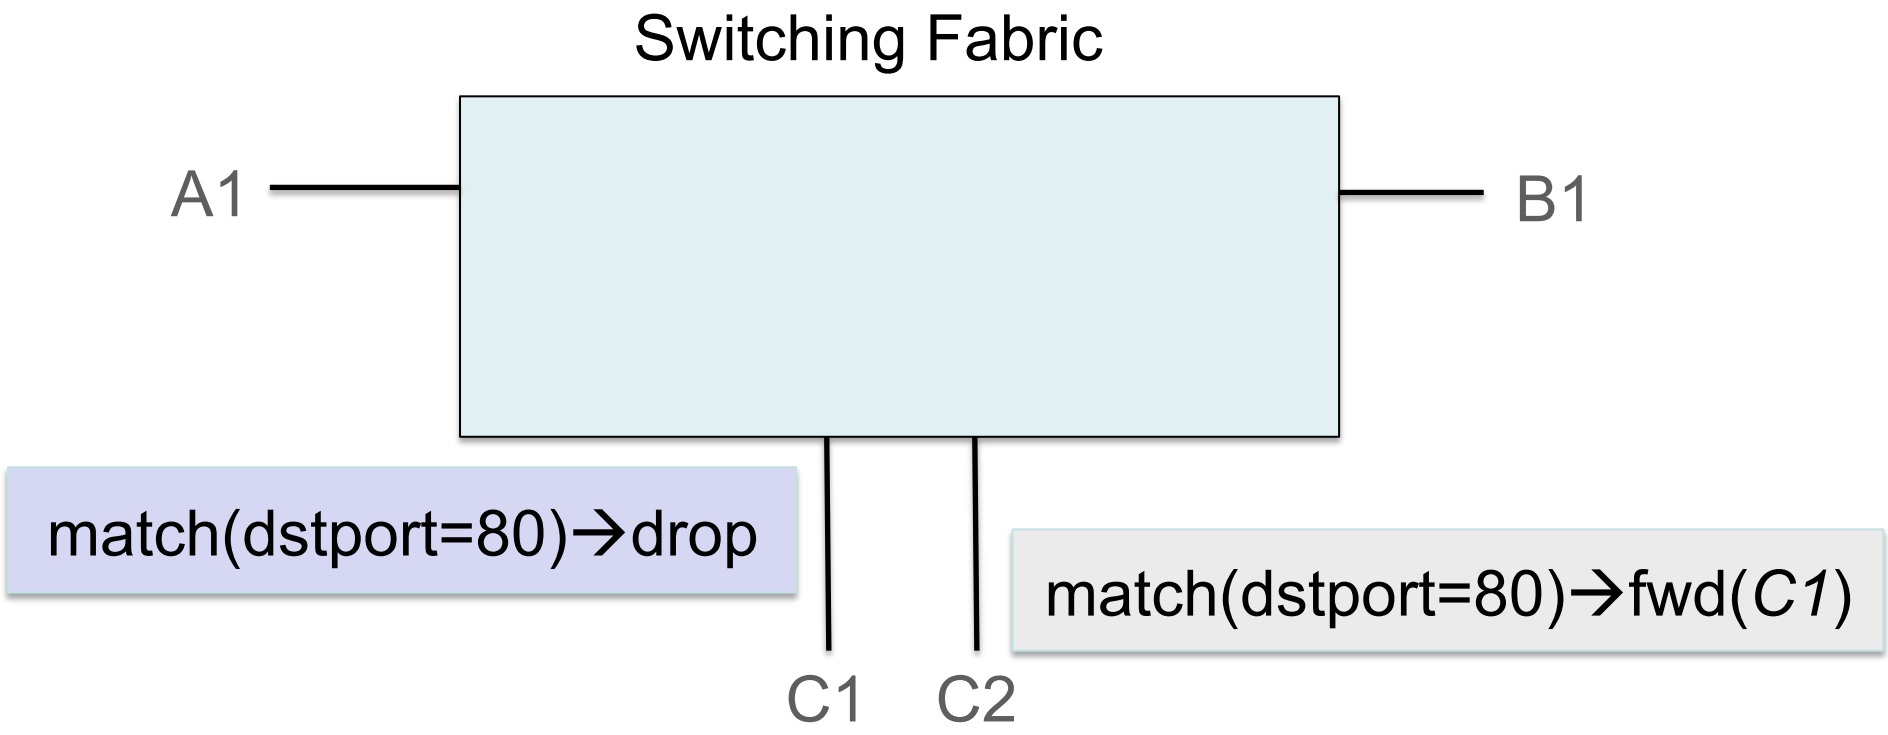

AS A & C directly program the SDX Switch

Conflicting Policies

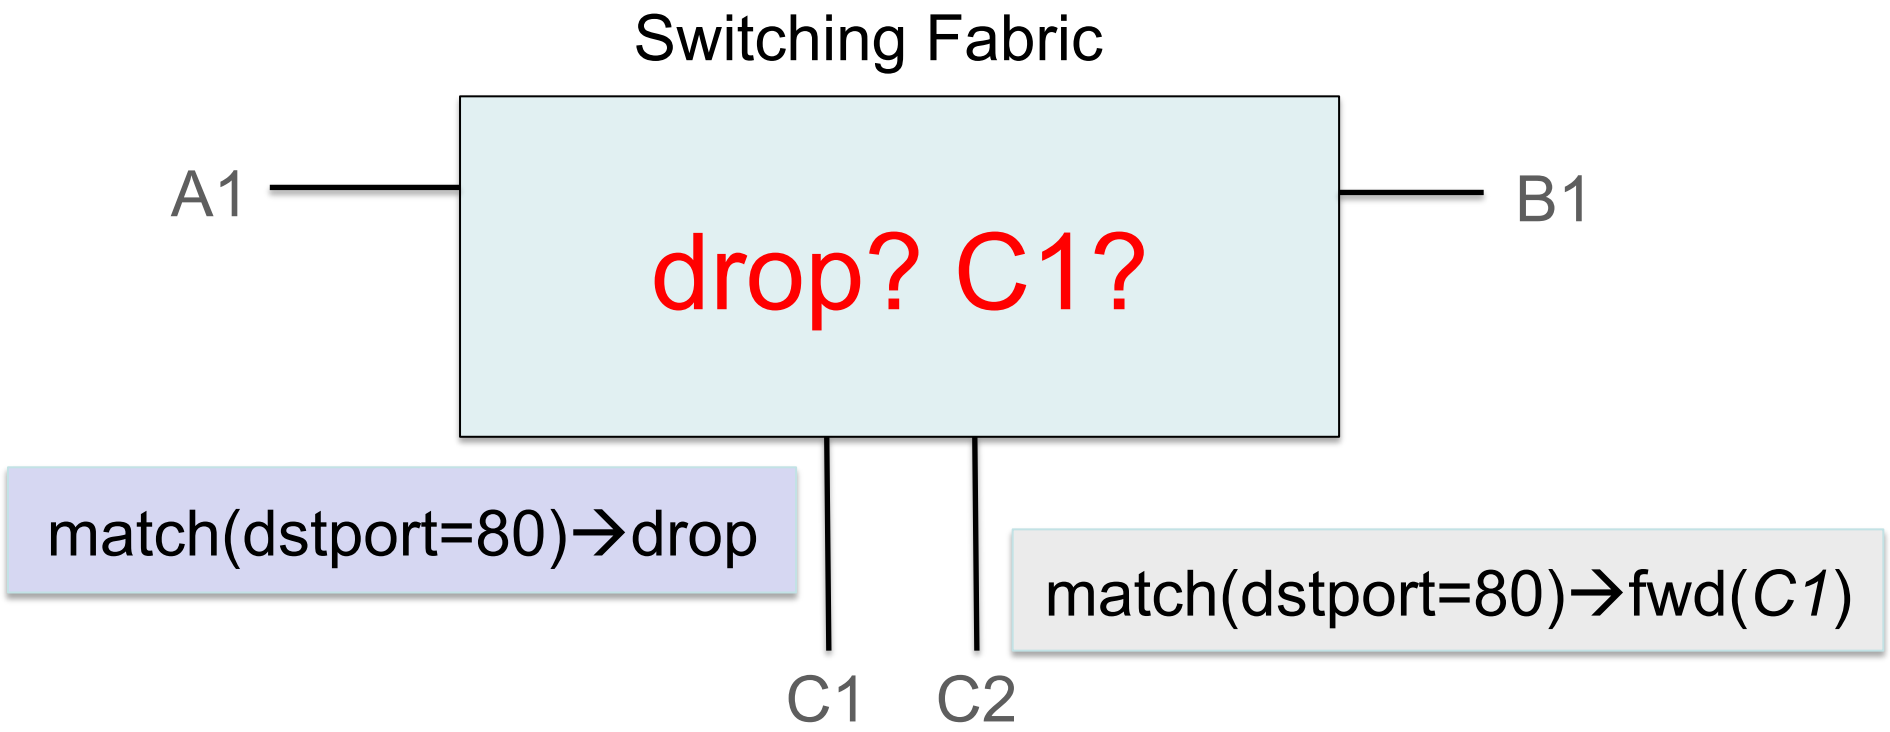

How to restrict participant's policy to traffic it sends or receives?

Virtual Switch Abstraction

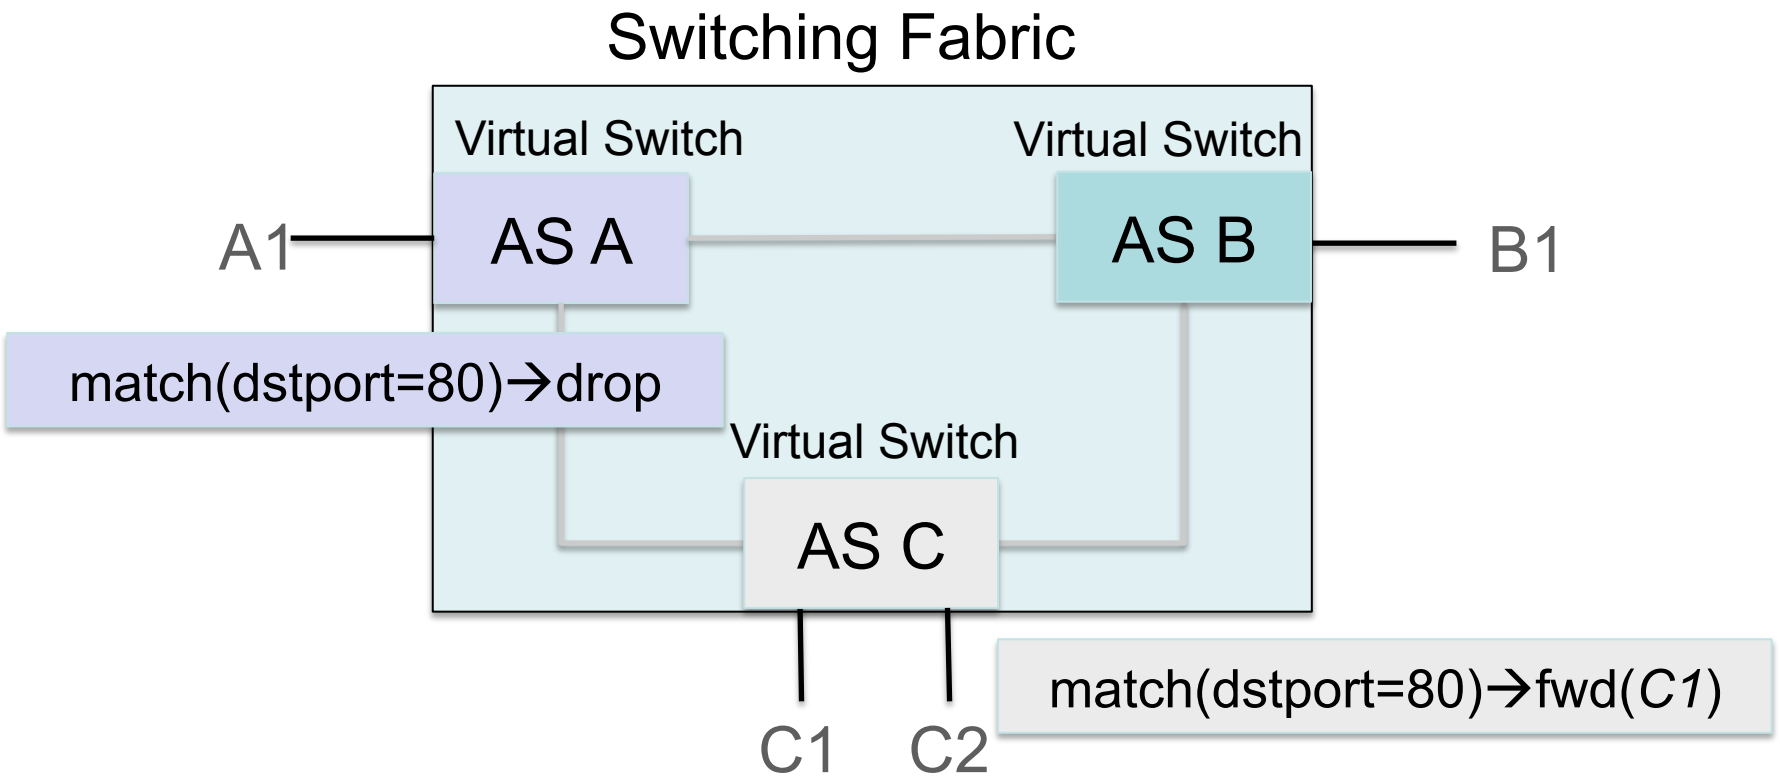

Each AS writes policies for its own virtual switch

Combining Participant's Policies

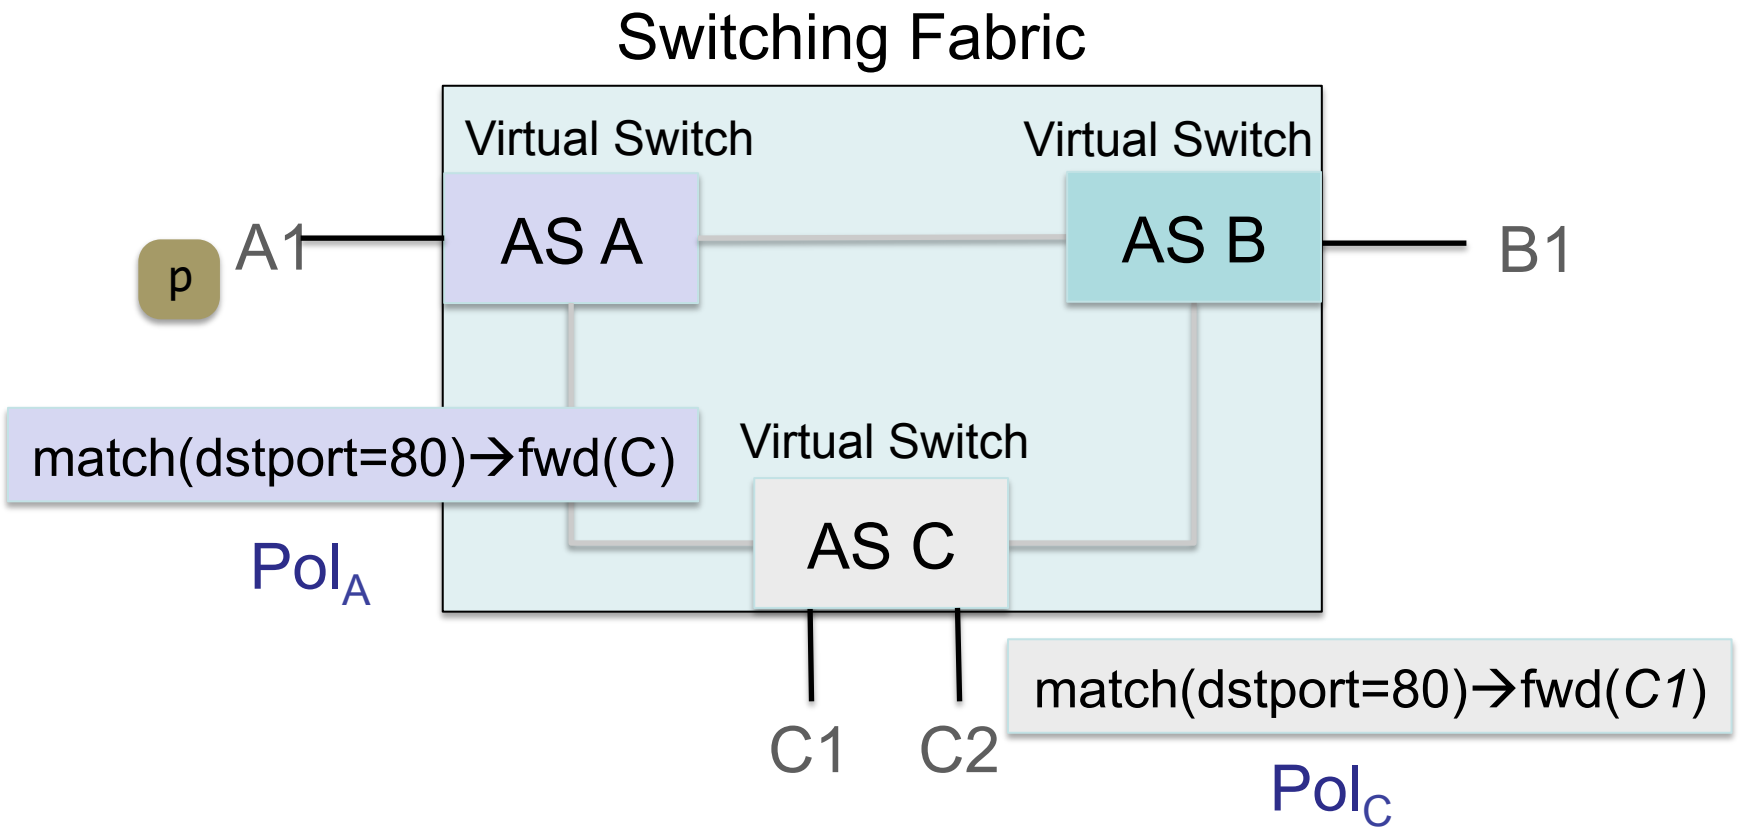

$$Policy(p) = Pol_A \rightarrow Pol_C$$

Building SDX is Challenging

- Programming **abstractions**
  - How networks define SDX policies and how are they combined together?
- **Interoperation** with BGP
  - How to provide flexibility w/o breaking global routing?
- **Scalability**
  - How to handle policies for hundreds of peers, half million prefixes and matches on multiple header fields?

Requirement: Forwarding Only Along BGP Advertised Routes

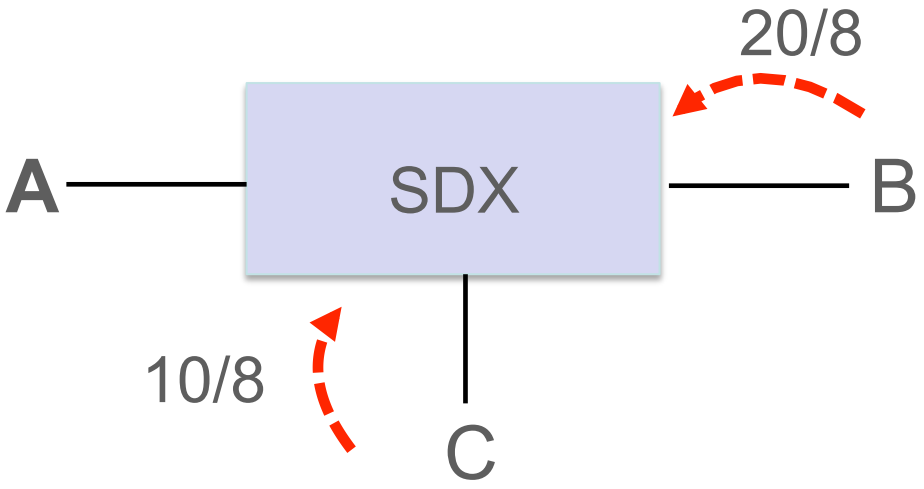

match(dstport=80) → fwd(*C*)

Ensure ‘p’ is not forwarded to C

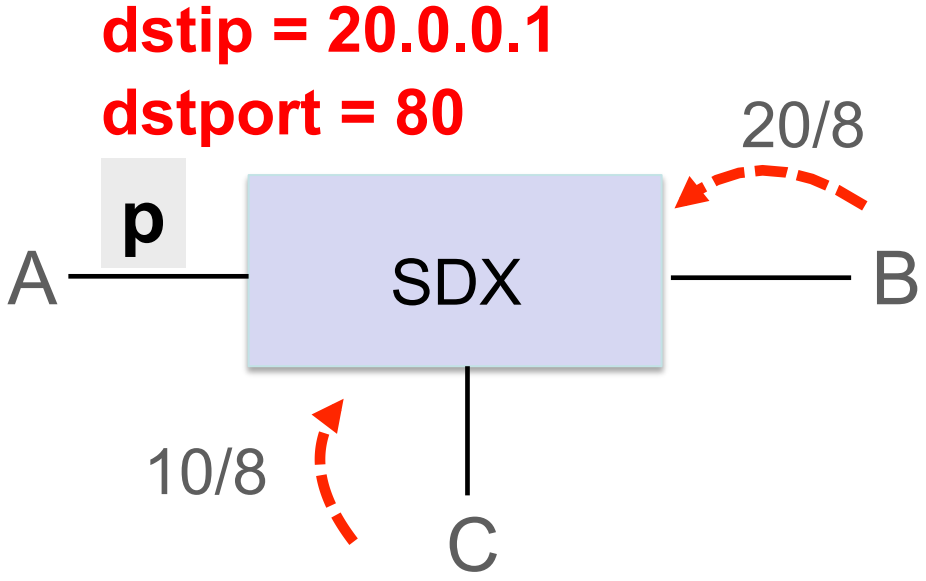

match(dstport=80) → fwd(*C*)

Solution: Policy Augmentation

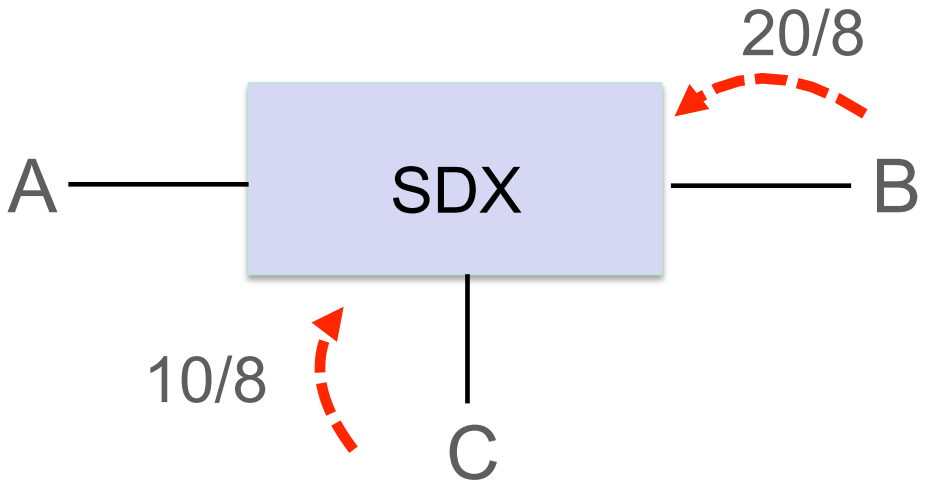

(match(dstport=80) && **match(dstip = 10/8)**)→ fwd(*C*)

Building SDX is Challenging

- Programming **abstractions**
  - How networks define SDX policies and how are they combined together?
- **Interoperation** with BGP
  - How to provide flexibility w/o breaking global routing?
- **Scalability**
  - How to handle policies for hundreds of peers, half million prefixes and matches on multiple header fields?

Scalability Challenges

- **Reducing Data-Plane State:** Support for all forwarding rules in (limited) switch memory

- **Reducing Control-Plane Computation:** Faster policy compilation

Scalability Challenges

- **Reducing Data-Plane State:** Support for all forwarding rules in (limited) switch memory
  millions of flow rules possible

- **Reducing Control-Plane Computation:** Faster policy compilation
  policy compilation could take hours

Reducing Data-Plane State: Observations

- Internet routing policies defined for **groups of prefixes**.*
- **Edge routers** can handle matches on hundreds of thousands of IP prefixes.

*Feamster et al., *Guidelines for Interdomain TE, CCR 2003*

Reducing Data-Plane State: Solution

Group prefixes with similar forwarding behavior

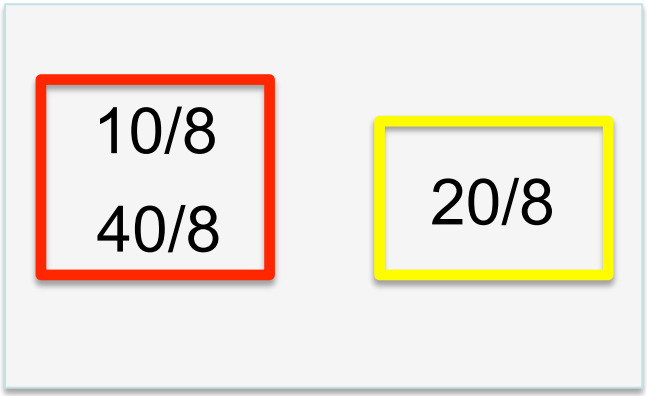

SDX Controller

Reducing Data-Plane State: Solution

Advertise one BGP next hop for each such prefix group

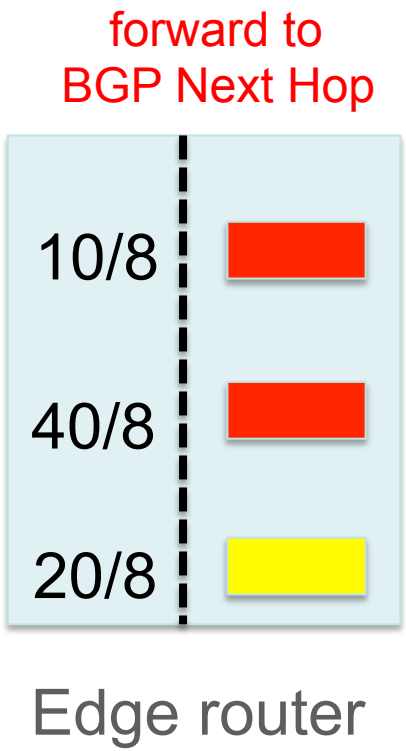

Reducing Data-Plane State: Solution

Flow rules at SDX match on BGP next hops

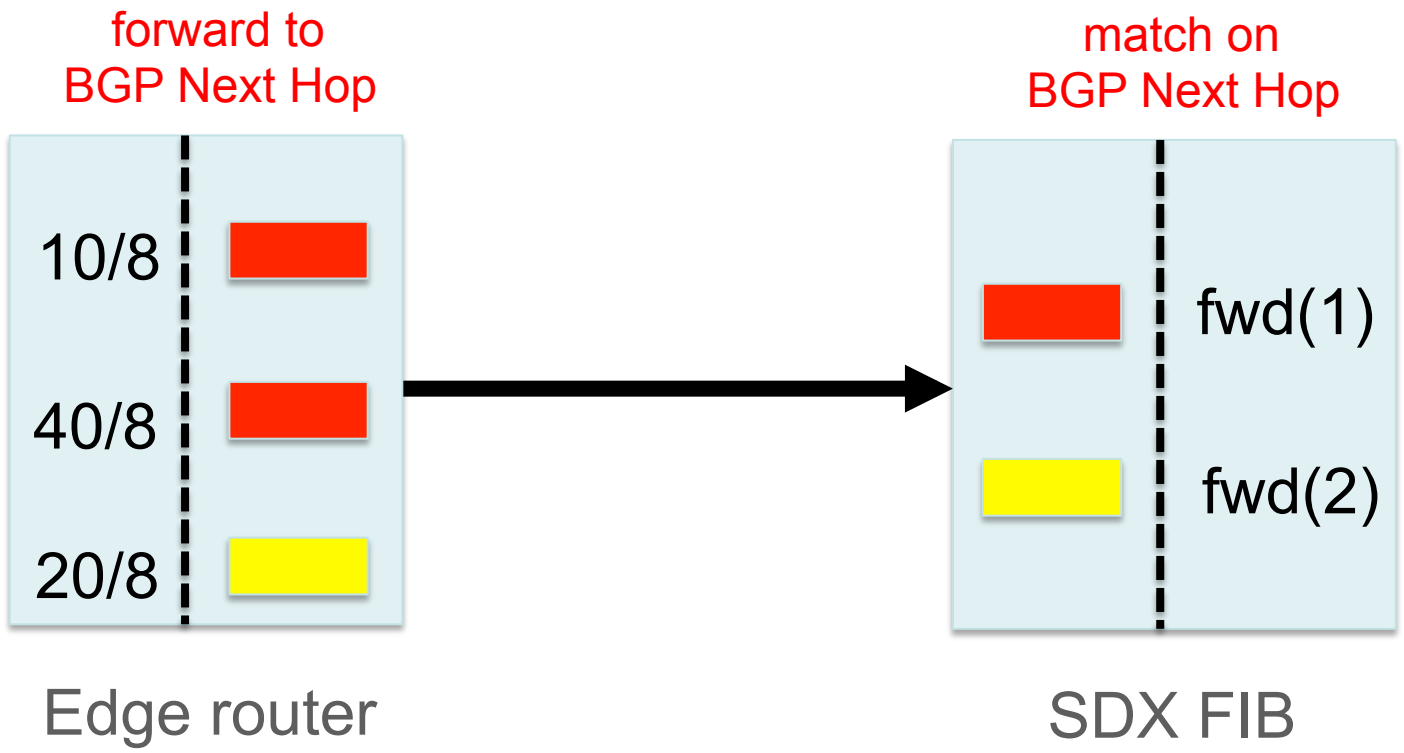

Reducing Data-Plane State: Solution

For hundreds of participants' policies,

few *millions* → < 35K

flow rules

Reducing Control-Plane Computation

- **Initial policy compilation time**
  - Leveraged domain-specific knowledge of policies
  - Hundreds of participants requires **< 15 minutes**
- **Policy recompilation time**
  - Leveraged bursty nature of BGP updates
  - Most recompilation after a BGP update **< 100 ms**

SDX Testbed

- Mininet-based Testbeds
  - Uses Transit Portal
  - Emulates edge routers

- Check out our demo
  - Application specific peering
  - Inbound traffic engineering

- Github repo: https://github.com/sdn-ixp/sdx/

Summary

- **SDN-based exchange (SDX)** is promising for fixing Internet routing
- Solved various challenges in building a real deployable SDX
- Many open research problems, both for building and using SDX