



# Coupled Congestion Control for RTP Media


Safiqul Islam, Michael Welzl, Stein Gjessing and Naeem Khademi

Department of Informatics


University of Oslo

R / T E

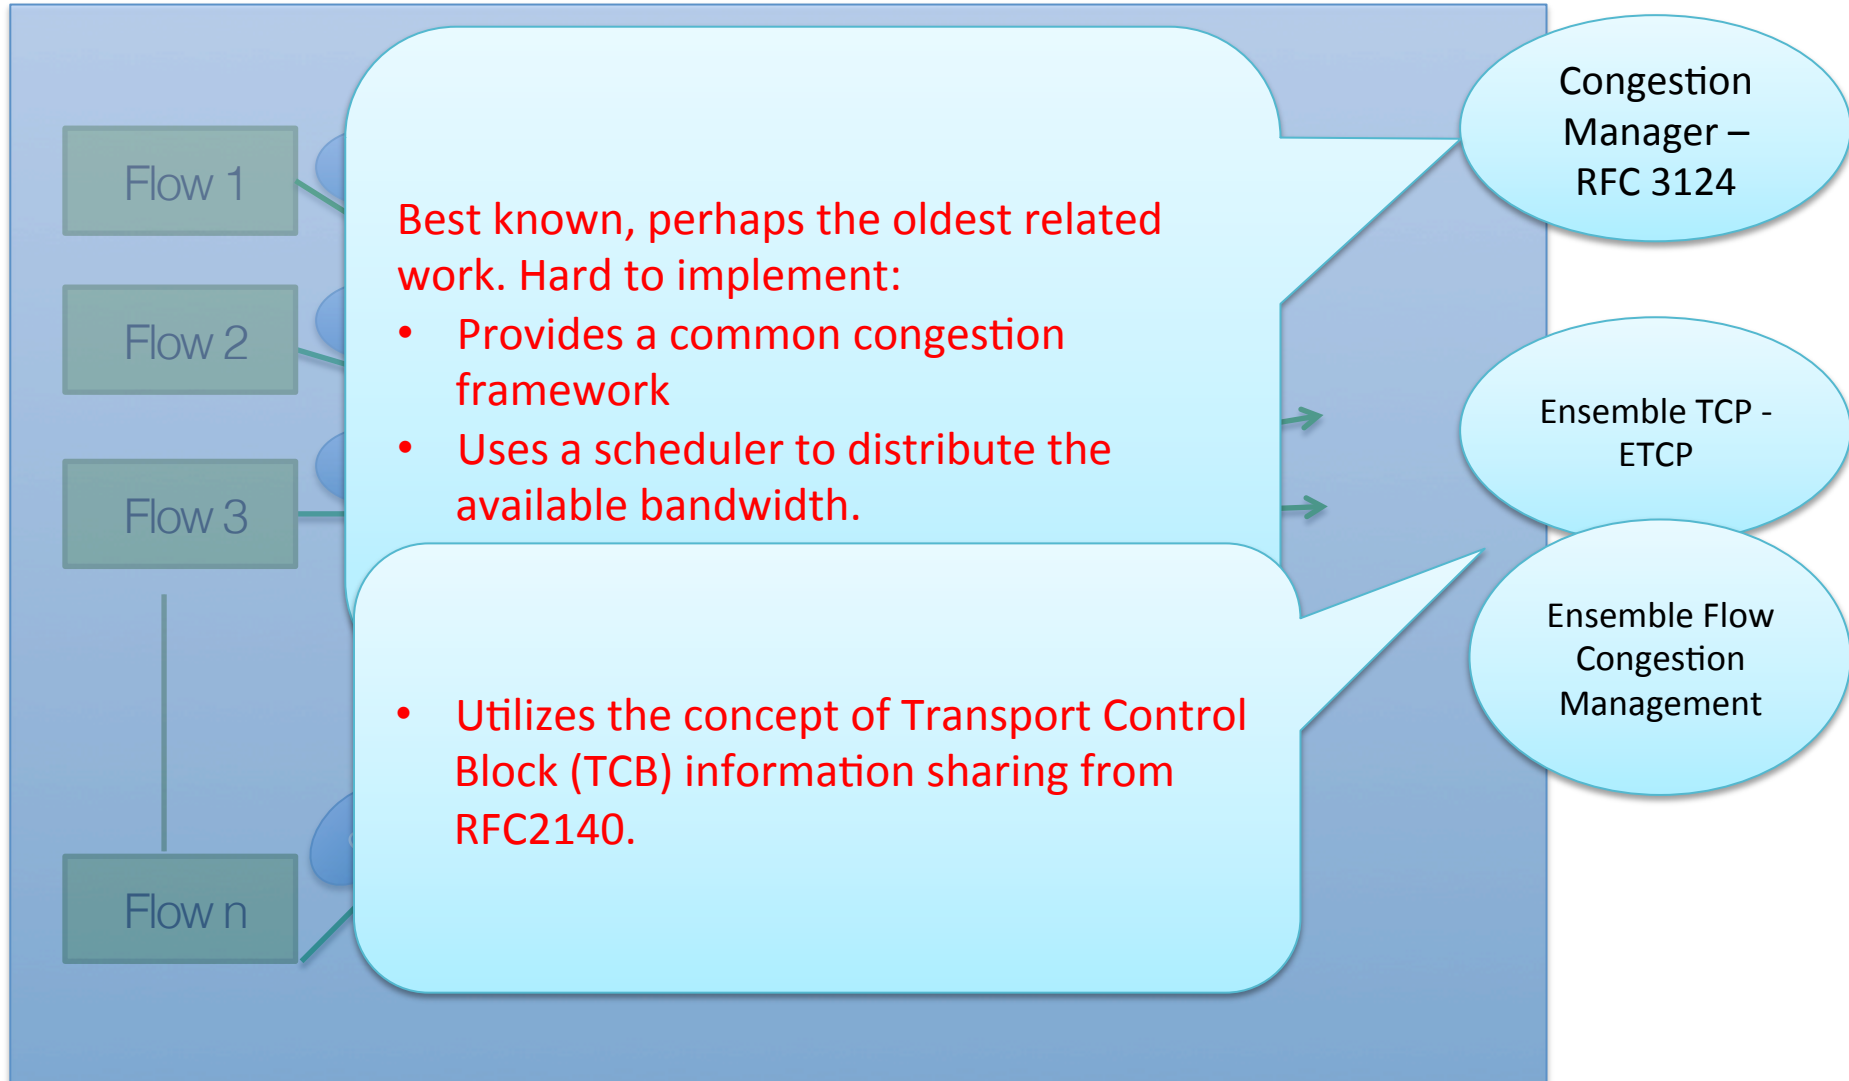
# Problem statement



# Problem statement – cont.




# Coupled CC – related work



# Shared bottlenecks

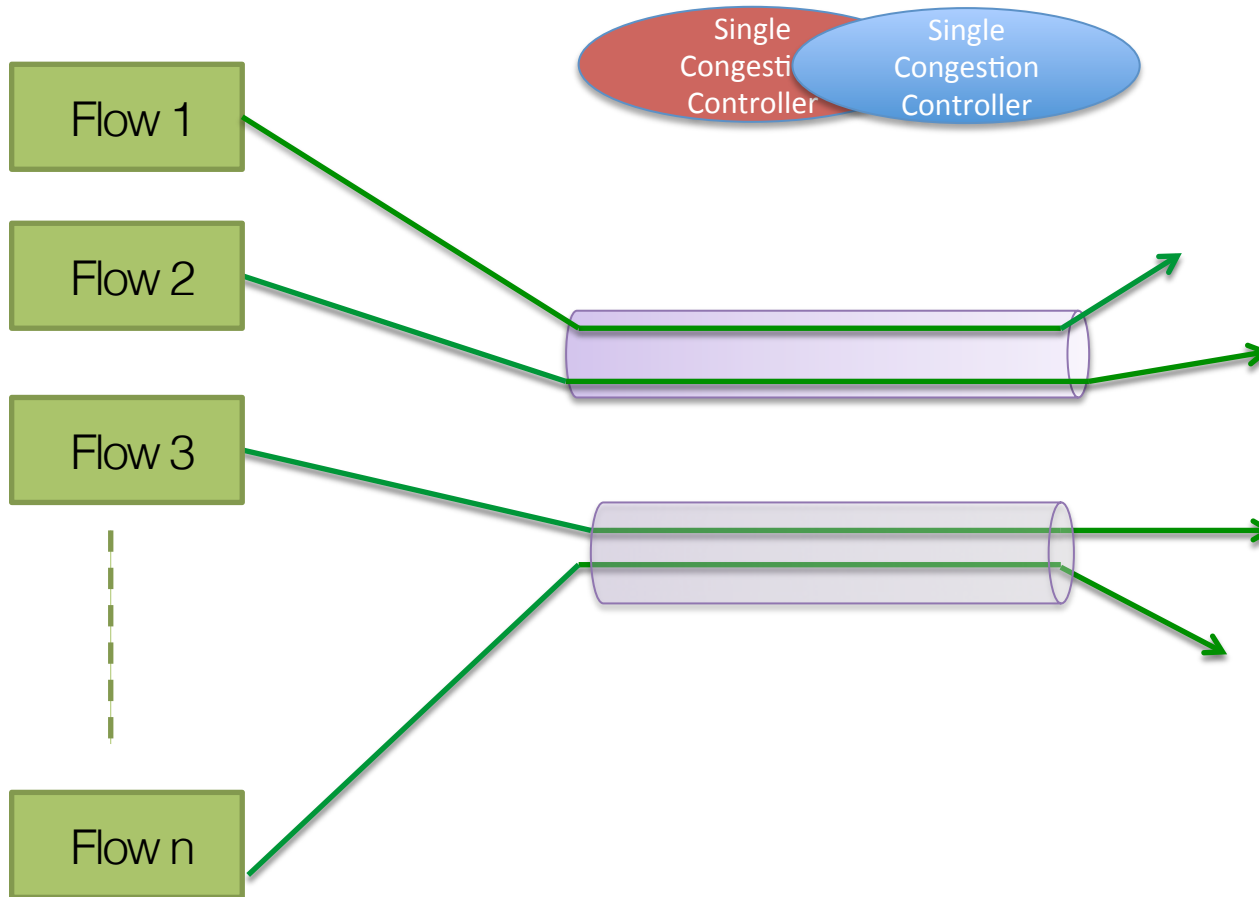
- Coupled CC only makes sense across a common bottleneck
  - This was ignored in prior work
  - But how to know?
- 1. Multiplexing (same 5-, actually 6-tuple)
  - Fits rtcweb (coupled-cc proposed in rmcats) – but only for same source/destination hosts
- 2. Configuration (e.g. common wireless uplink)
- 3. Measurement
  - Never 100% reliable, but: different receivers possible!
  - Historically considered impractical, but recent work:  
*David Hayes, Simone Ferlin-Oliveira, Michael Welzl: "Practical Passive Shared Bottleneck Detection Using Shape Summary Statistics", accepted for publication, IEEE LCN 2014, 8-11 September 2014*

# Coupled CC



This might work well for flows having same tuple but what happens when bottleneck changes?

# Coupled CC




- Add another CC?
- Add another scheduler?
- What about previous CC. state?

What we think:

- Better to have a simple algorithm that loosely couples existing cc with minimal changes.

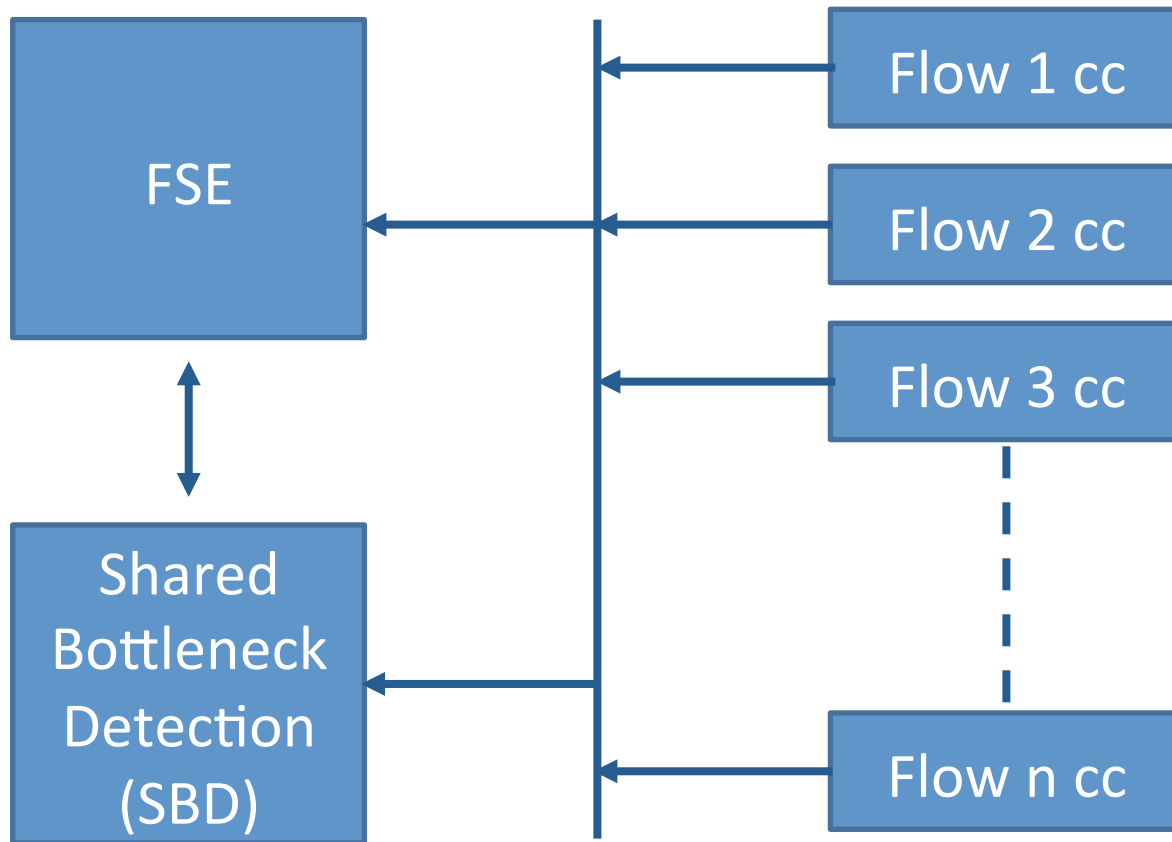
# Coupled CC



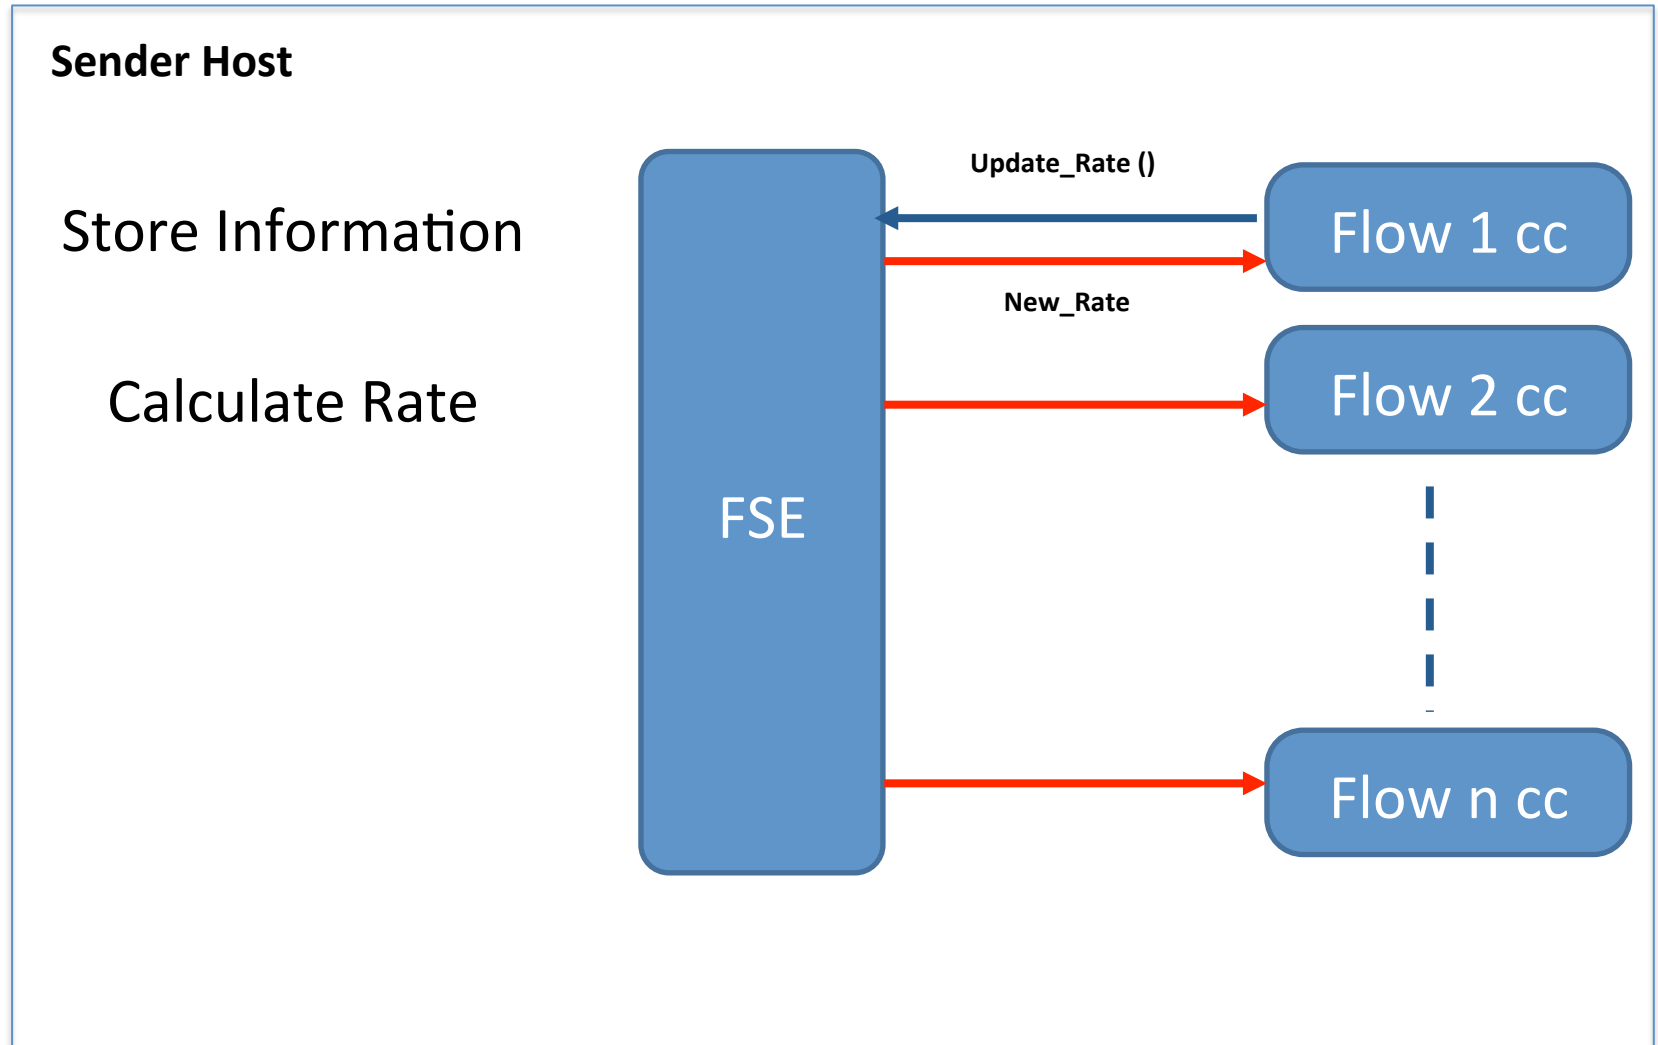


# The Flow State Exchange (FSE)

Sender Host



# The Flow State Exchange (FSE)

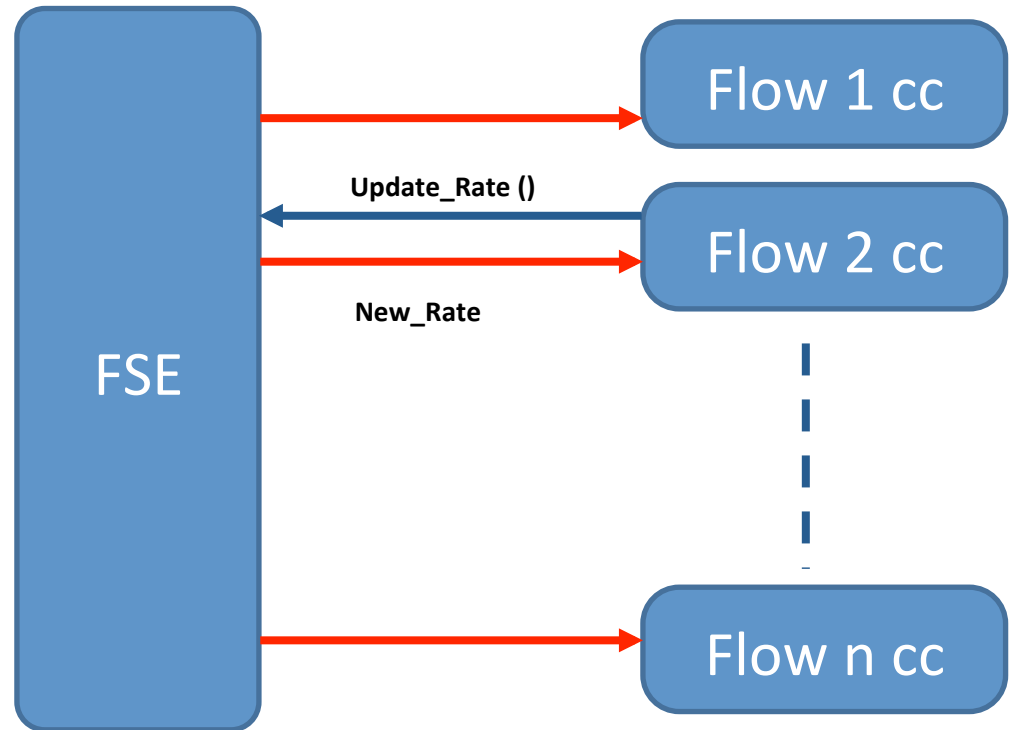


# The Flow State Exchange (FSE)

Sender Host

Store Information

Calculate Rate



# Simple algorithm

- Every time the congestion controller of a flow determines a new sending rate, the flow calls UPDATE

- FSE updates the sum of all rates, calculates the sending rates for all the flows and distributes them

.....▶

for all flows  $i$  in FG do

$FSE\_R(i) = (P(i) * \sum CR) / \sum P$

send  $FSE\_R(i)$  to the flow  $i$

end for


- Results were not good
  - Details are in the paper

# Updated algorithm

Idea: reduce the rate on congestion as one flow.

- No congestion: increase the aggregate by  $I/N$  where  $I$  is the increase factor.
- Congestion: Proportionally reduce the rate to emulate the congestion response of one flow.
  - Avoid over-reacting: set a time ( $2RTTs$ ) to react only once in the same loss event.


# Some simulation results



- Implemented in ns-2
- Two rate based protocols:
  - Rate Adaptation Protocol (RAP)
  - TCP Friendly Transport Protocol (TFRC)
- Bottleneck – 10 Mbps, Queue-length – 62 Packets (1/2 BDP), Packet Size – 1000 Bytes, RTT – 100 ms
- All tests (except when x-axis = time) ran for 300 seconds, carried out 10 times with random start times picked from first second; stddev consistently very small (  $\leq 0.2\%$  )

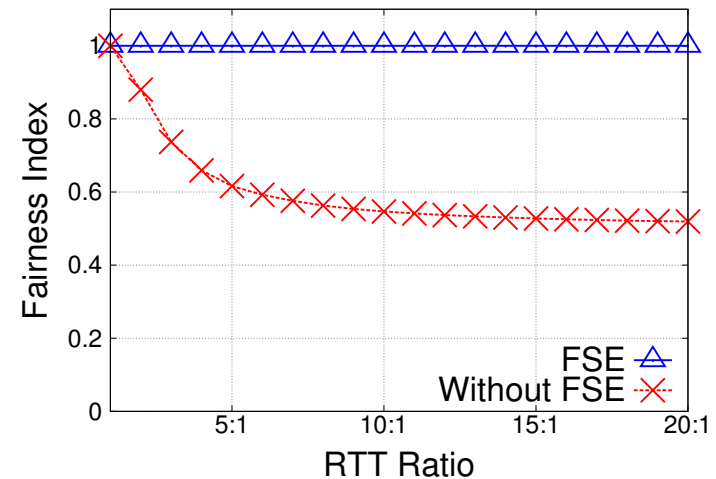
# Evaluation – prioritization and fairness

## Prioritization

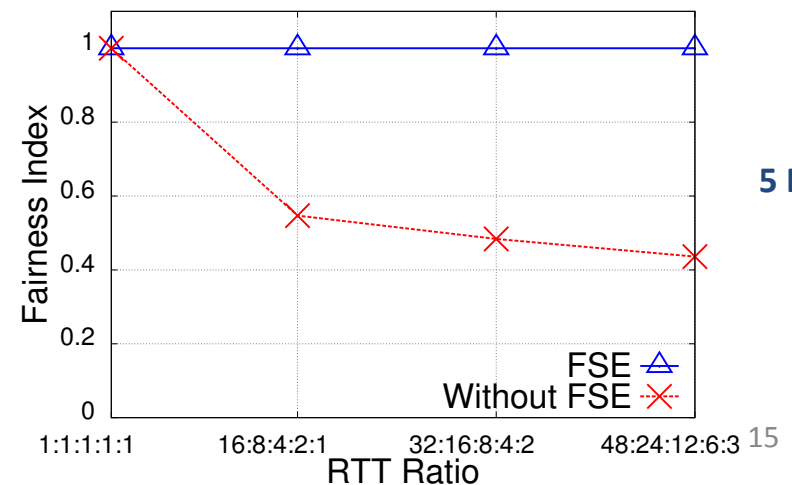


Priority of flow 1 increased over time

## Fairness




2 Flows




5 Flows

# Evaluation – FSE controlled flows competing with synthetic traffic

- TMIX synthetic traffic, taken from 60 minute trace of campus traffic at the University of Carolina [TCP Evaluation suite]
  - We used the pre-processed version of this traffic which is adapted to provide an approximate load of 50%











**TFRC**

Average Queue




Receiver makes assumptions about sending rate (expected length of loss interval) → loss event ratio  $p$  calculation wrong → sender too aggressive

Packet Loss Ratio




Link Utilization




# Different max Queue Lengths



**RAP**



**TFRC**



**10 Flows**




**15 Flows**

# How to evaluate app-limited flows?


- Not easy: who is in control?
- RMCAT codec model not available yet
- From a transport point of view, the send buffer can either run empty or not, with variations in how quickly changes between these two states occur
  - We used a non-reacting video trace of a person talking in a video conference with a well-known H264 encoder (X264) to steer the app sending rate
    - I-frame in the beginning, rest was mostly P-frames

# Evaluation – an application limited flow and one greedy flow (RAP)

FSE




Without FSE



FSE-controlled flows proportionally reduce the rate in case of congestion; without FSE, synchronization causes app-limited flow to over-react

# Using priorities to “protect” the app-limited from the greedy flow (RAP)



High-priority (1) application limited flow #1 is hardly affected by a low-priority (0.2) flow #2 as long as there is enough capacity for flow 1

# Summary

- Coupled congestion control via Flow State Exchange
  - Currently proposed for WebRTC in the RMCAT group
  - Satisfies the requirements of controllable fairness with prioritization
  - Reduces queue delay and packet loss without significantly affecting throughput
- Future work
  - Test our method in Chromium (Google's CC)
  - To incorporate WebRTC's data channel, we will investigate coupling with window-based protocol too

That's all !!

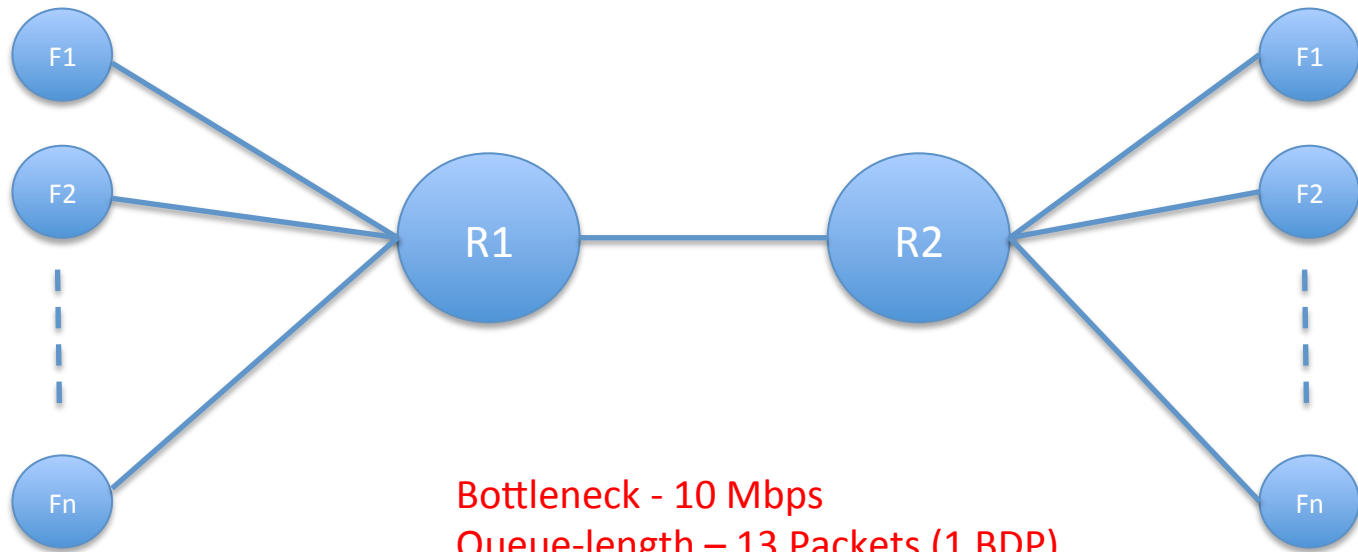
**Questions?**



# Backup slides




# Experimental setup (simple algorithm)




Bottleneck - 10 Mbps  
Queue-length – 13 Packets (1 BDP)  
Packet Size – 1000 Bytes  
Senders have always data to send

# Results – simple algorithm


RAP




TFRC



Average  
Queue  
Length




Packet  
Loss  
Ratio




Why?

# What's going on? (simple algorithm)



With FSE



Without FSE

- Queue drains more often without FSE
  - Should emulate the congestion response of one flow
    - FSE: 2 flows with rate  $X$  each; one flow halves its rate:  $2X \rightarrow 1\frac{1}{2}X$
    - When flows synchronize, both halve their rate on congestion, which halves the aggregate rate
    - We want that !  $2X \rightarrow 1X$