

A Decomposition-Based Architecture for Distributed Virtual Network Embedding

Flavio Esposito

Advanced Technologies



St.Louis, MO

Ibrahim Matta

Computer Science Department



Imagine a world without knobs



We have knobs to tune what's important to us



House's Behavior (Electrical)
oven, air conditioning, ...



Car's Behavior (Mechanical)
speed, music volume, ...



Organ's Behavior (Biological)
diabetes, dopamine, hormones, ...

Wouldn't it be great to have knobs for protocols too?



Algorithmic tradeoffs

e.g., converge speed vs performance optimality

Design tradeoffs

e.g., centralized or distributed

Implementation tradeoffs

e.g., support for delay or BW sensitive app

Ad-hoc protocol development prevents quick adaptation

Small Cloud-based **Startup**

Adaptation is the key to survive
code to customers needs quickly

Medium-size **Enterprise**

Slower growth: can't afford clean-slate
too much to change: this one works!
...till when?

Big Cloud **Provider**

Ossification (see Internet)
only few dominate the cloud market



Need for a Cloud easier and cheaper to manage

*“Management is responsible for 80% of IT budget”
“responsible for 62% of outages”*



Millions of dollars are spent in support of many protocols and apps with different requirements



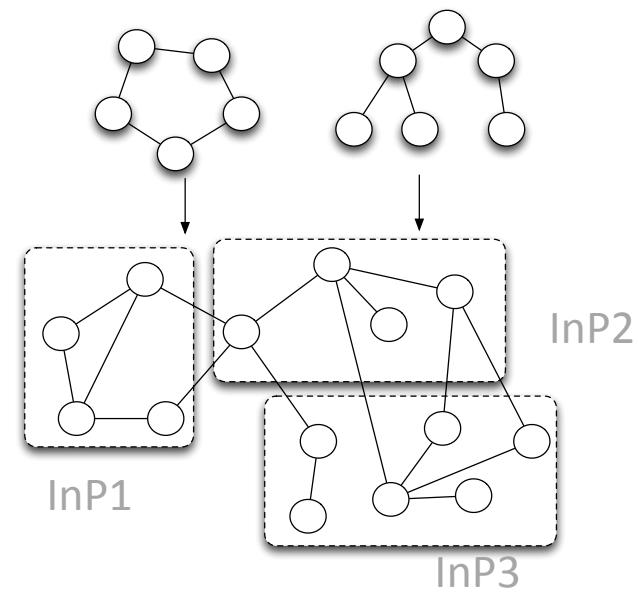
Virtual Network Embedding is a Constrained Graph Matching



Shared
Infrastructure

Virtual
networks

Possibly multiple:
Service Providers (SPs)
Infrastructure Providers (InPs)



Use our architecture to customize distributed VN embedding solutions

Apply decomposition techniques to VN Embedding solved as network utility maximization problem.

Obtain distributed algorithms with different goals that suit needs of specific applications



Can you model your problem as Network Utility Maximization?

Different instantiation of the optimization problem

Different policies

Different implications

Alternative VN Embedding problem representations

Different decompositions

Embedding algorithms

Engineering tradeoffs



Our Contributions (Talk Outline)



To separate policies from
VN embedding mechanisms

Unifying VN Embedding Architecture
identifying the invariances of the problem

To adapt embedding
to providers' goals

Design new and subsume embedding solutions
instantiating different decomposition policies

To enable experiments with
new VN embedding policies

Tradeoff analysis over simulation & prototype
implemented over a local Linux-based testbed

Identification of Mechanisms



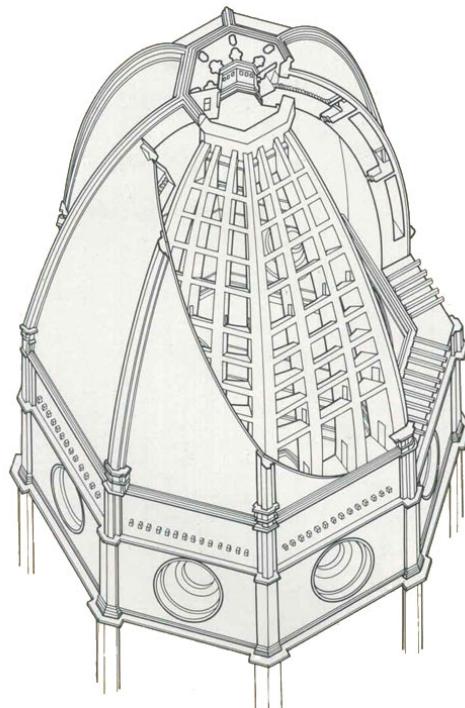
To separate policies from
VN embedding mechanisms

Unifying VN Embedding Architecture
Identifying the invariances of the problem

Design new and subsume embedding solutions
Instantiating different decomposition policies

Tradeoff analysis over simulation & prototype
implemented over a local Linux-based testbed

An architecture identifies the invariances in a problem



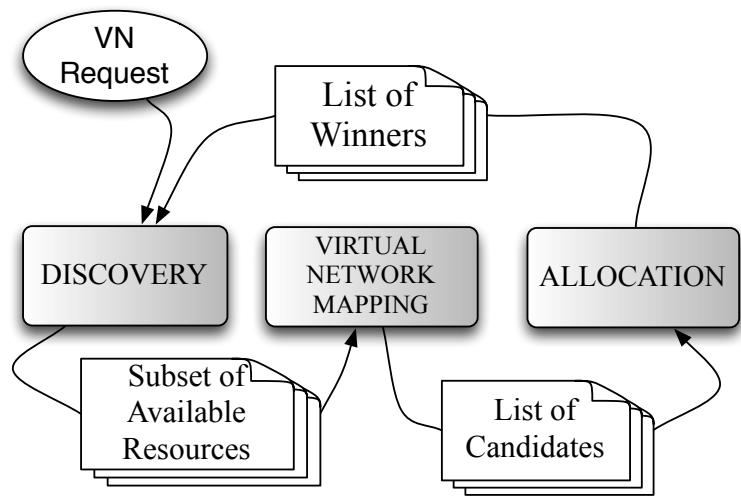
What is an architecture?

1. Identify **mechanisms**
(invariances)
2. Identify **who does what**
(separation of functionalities)

What is a policy? (variant aspect of a mechanism)

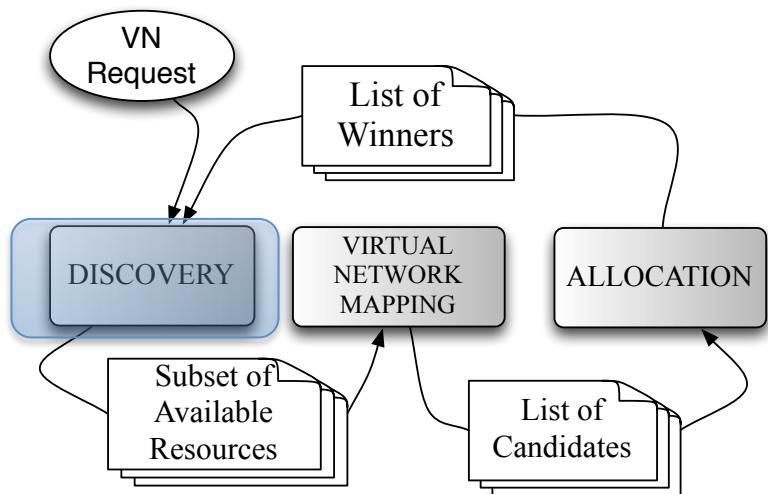
Provider policies
to host virtual nodes & links

Three common mechanisms interfaced by binding constraints

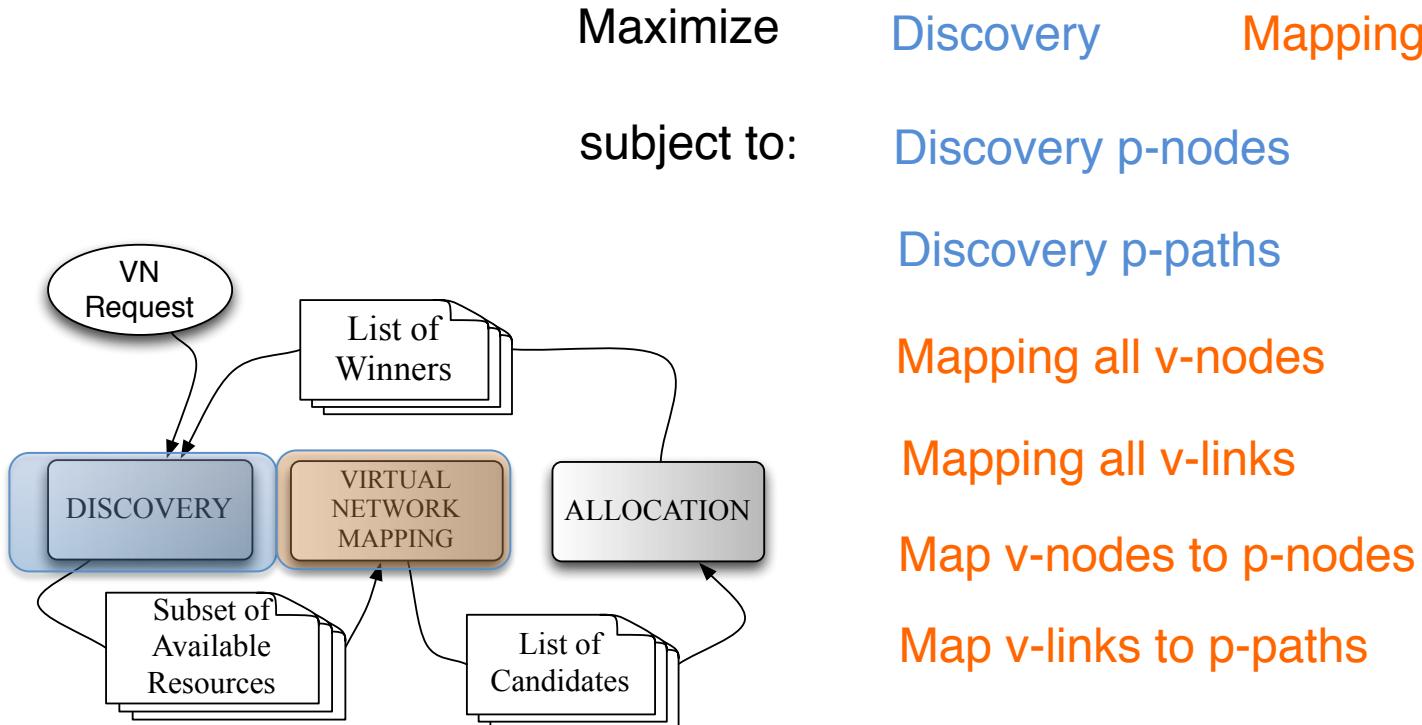


Three common mechanisms interfaced by binding constraints

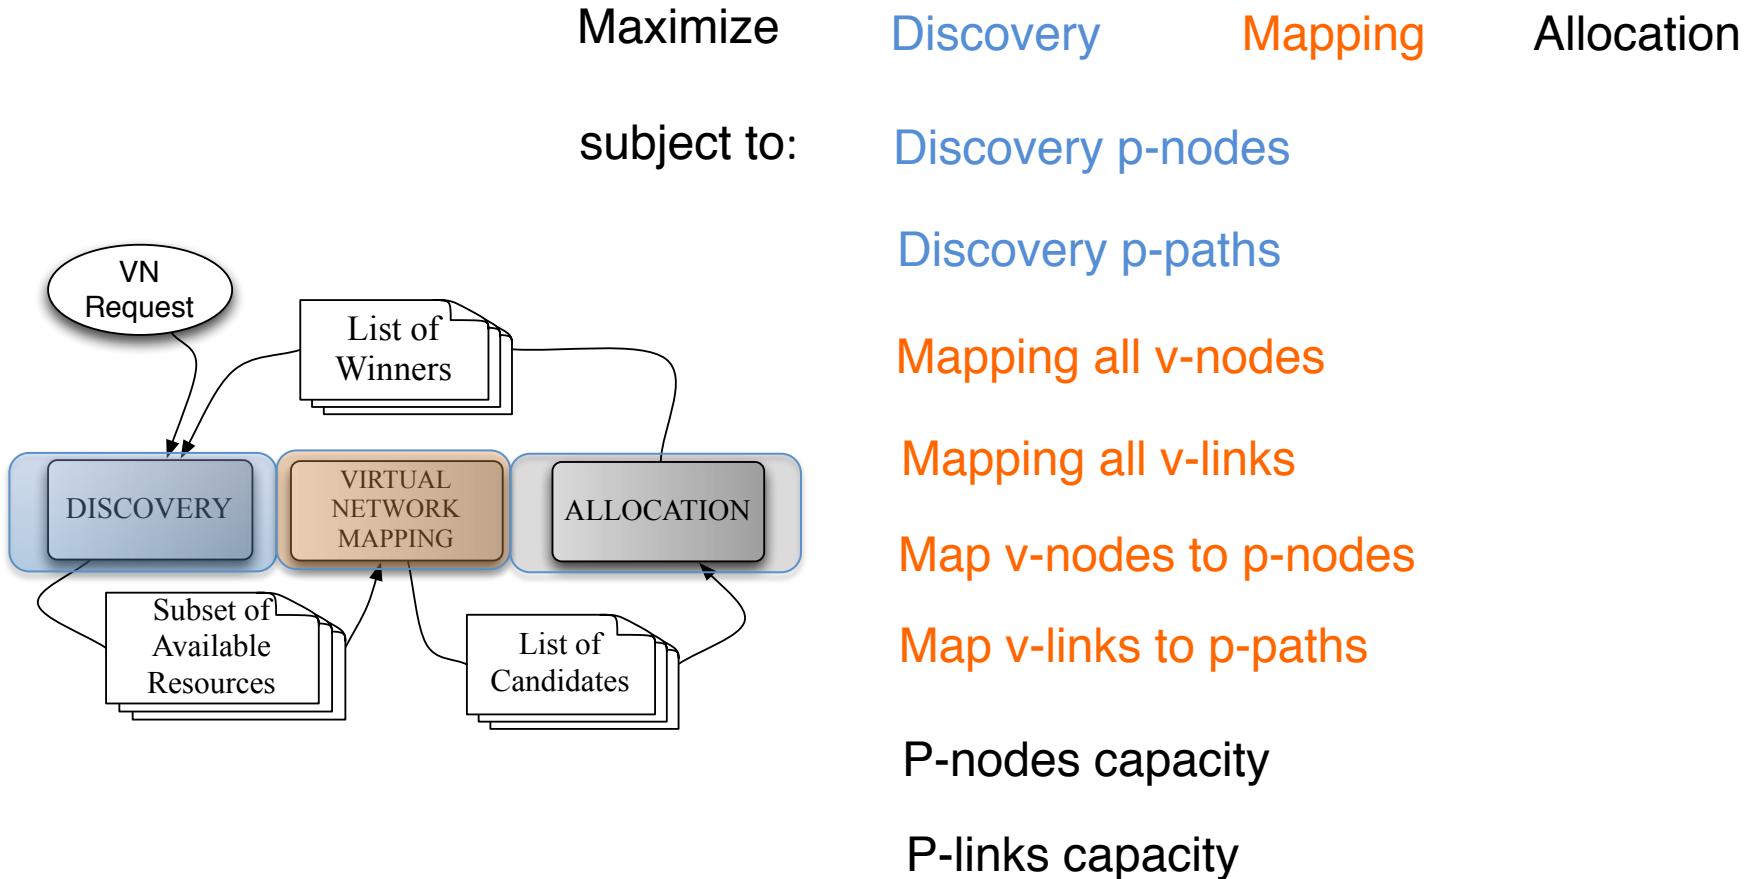
Maximize Discovery
subject to: Discovery p-nodes
 Discovery p-paths



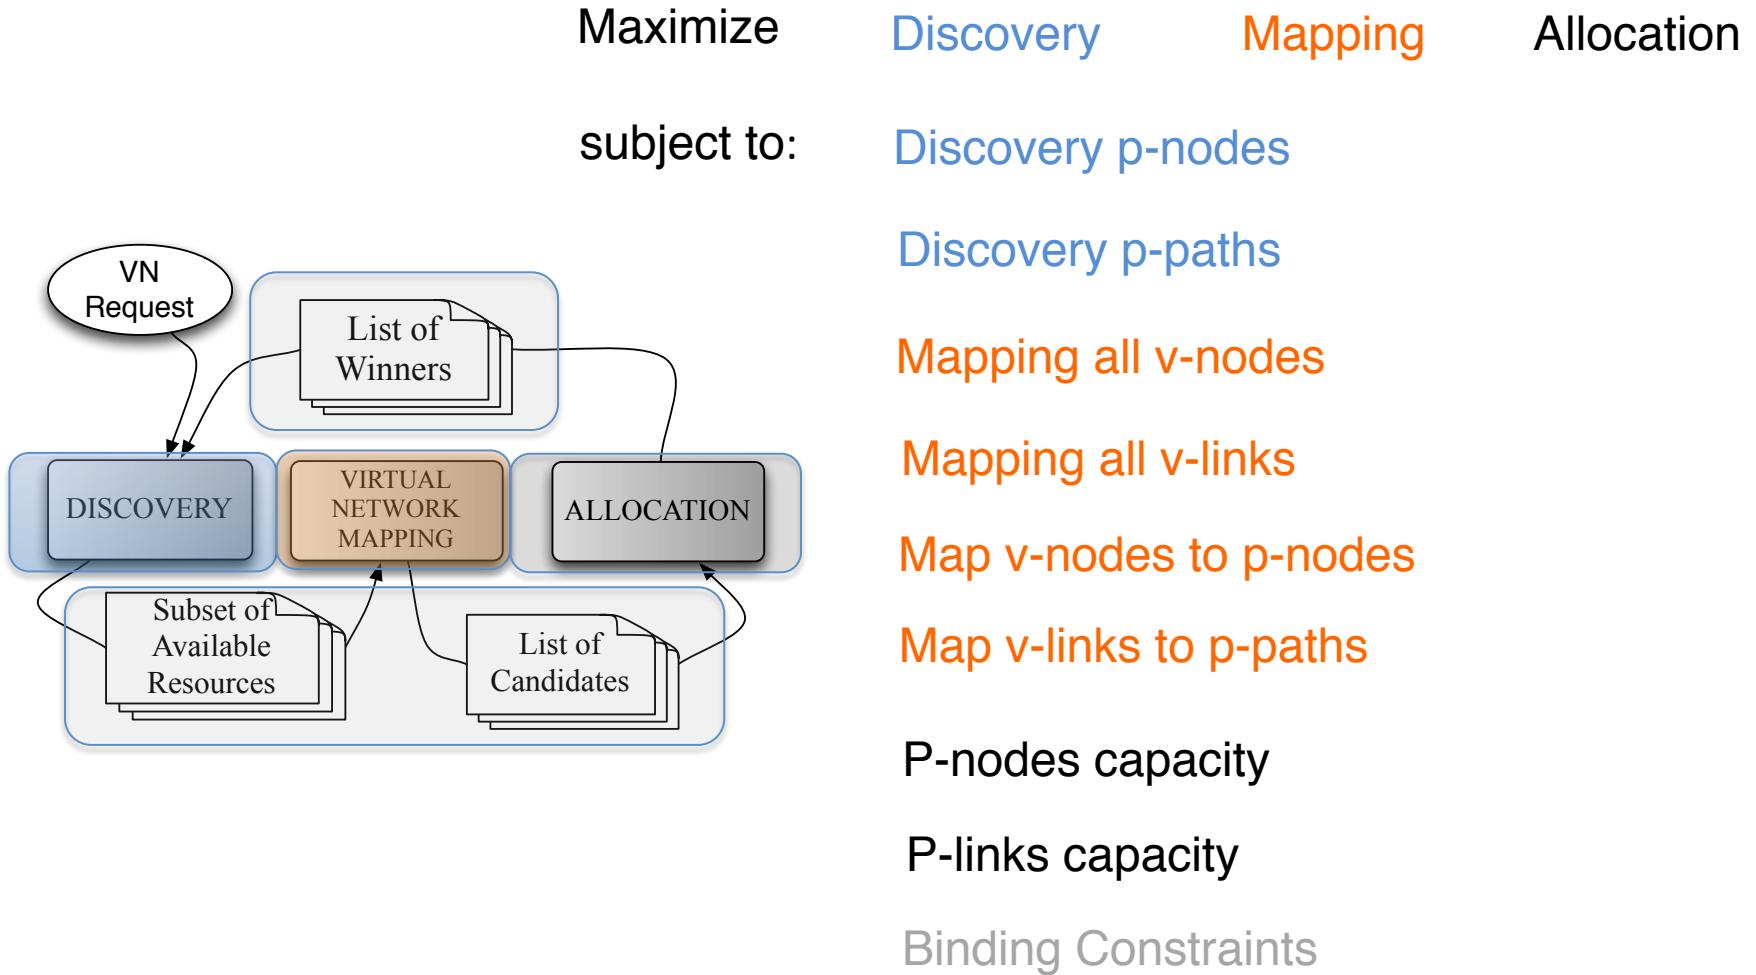
Three common mechanisms interfaced by binding constraints



Three common mechanisms interfaced by binding constraints



Three common mechanisms interfaced by binding constraints



Different decompositions for different distributed (VN embedding) solutions



To separate policies from VN embedding mechanisms

Unifying VN Embedding Architecture identifying the invariances of the problem

To adapt embedding to providers' goals

Design new and subsume embedding solutions
instantiating different decomposition policies

Enable experiments with new embedding policies

Tradeoff analysis over simulation & prototype implemented over a local Linux-based testbed

The structure of the problem is rich
so many decompositions are possible

Primal Decomposition

Dual Decomposition

The structure of the problem is rich
so many decompositions are possible

Solve fixing **Variables:** Primal Decomposition

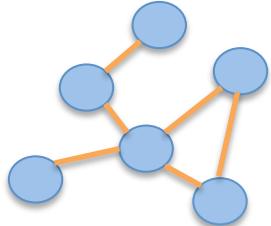
Solve relaxing **Constraints:** Dual Decomposition

The structure of the problem is rich
so many decompositions are possible

Solve fixing **Variables**: **Primal Decomposition**
Master: assigns ph.resources to subproblems
Subproblems: embed given available resources

Solve relaxing **Constraints**: **Dual Decomposition**
Master: sets prices of all virtual resources
Subproblems: embed given resource prices

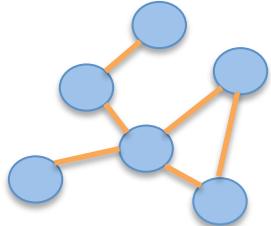
Primal: decompose fixing decision variables



Fix nodes & links
decision variables

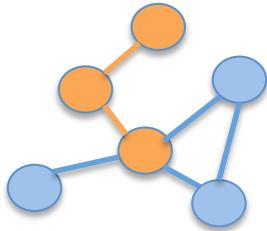
Optimize virtual nodes then links
using the optimal node embedding
see e.g. [9,20]

Primal: decompose fixing decision variables



Fix nodes & links
decision variables

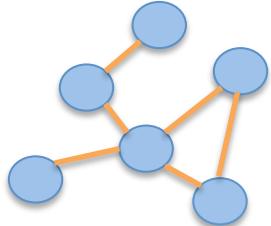
Optimize virtual nodes then links
using the optimal node embedding
see e.g. [9,20]



Fix sub-VN
decision variables

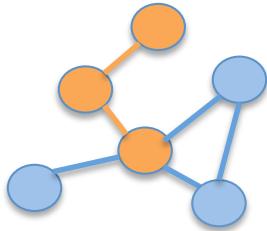
Optimize nodes & links of first VN partition
then nodes & links of the next VN partition
see e.g. [14,15]

Primal: decompose fixing decision variables



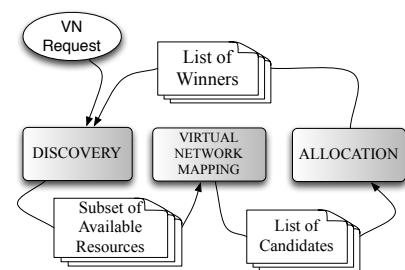
Fix nodes & links
decision variables

Optimize virtual nodes then links
using the optimal node embedding
see e.g. [9,20]



Fix sub-VN
decision variables

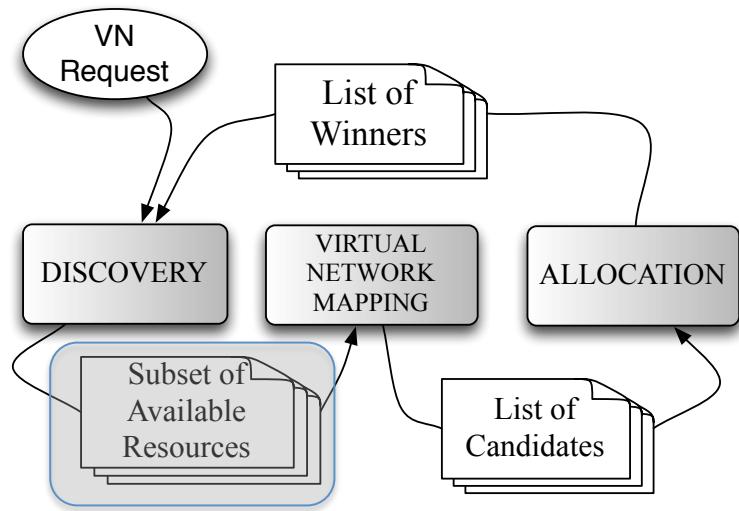
Optimize nodes & links of first VN partition
then nodes & links of the next VN partition
see e.g. [14,15]



Fix subproblems
decision variables

Discovery and/or Mapping and/or Allocation
Optimize only some embedding mechanisms
see e.g. [6,18]

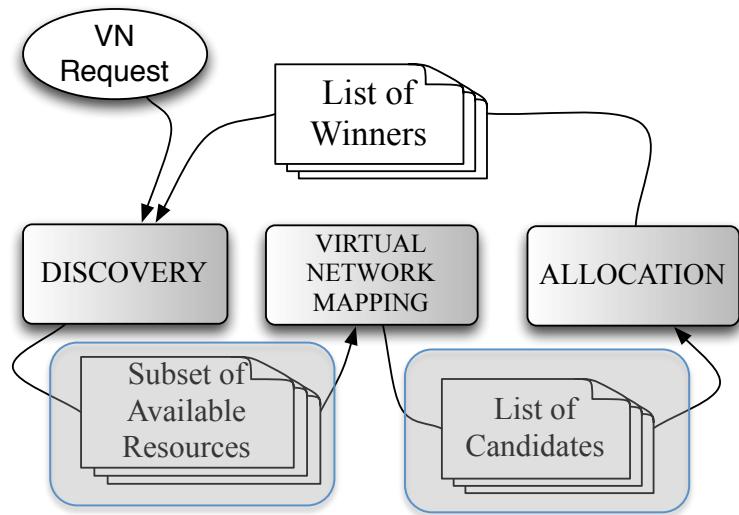
Dual: decompose relaxing constraints



Relax binding constraints between
Discovery and **VN Mapping** phase

see e.g. [2,12]

Dual: decompose relaxing constraints



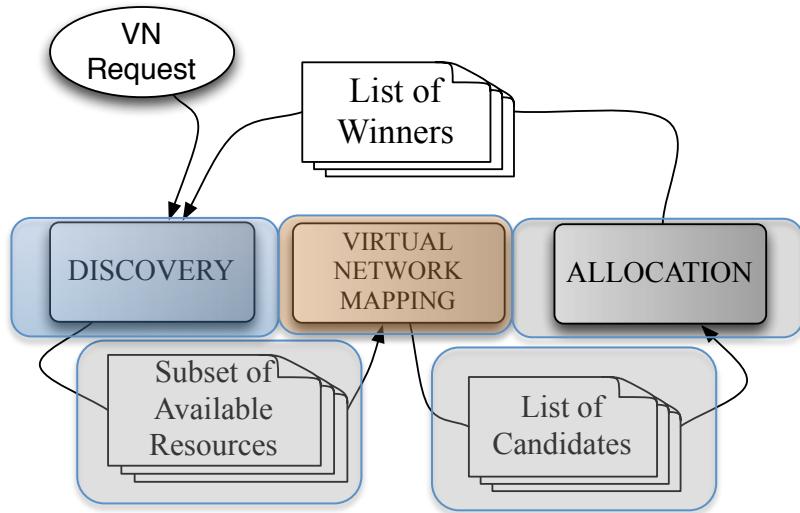
Relax binding constraints between
Discovery and **VN Mapping phase**

see e.g. [2,12]

Relax binding constraints between
VN Mapping & Allocation phase

see e.g. [14]

Dual: decompose relaxing constraints



Relax binding constraints between
Discovery and **VN Mapping phase**

see e.g. [2,12]

Relax binding constraints between
VN Mapping & Allocation phase

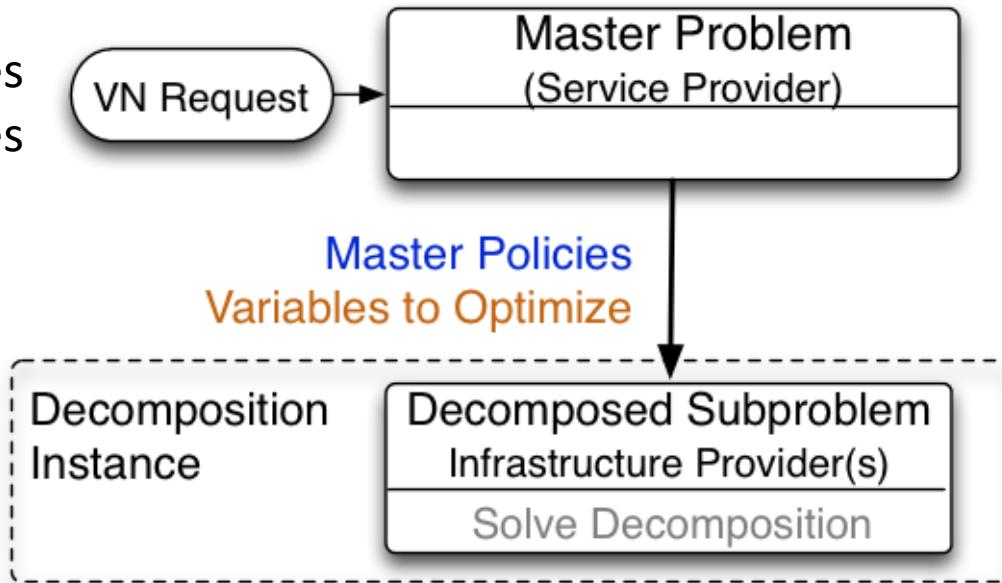
see e.g. [14]

Relax any non-binding constraints

[your papers here]

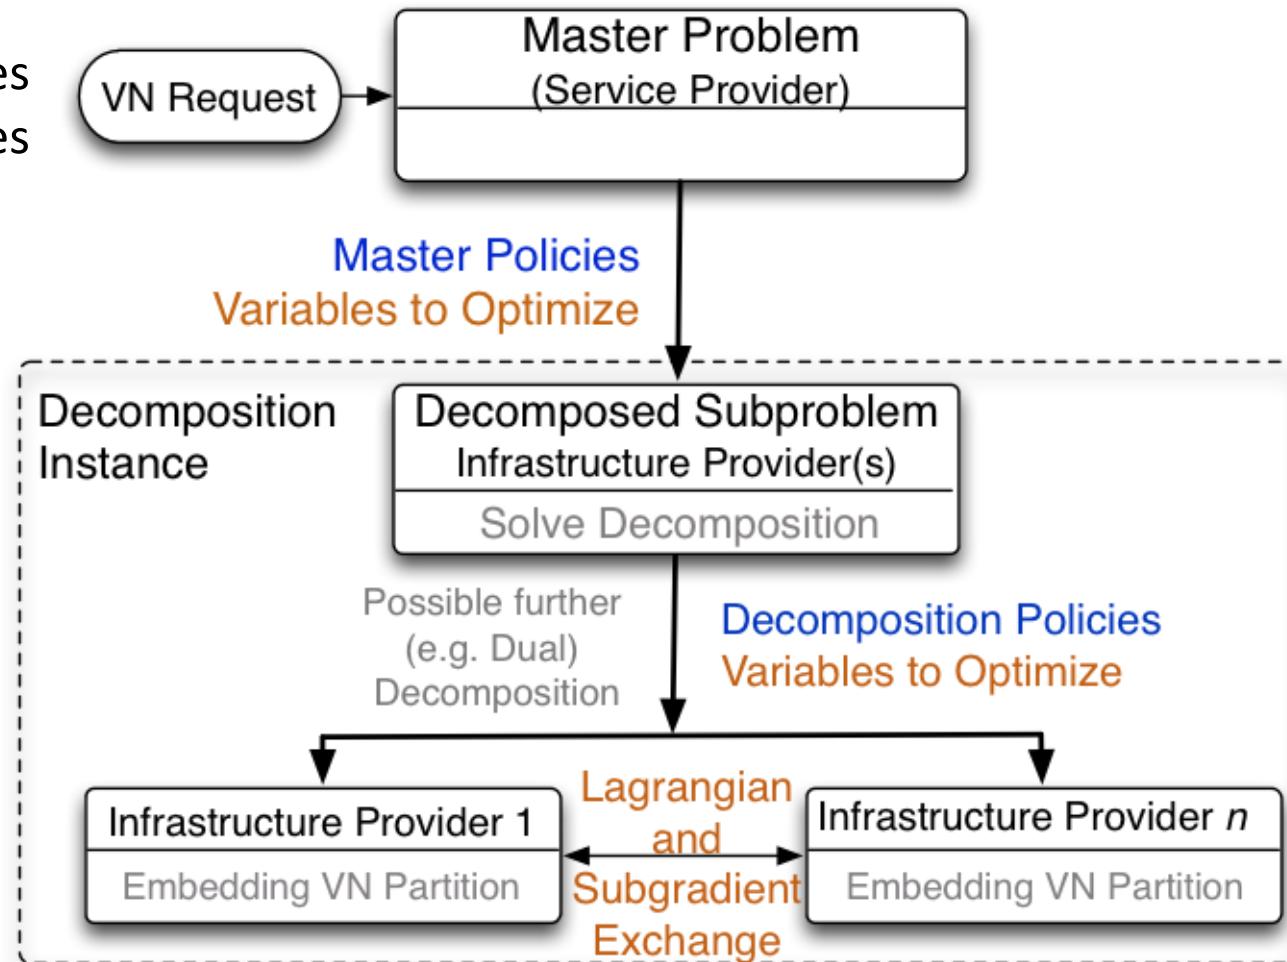
Use primal and dual to design your VN embedding solution

Service Provider instantiates according to its policies



Use primal and dual to design your VN embedding solution

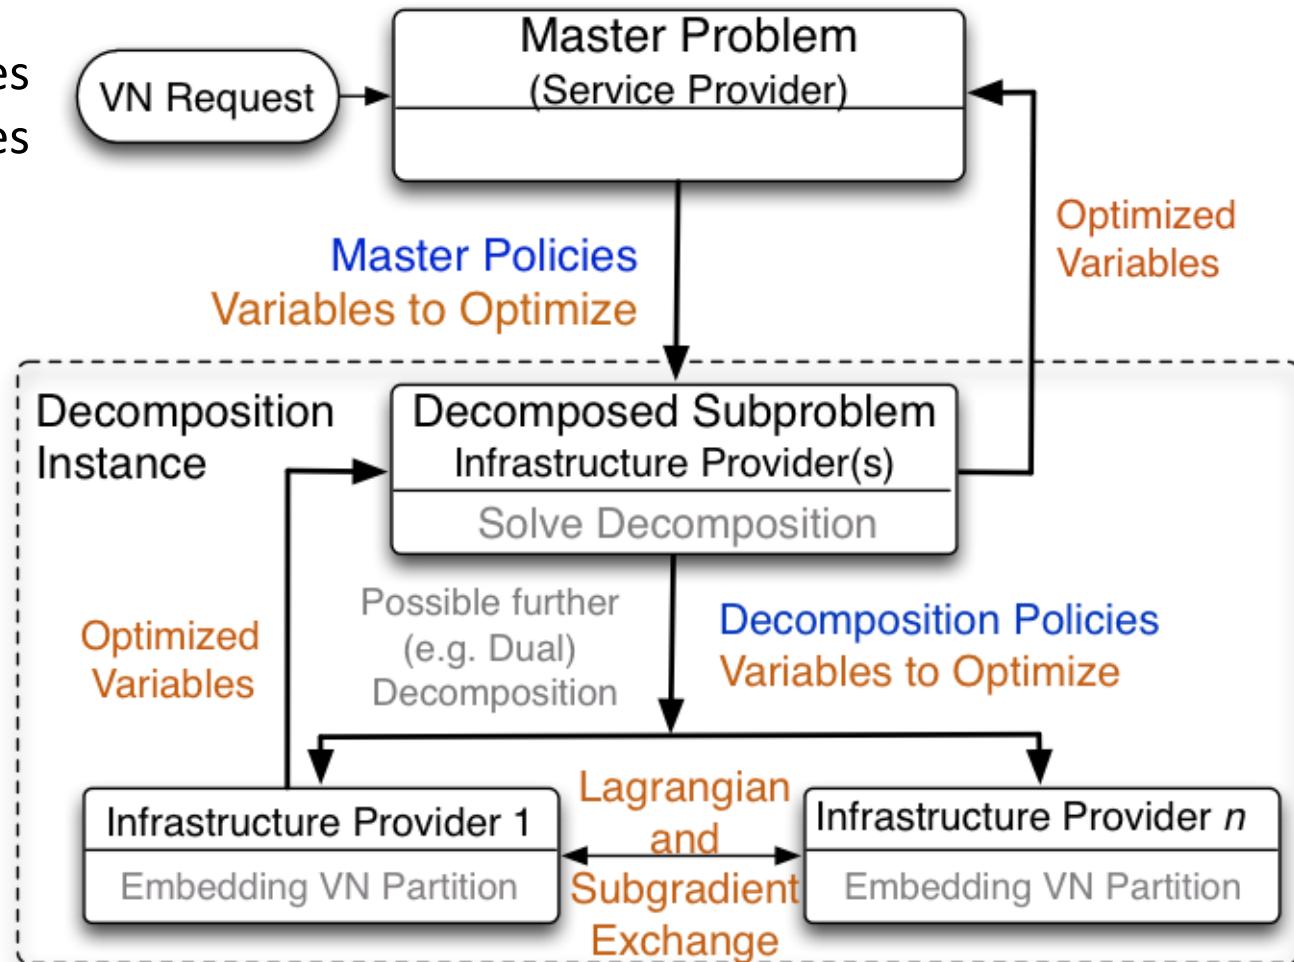
Service Provider instantiates according to its policies



Infrastructure Providers
Solve it (decomposing)

Use primal and dual to design your VN embedding solution

Service Provider instantiates according to its policies



Infrastructure Providers Solve it (decomposing)

Optimal are returned to SP that releases next VN

Testing different decompositions over our VN embedding Prototype



To separate policies from
VN embedding mechanisms

Unifying VN Embedding Architecture
Identifying the invariances of the problem

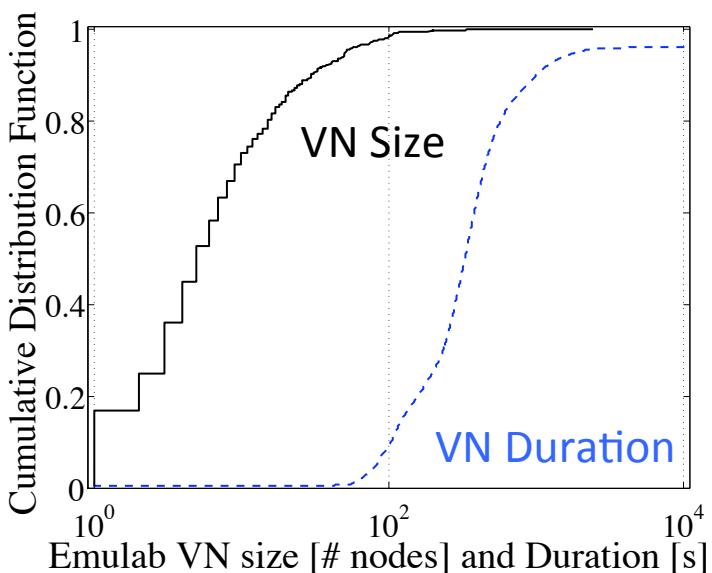
To adapt embedding
to providers' goals

Design new and subsume embedding solutions
Instantiating different decomposition policies

Enable experiments with
new embedding policies

Tradeoff analysis over simulation & prototype
Implemented over a local Linux-based testbed

8 Years of real Emulab VN requests used in our CPLEX Simulations



62,000 Emulab requests '01 to '09

Virtual Topologies

Range of VN size [1, 100]

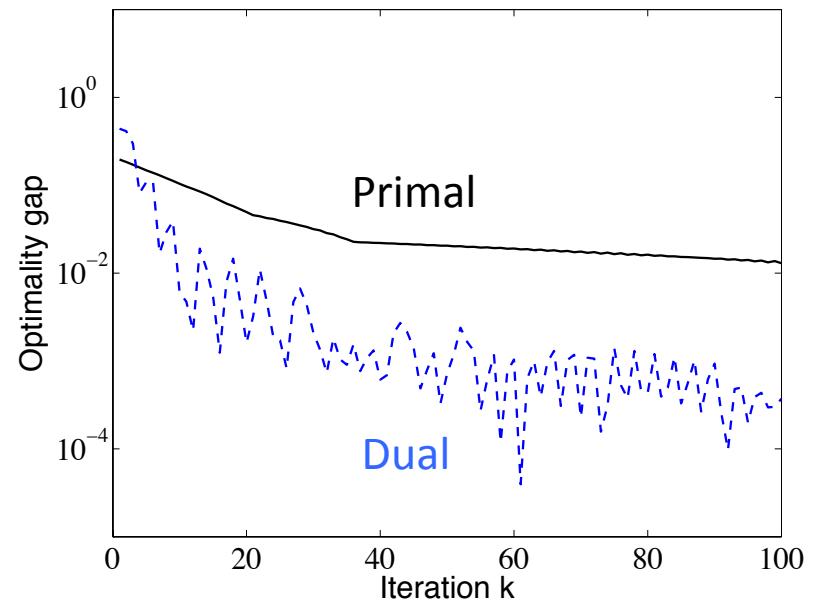
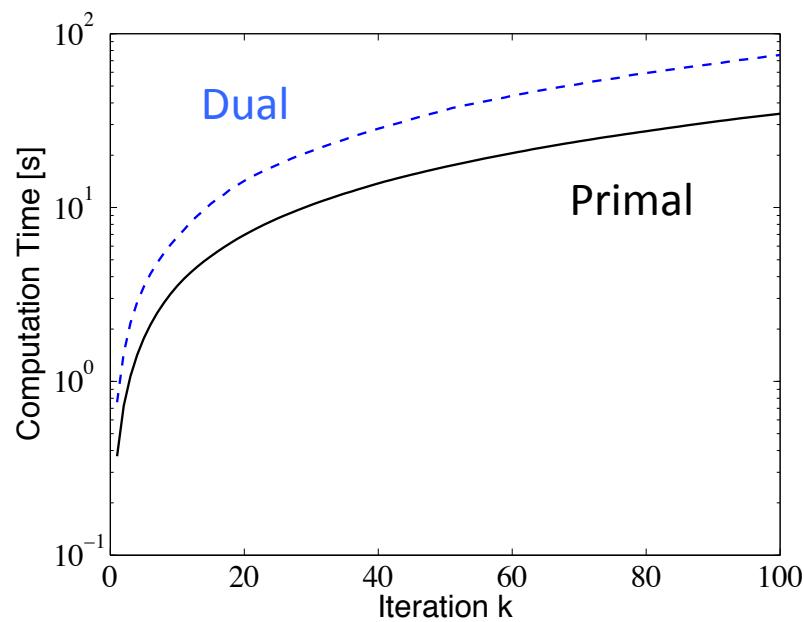
Physical Topologies

Waxman and Barabasi-Albert with BRITE

Performance Metrics

Computational Time, Optimality

Primal has lower computation time
Dual has lower optimality gap



CPLEX-based VN Embedding Simulator

We release our VINEA testbed for VINO objects instantiation

VIrtual Network Embedding Architecture



VIrtual Network Objects (VINO)

We release our VINEA testbed for VINO objects instantiation

VIrtual Network Embedding Architecture



VINEA nodes join a private overlay
after authentication and policy exchange



InPs run the VN Mapping Protocol
asynchronous consensus-based

Virtual resources are reserved
using vSwitch and Linux traffic control

VIrtual Network Objects (VINO)

We release our VINEA testbed for VINO objects instantiation

Virtual Network Embedding Architecture



Virtual Network Objects (VINO)

Discovery

VINEA nodes join a private overlay
after authentication and policy exchange

VN Mapping

InPs run the VN Mapping Protocol
asynchronous consensus-based

Allocation

Virtual resources are reserved
using vSwitch and Linux traffic control

The embedding protocol runs on a private overlay

Virtual Network Embedding Architecture



Virtual Network Objects

Discovery

VINEA nodes join a private overlay after authentication and policy exchange

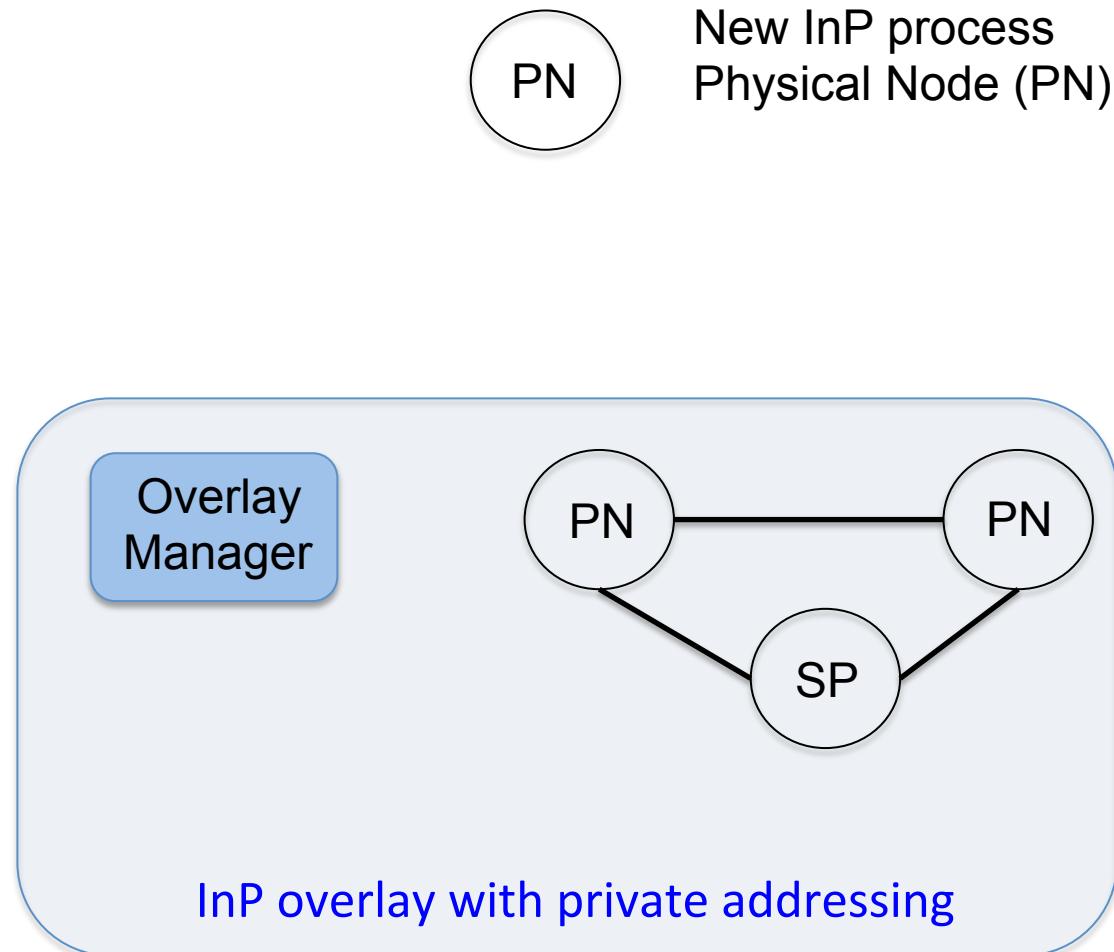
VN Mapping

InPs run the VN Mapping Protocol asynchronous consensus-based

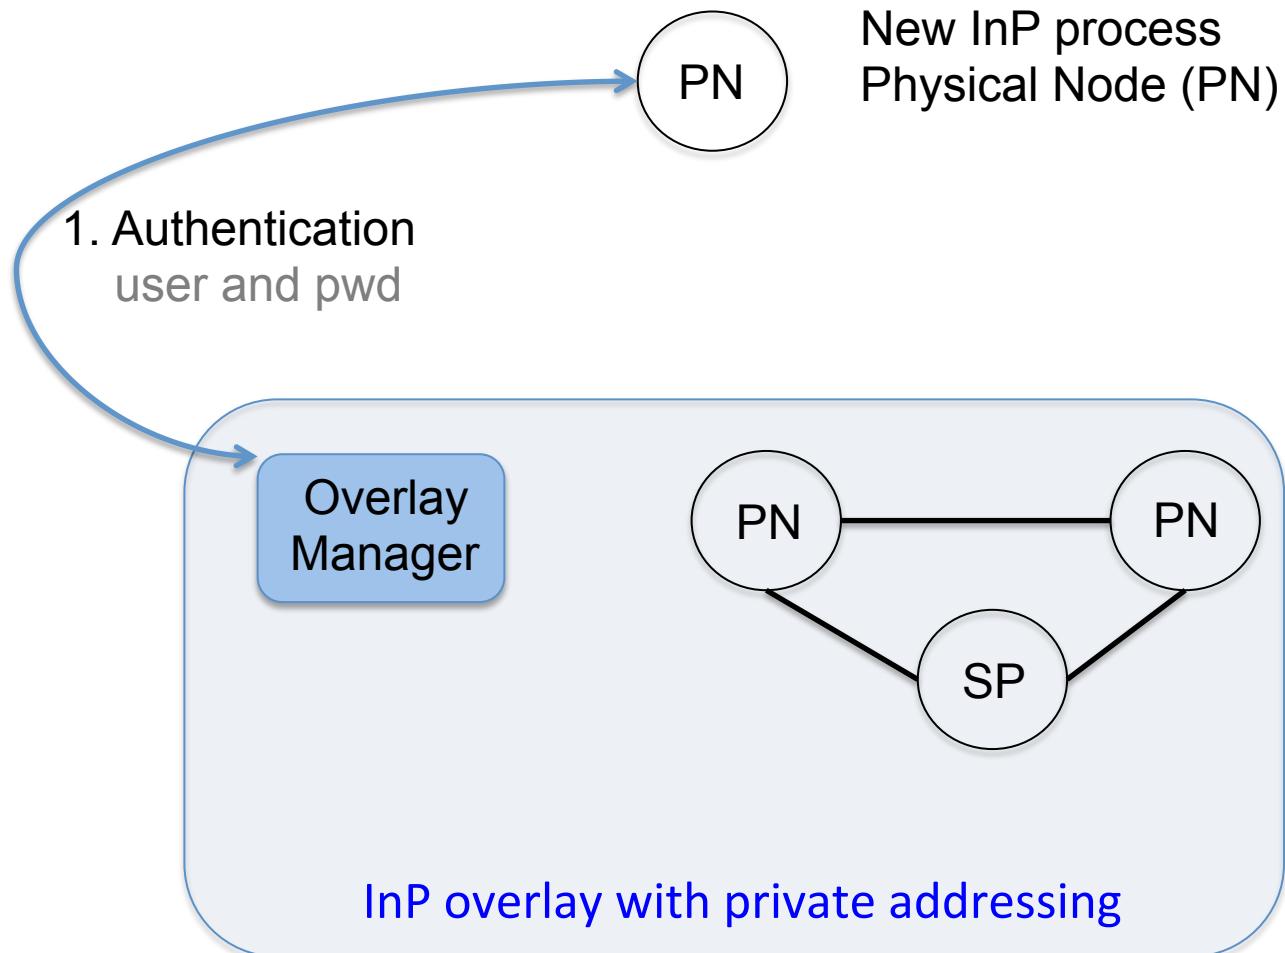
Allocation

Virtual resources are reserved using vSwitch and Linux traffic control

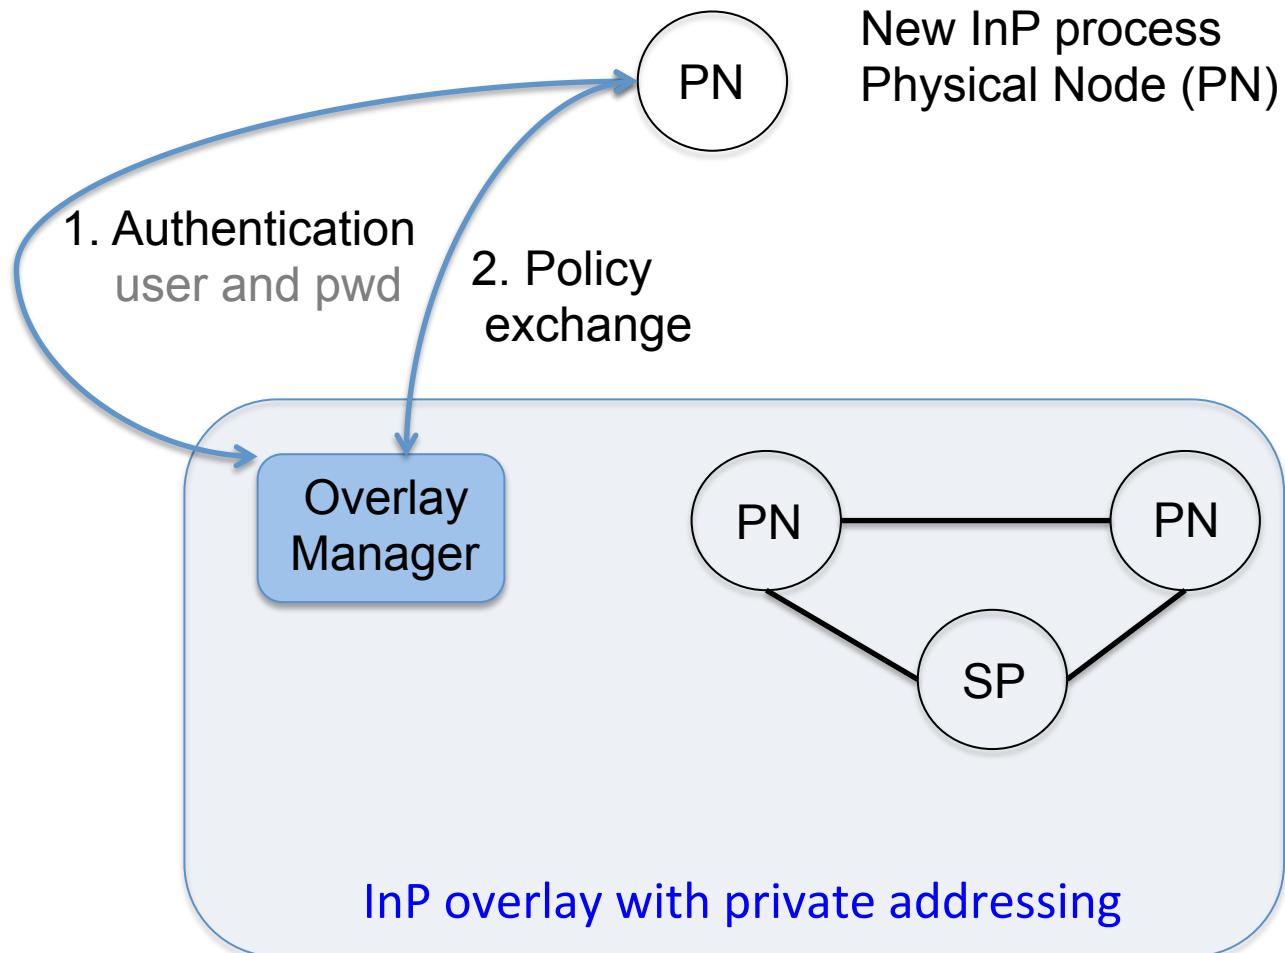
Resource Discovery: SP and InPs processes enroll in a private overlay



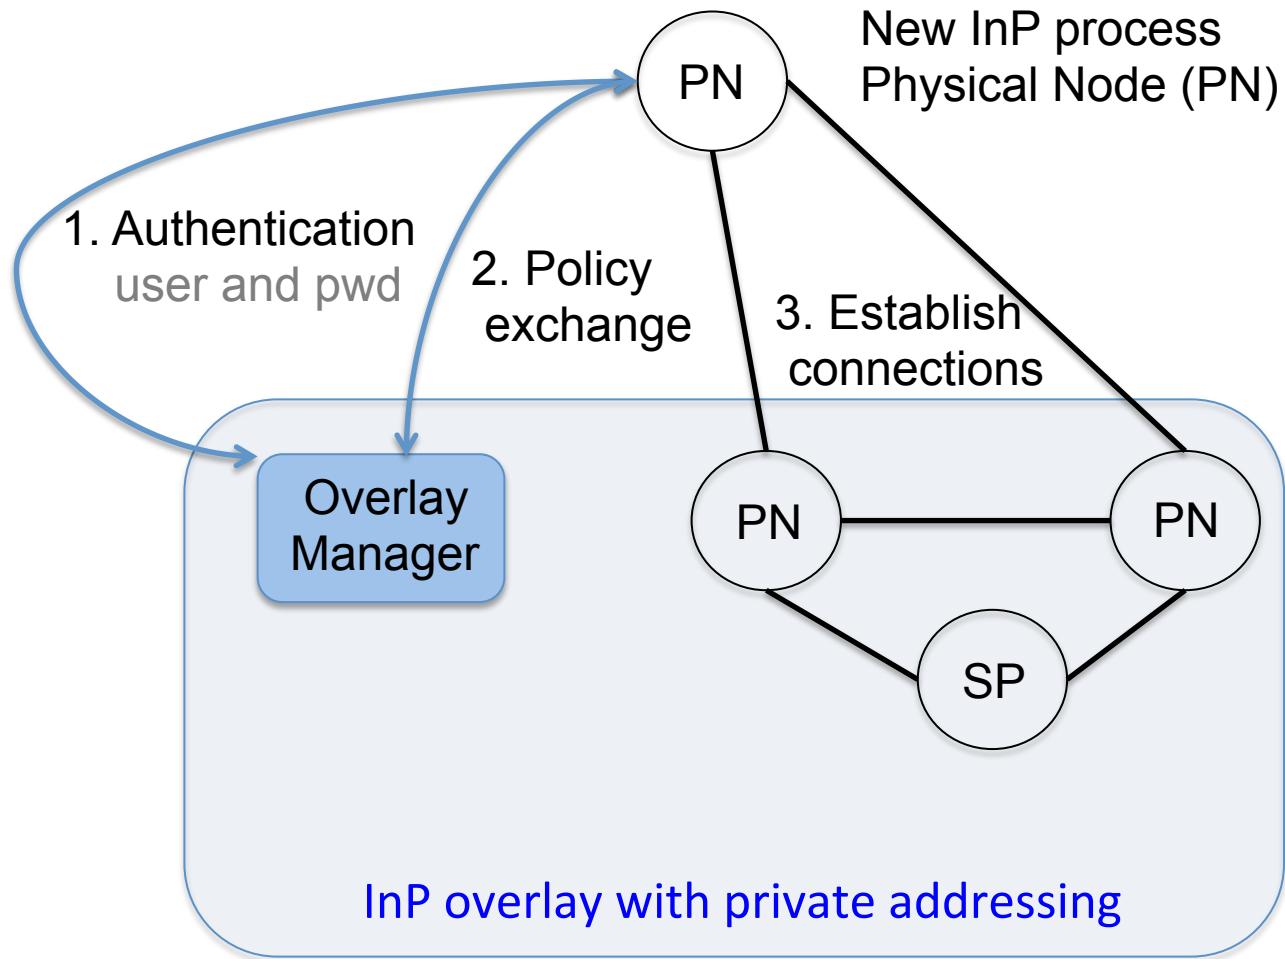
Resource Discovery: SP and InPs processes enroll in a private overlay



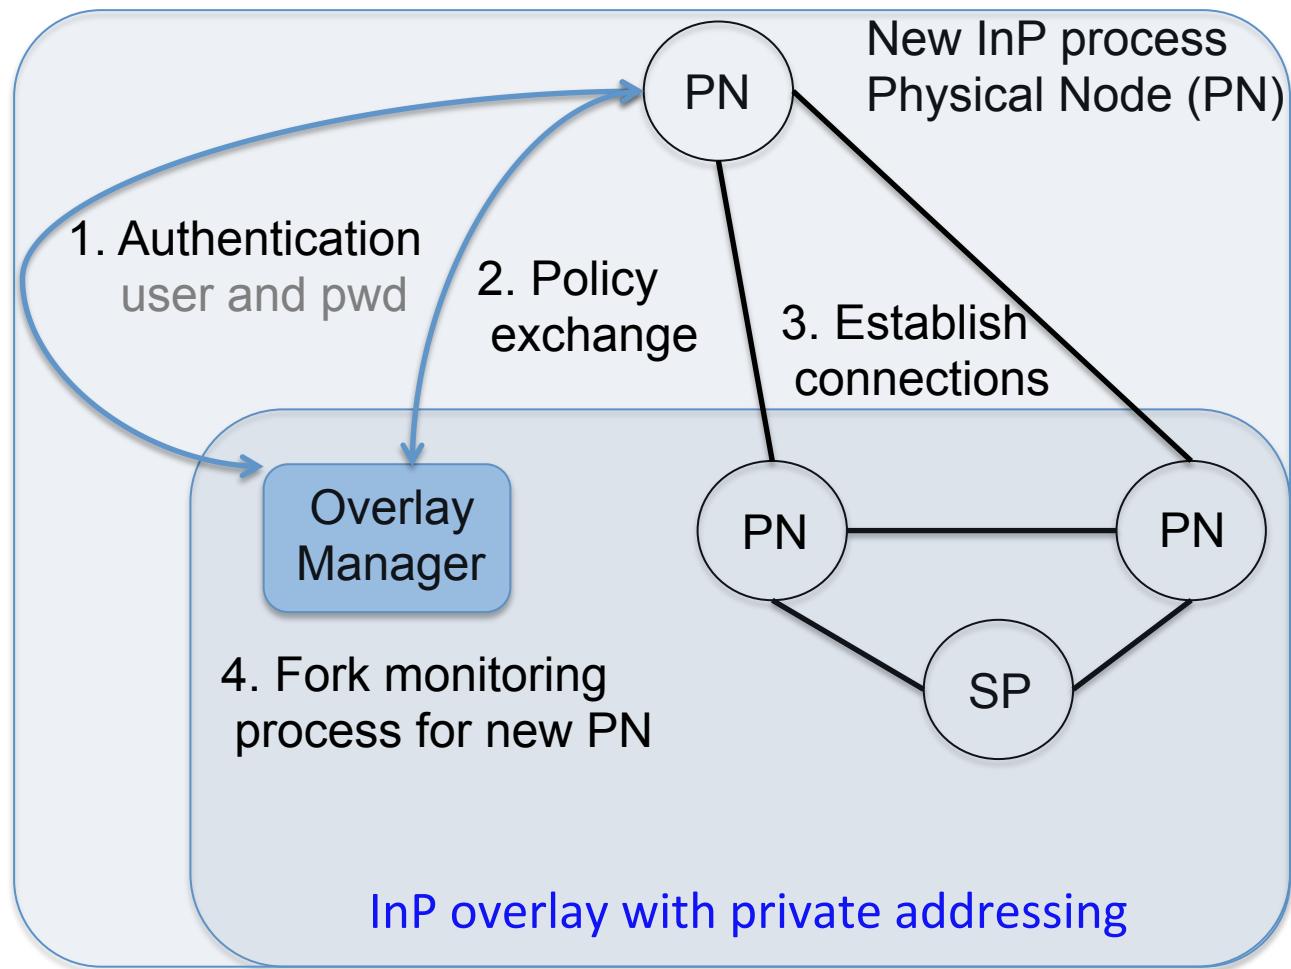
Resource Discovery: SP and InPs processes enroll in a private overlay



Resource Discovery: SP and InPs processes enroll in a private overlay



Resource Discovery: SP and InPs processes enroll in a private overlay



VN Mapping: physical nodes decide which hosts paying the lowest price

Virtual Network Embedding Architecture



Virtual Network Objects

Discovery

VINEA nodes join a private overlay after authentication and policy exchange

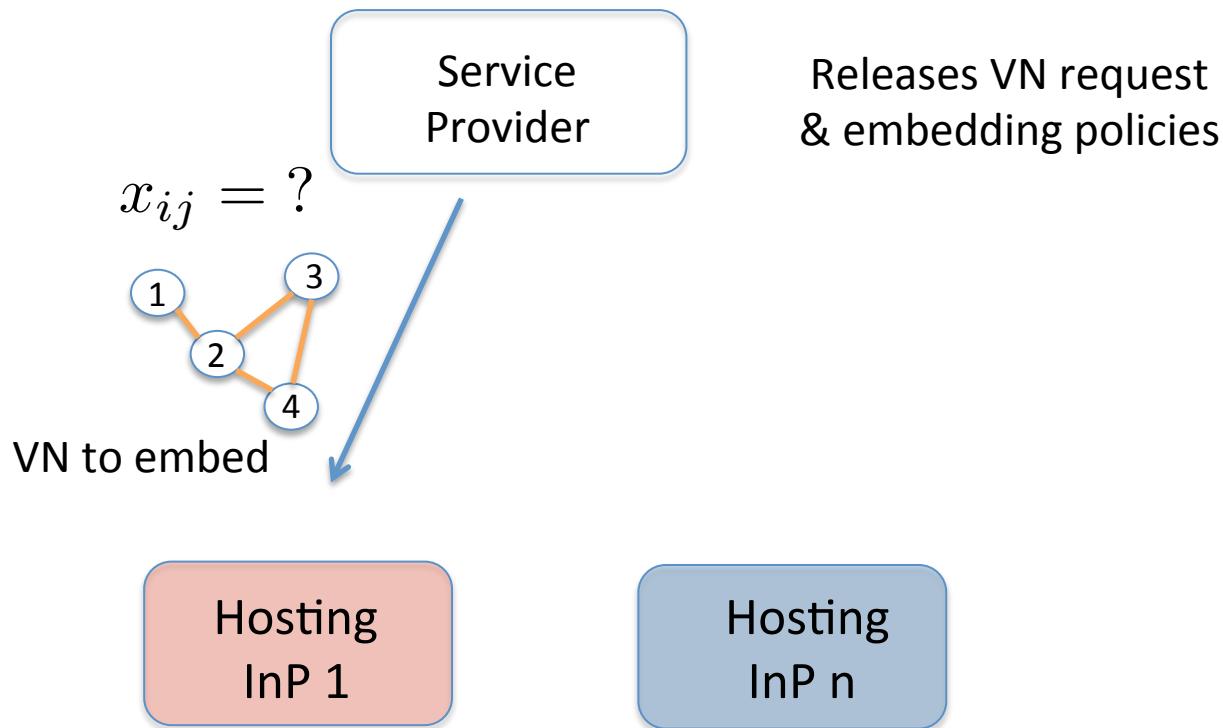
VN Mapping

InPs run the VN Mapping Protocol asynchronous consensus-based

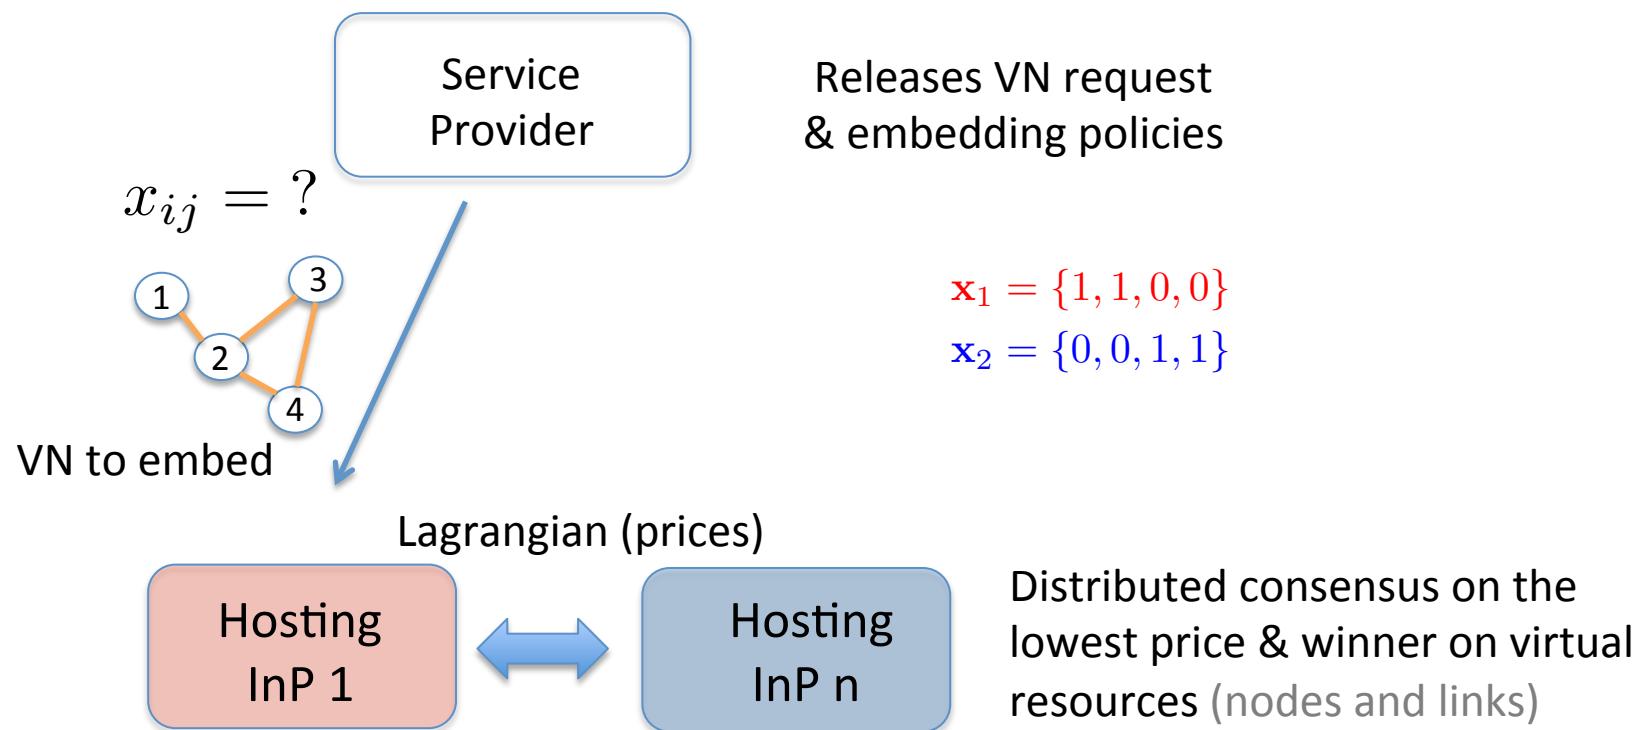
Allocation

Virtual resources are reserved using vSwitch and Linux traffic control

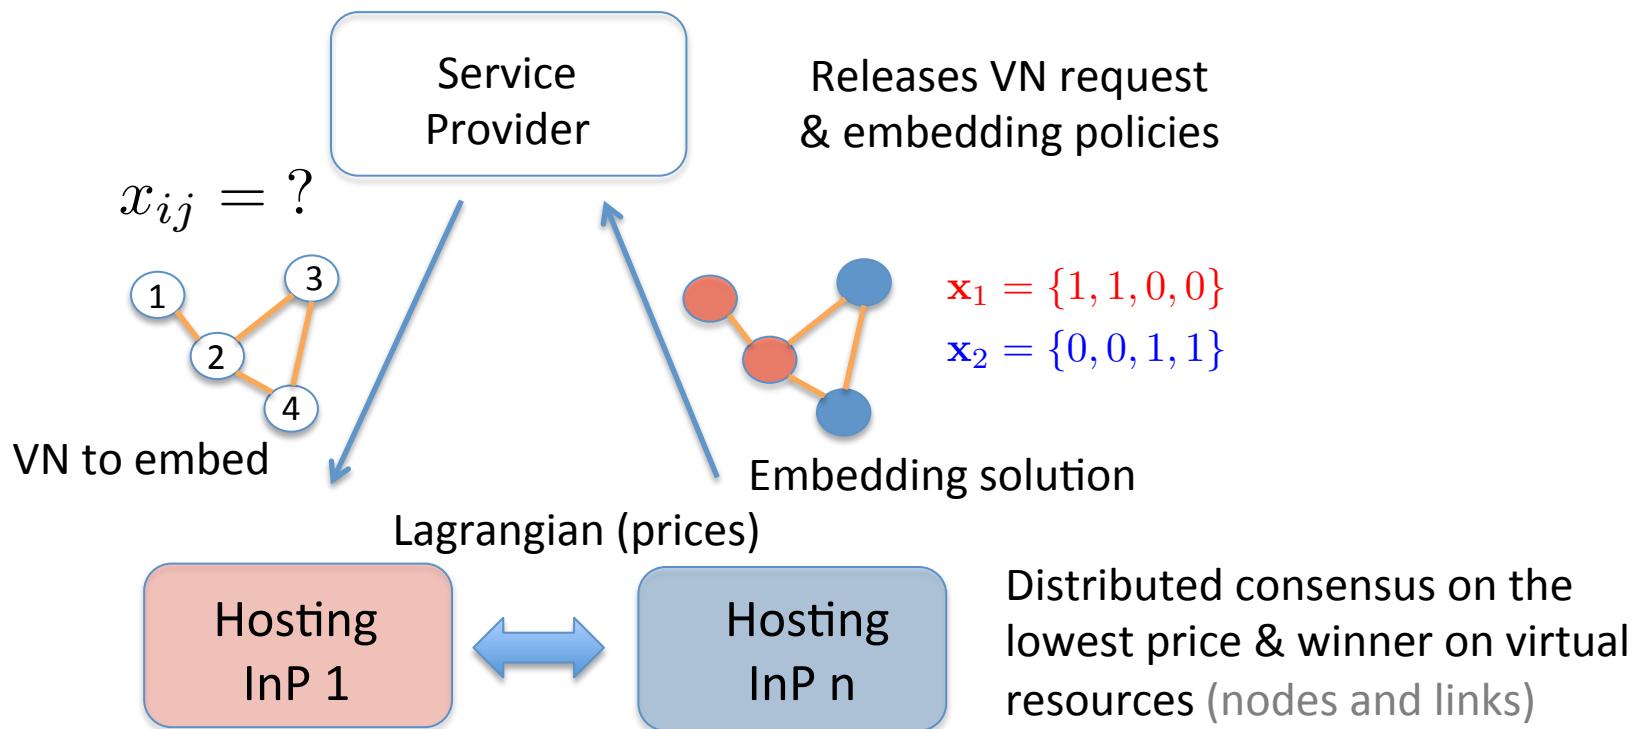
VN Mapping: physical nodes decide which hosts paying the lowest price



VN Mapping: physical nodes decide which hosts paying the lowest price



VN Mapping: physical nodes decide which hosts paying the lowest price



Allocation: User-space vHost, Kernel vSwitches & BW is reserved with Linux tc

Virtual Network Embedding Architecture



Virtual Network Objects

Discovery

VINEA nodes join a private overlay after authentication and policy exchange

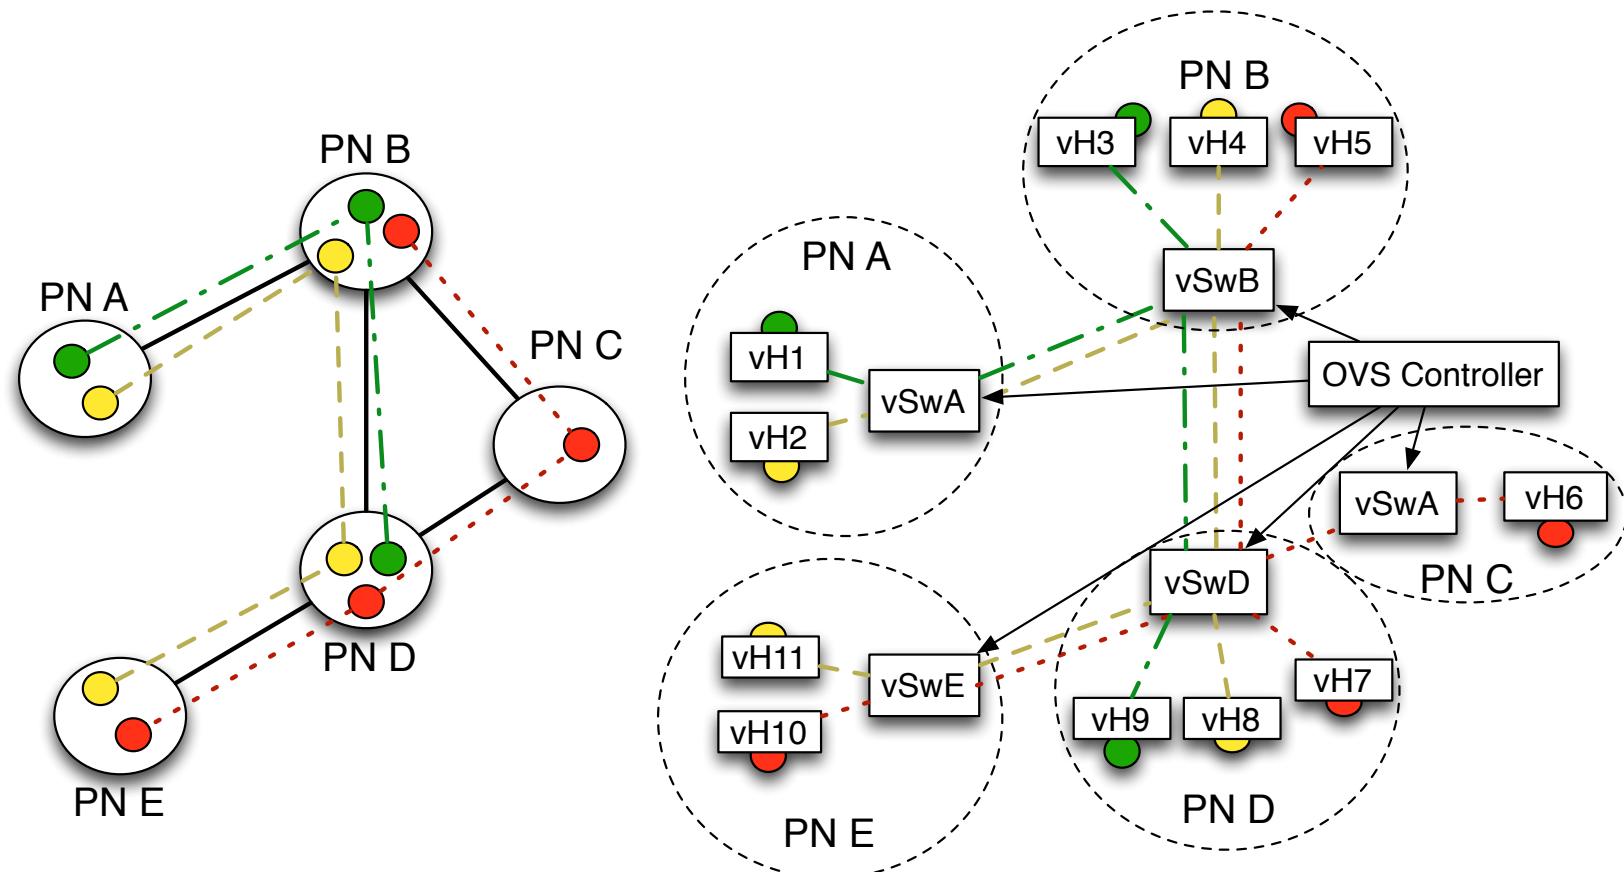
VN Mapping

InPs run the VN Mapping Protocol asynchronous consensus-based

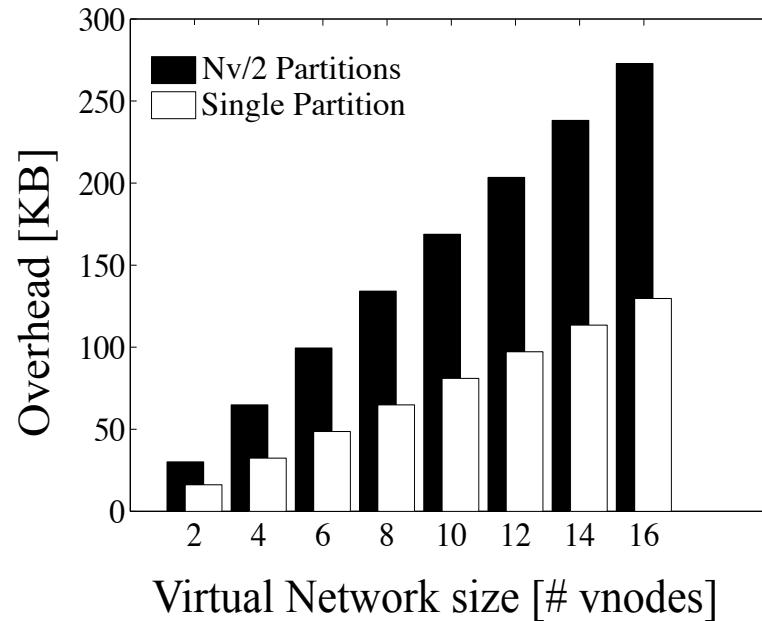
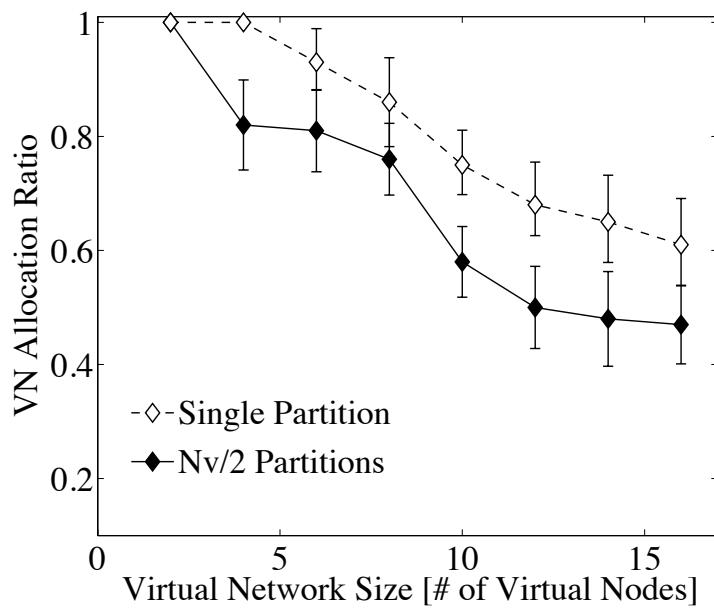
Allocation

Virtual resources are reserved using vSwitch and Linux traffic control

Allocation: User-space vHost, Kernel vSwitches & BW is reserved with Linux tc



Partitioning lowers the allocation ratio and increases the network overhead



VN embedding Prototype

Existing solutions [14] require VN partitioning

What to remember from this talk?



Ad-hoc protocols
prevent adaptation

What to remember from this talk?



Ad-hoc protocols
prevent adaptation



Decomposition is a useful tool for distributed
cloud algorithm design and tradeoff analysis

What to remember from this talk?



Ad-hoc protocols
prevent adaptation



Decomposition is a useful tool for distributed
cloud algorithm design and tradeoff analysis

Chianti Merlot Chardonnay



Use our VINEA testbed to test
your own embedding policies
<http://csr.bu.edu/vinea>

A Decomposition-Based Architecture for Distributed Virtual Network Embedding

Flavio Esposito



Ibrahim Matta



Thank you

Computer
Science
Research

Virtual
Network
Embedding
Architecture

<http://csr.bu.edu/vinea>

Boston
University