

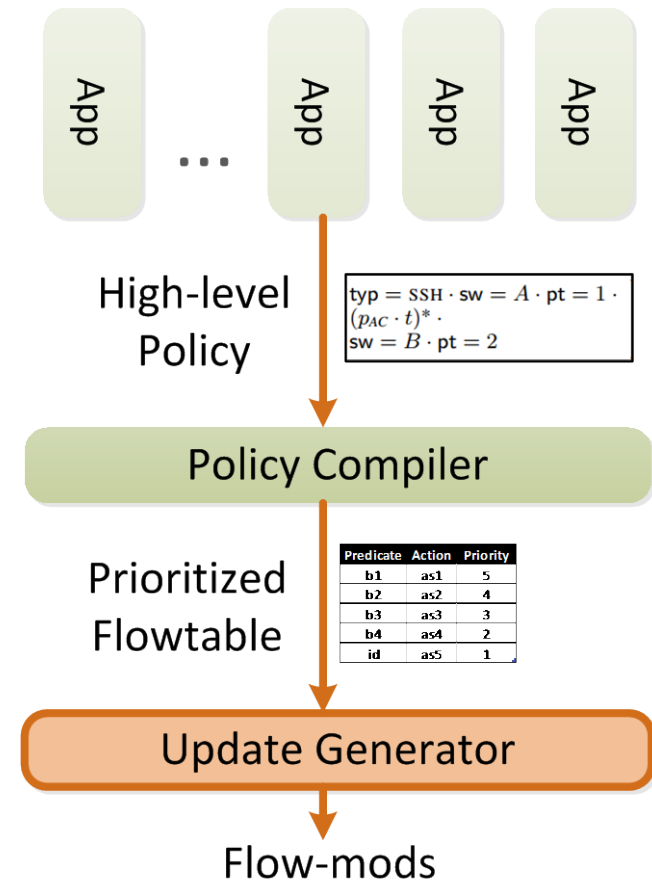
Xitao Wen, Chunxiao Diao, Xun Zhao, Yan Chen, Li Erran Li,
Bo Yang, Kai Bu

*Northwestern University, Tsinghua University, Bell Labs Alcatel-Lucent,
Zhejiang University*

Compiling Minimum Incremental Update for Modular SDN Languages

Motivation

- Flowtable update bottleneck
 - 10s to 100s of rule edits per second
 - Full refresh of 5K entries takes minutes
- **Goal:** minimizing update size to speed up flowtable update
 - Only update the “diff”



Problem Statement

- Update rules whose content **or priority** changes
 - 3 rule adds + **2 priority updates**
- Priority updates contribute **over 90%** in average!


Pattern	Priority
<1, 2>	3
<*, 2>	2
<*, *>	1



Old



Pattern	Priority
<1, 2>	5
<2, *>	4
<1, *>	3
<*, 2>	3
<3, *>	2
<*, *>	1

New



 Unmodified fields
 Modified fields

Question: How can we minimize priority updates?

Diff with Reassigned Priority

- **Idea:** Modify priorities assigned by compiler
- **Challenges:**
 - Constraint 1:* New priority assignment MUST observe rule dependency
 - *Solution:* Minimum dependency construction
 - Constraint 2:* New priority values MUST be integers within [0, 65535]
 - *Solution:* Priority gap maintenance

Pattern	Priority
<1, 2>	3
<*, 2>	2
<*, *>	1

Old



Pattern	Priority
<1, 2>	5 -> 3
<2, *>	4 -> 2.5
<1, *>	3 -> 2
<*, 2>	3 -> 2
<3, *>	2 -> 1.5
<*, *>	1

New



3 rule edits!

Dependency Constraint

- **Constraint 1**: rule dependency
- Dependency inferred from priority value is problematic

	Pattern	Priority
A	$\langle 1, 2, * \rangle$	5
B	$\langle *, 2, 3 \rangle$	4
C	$\langle 1, *, 4 \rangle$	4
D	$\langle 1, *, 3 \rangle$	3
E	$\langle *, *, 4 \rangle$	3
F	$\langle *, *, 3 \rangle$	2

Old



	Pattern	Priority
A	$\langle 1, 2, * \rangle$	5
B	$\langle *, 2, 3 \rangle$	4
G	$\langle *, 2, 4 \rangle$	3
C	$\langle 1, *, 4 \rangle$	2
D	$\langle 1, *, 3 \rangle$	3
E	$\langle *, *, 4 \rangle$	1
F	$\langle *, *, 3 \rangle$	2



New

Dependency Constraint

- Minimum dependency can be inferred from rule patterns

	Pattern	Priority
A	$\langle 1, 2, * \rangle$	5
B	$\langle *, 2, 3 \rangle$	4
C	$\langle 1, *, 4 \rangle$	4
D	$\langle 1, *, 3 \rangle$	3
E	$\langle *, *, 4 \rangle$	3
F	$\langle *, *, 3 \rangle$	2

Old





	Pattern	Priority
A	$\langle 1, 2, * \rangle$	5
B	$\langle *, 2, 3 \rangle$	4
G	$\langle *, 2, 4 \rangle$	3
C	$\langle 1, *, 4 \rangle$	2
D	$\langle 1, *, 3 \rangle$	3
E	$\langle *, *, 4 \rangle$	1
F	$\langle *, *, 3 \rangle$	2

New

Dependency Constraint

- With minimum dependency graph, one can always generate minimum-size flowtable update if priority value is continuous


	Pattern	Priority
A	<1, 2, *>	5
B	<*, 2, 3>	4
C	<1, *, 4>	4
D	<1, *, 3>	3
E	<*, *, 4>	3
F	<*, *, 3>	2



	Pattern	Priority
A	<1, 2, *>	5
B	<*, 2, 3>	4
G	<*, 2, 4>	3 → 4.5
C	<1, *, 4>	2 → 4
D	<1, *, 3>	3
E	<*, *, 4>	1 → 3
F	<*, *, 3>	2

Take-away: Minimum dependency helps eliminate priority updates

How to obtain minimum dependency

 Restored from prioritized flowtable after compilation

- Incurs complicated header space computation

■ Constructed along with compilation

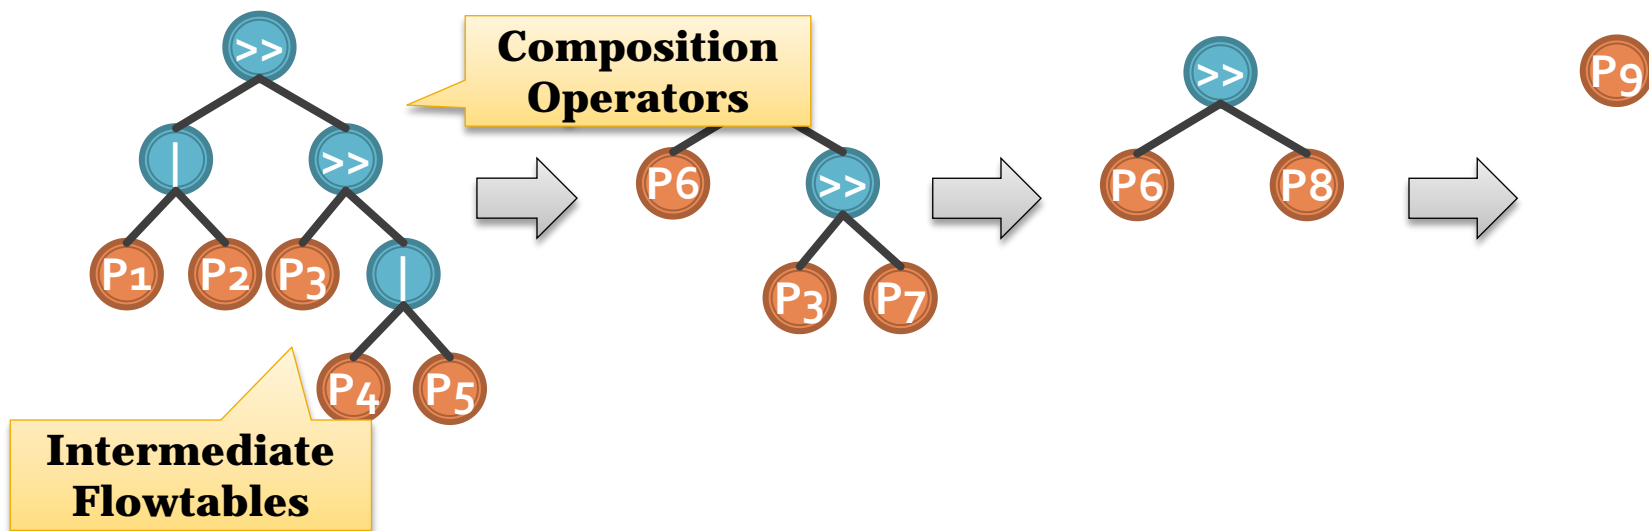
- Rule dependency can be recursively inferred from policy composition process



- Incurs little additional overhead over compilation

Incremental Dependency Construction

- Recursive construction of minimum dependency
 - Extend intermediate flowtables with dependency graph
 - Keep track of dependency during compilation
 - Parallel composition
 - Sequential composition



Incremental Dependency Construction

- Infer dependency for parallel composition
 1. Graph cross-product
 2. Intersection tests
 3. Special treatment for compiler-specific data structure
- Sequential composition is similar except for the actions of the first operands


Maintaining Priority Value Distribution

- **Constraint 2**: discrete priority values
 - Integers ranging [0-65535] for OpenFlow
 - If new rule is inserted between adjacent priority values, we have to shift existing rules to make room for them
- **Problem Statement**
 - Assign priority values for priority levels
 - Objective: minimize the estimation of priority shifts
- **Online strategy**
 - Undetermined future policy update sequence

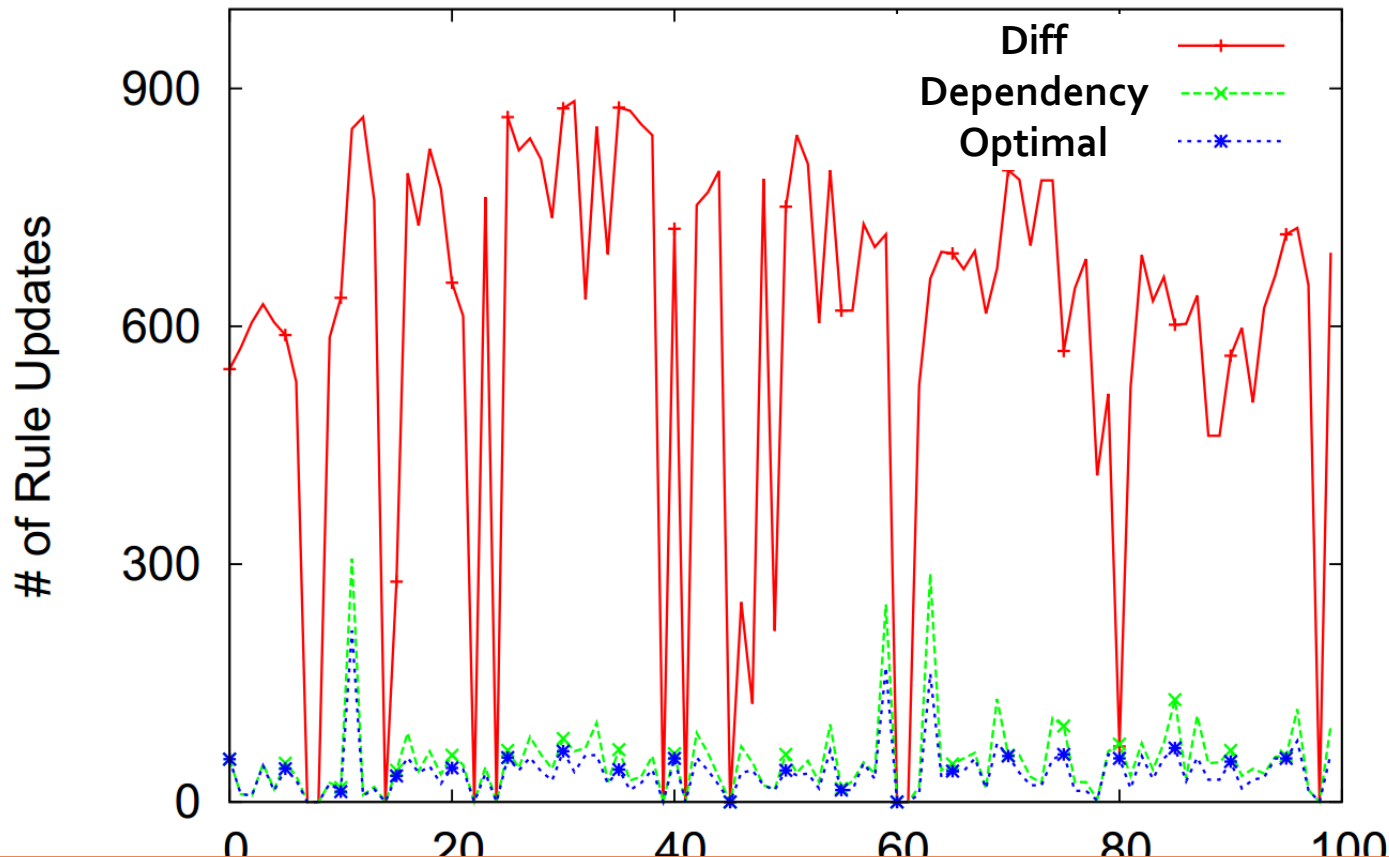
Pattern	Priority
<1, 2>	5 -> 3
<2, *>	4 -> 2.5
<1, *>	3 -> 2
<*, 2>	3 -> 2
<3, *>	2 -> 1.5
<*, *>	1

K-factor strategy

- Key idea: proactively maintains the ratio between max priority gap and min priority gap
- Initiation
 - Distribute priority levels evenly on [0-65535]
- Invariance
 - Keep lengths of all gaps between $[1/k, k] * \text{mean length}$, where k is a parameter
- Cost
 - Amortized: $O(1)$ per update



Evaluation



- *Eliminates almost all priority updates*
- *10x smaller on average compared to diff*

Conclusion

- Minimum incremental update framework comprising of
 - Minimum update generation with dependency
 - K-factor strategy for priority gaps maintenance
- Future work
 - Dependency construction algorithms generic to policy languages
 - Lower bound of priority gap maintenance

Xitao Wen

xitaowen2015@u.northwestern.edu

Question?